



AWS Serverless

Prática do Dev ao SysAdmin

Fabian Da Silva
Senior Solutions Architect, AWS Partner Network
 @fabianmartins

Oct, 14, 2020

Falando de Serverless

Porque começamos com Alien Attack?



Padrão Produtor-Consumidor



Casos de uso relacionados

- Acompanhamento de vendas em tempo-real
- Detecção tem tempo-real de eventos a partir de beacons em lojas de departamento (Marketing/promoções em near real-time)
- Tratamento de transações financeiras em tempo-real (tendências, prevenção à fraudes...)
- Processamento de eventos de dispositivos em tempo-real (IoT, Mobile)
- Monitoramento de pacientes em hospitais em tempo-real
- Rastreamento de ativos/cargas em tempo-real
- Cenário de gamificação (onde um scoreboard viria a calhar?)
- ...

Requisitos da aplicação

Segurança
(RBAC)

Armazenamento
dos dados para
uso futuro

Processamento
em tempo-real

Minimizar o TCO
com Infraestrutura



Cloud Economics

Custo

Quanto paguei neste mês com infra-estrutura?

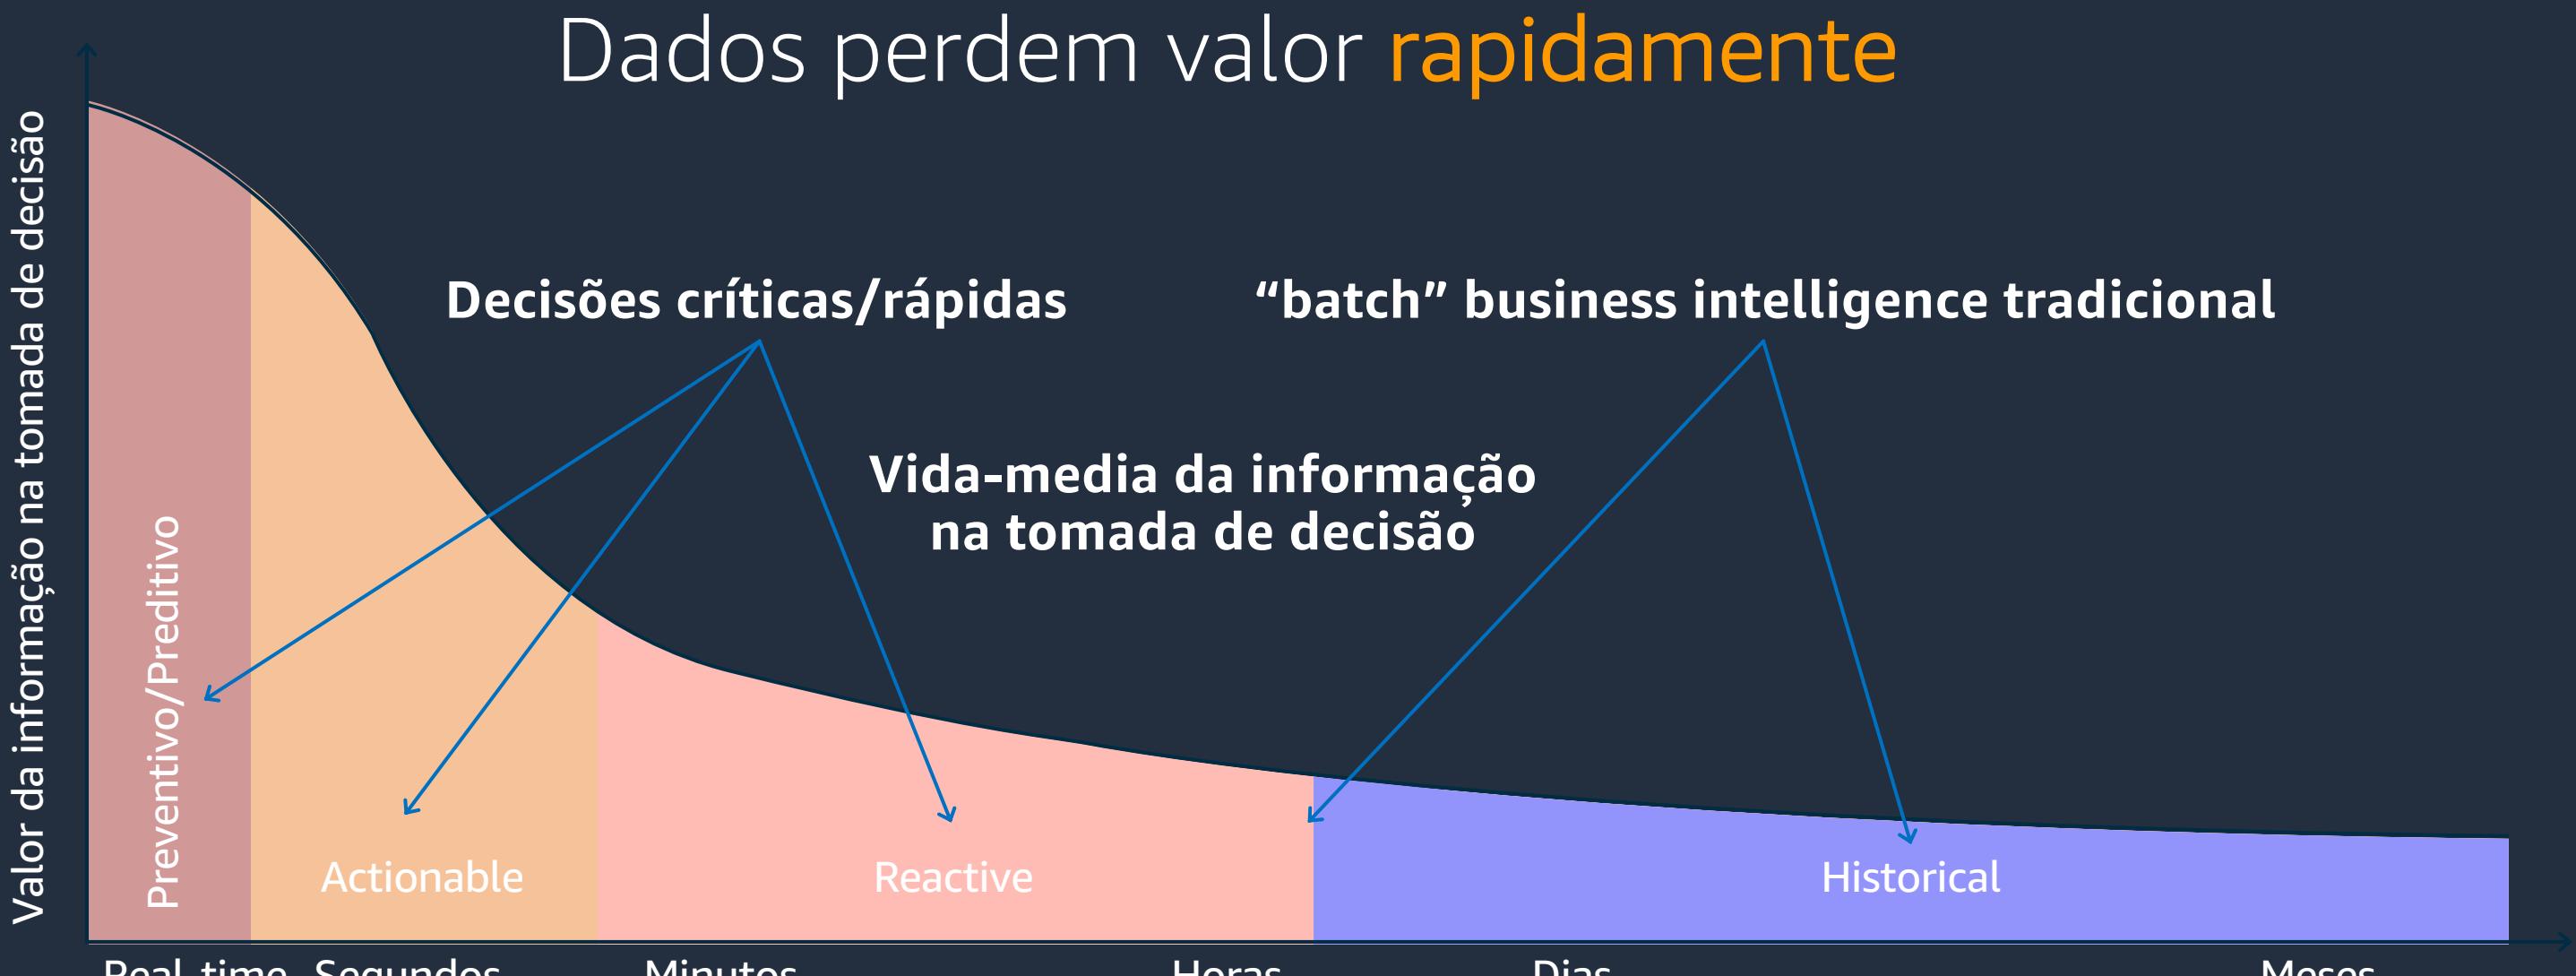
TCO (Total Cost of Ownership)

Qual o custo total para manter a minha infra-estrutura executando?

- Energia elétrica
- Rede
- Armazenamento
- Processamento
- Licenças
- Horas de suporte/manutenção
- Horas do time de Dev/Test/QA
- Custo de governança e compliance
- Quanto pago de idle?
- ... em todos os ambientes

Para comparar Cloud com on-premises, você precisa saber o TCO on-premises, e em geral não se sabe (excessiva alocação de custos em "*overhead*")

O que significa “real-time”?



Source: Perishable insights, Mike Gaultieri, Forrester

Arquitetura de alto-nível padrão



Vários caminhos...



Mas, afinal, o que é Serverless?



Sem infra-estrutura para provisionar,
Sem gerenciamento



Escalabilidade automática

Computação sem Servidores

Comunicação
Processamento
Armazenamento

Sem que eu precise gerenciar
(instalar, atualizar, escalar...)

Pague pelo valor adicionado



Altamente disponível e seguro



(Infraestrutura Virtualizada Tradicional + API + Padrões de Automação) X Engenheiros brilhantes + Camada de Abstração = Serverless

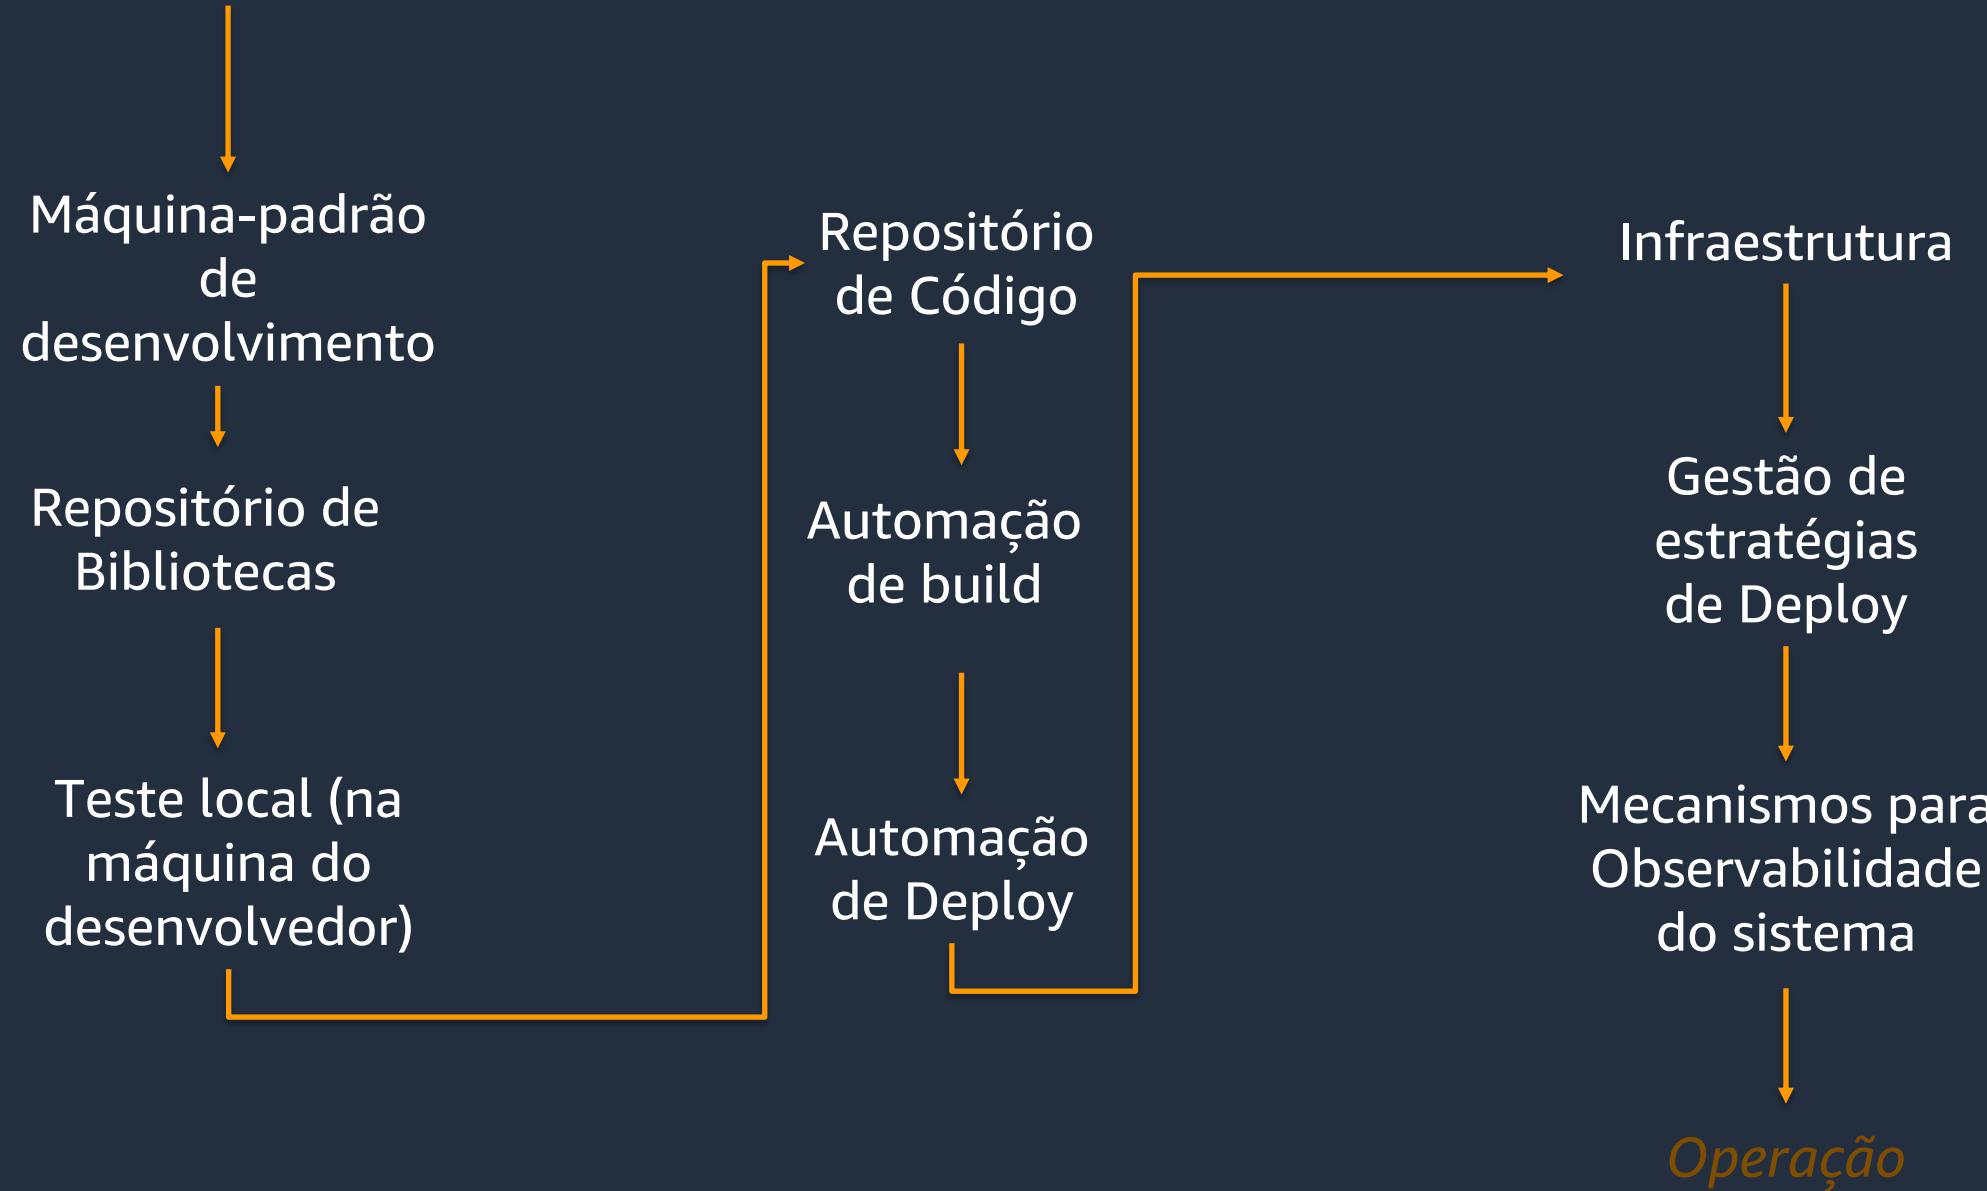
AWS Alien Attack

Como você construiria esse sistema?

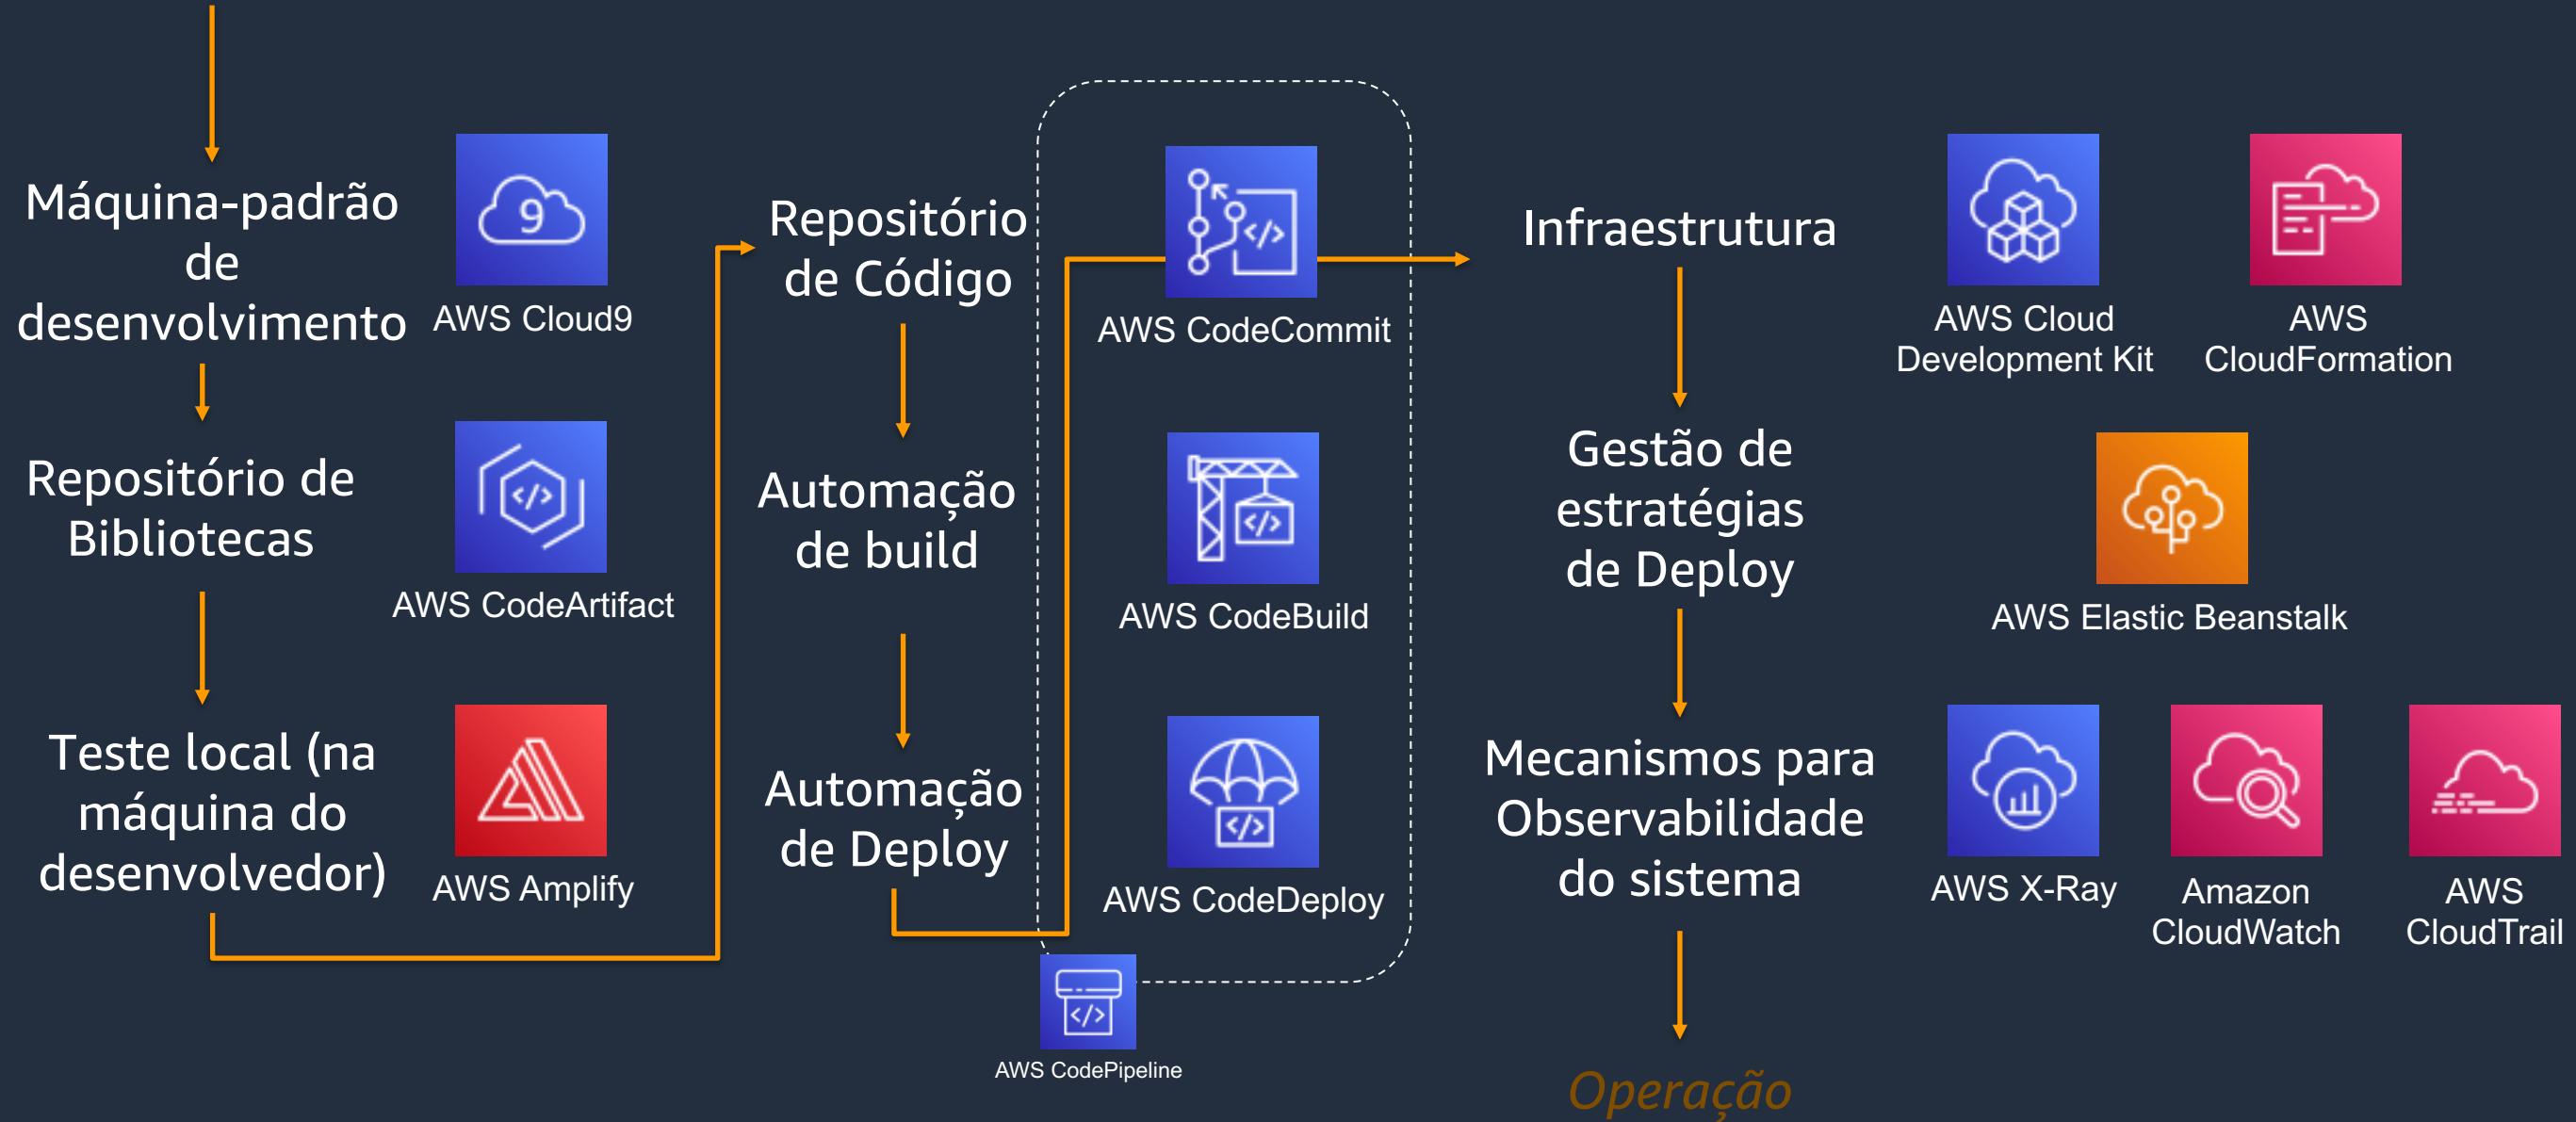
*Desenvolvimento
Infra-estrutura
Operação*

Desenvolvimento do Sistema

Todo time de desenvolvimento precisa...



Todo time de desenvolvimento precisa...



AWS Cloud9

Cloud-based integrated development environment



Programe usando o browser

IDE completa

Construa aplicações serverless

Colabore em tempo-real

AWS Cloud9

The screenshot shows the AWS Cloud9 IDE interface. On the left is a file browser with a dark theme, displaying a hierarchical file structure. The main area contains three tabs: 'config.sh', 'destroy.sh', and 'index.js'. The 'index.js' tab is active, showing a block of JavaScript code. Below the tabs is a terminal window titled 'bash - "ip-172-31-35-50"'. The terminal output shows the node version and a prompt from the environment.

```
// Copyright Amazon.com, Inc. or its affiliates. All Rights Reserved.
// SPDX-License-Identifier: MIT-0
/**
 * The purpose of this function is to retrieve session data if there are seats available for the user
 * If there is not, an error shall be returned.
 */
'use strict';

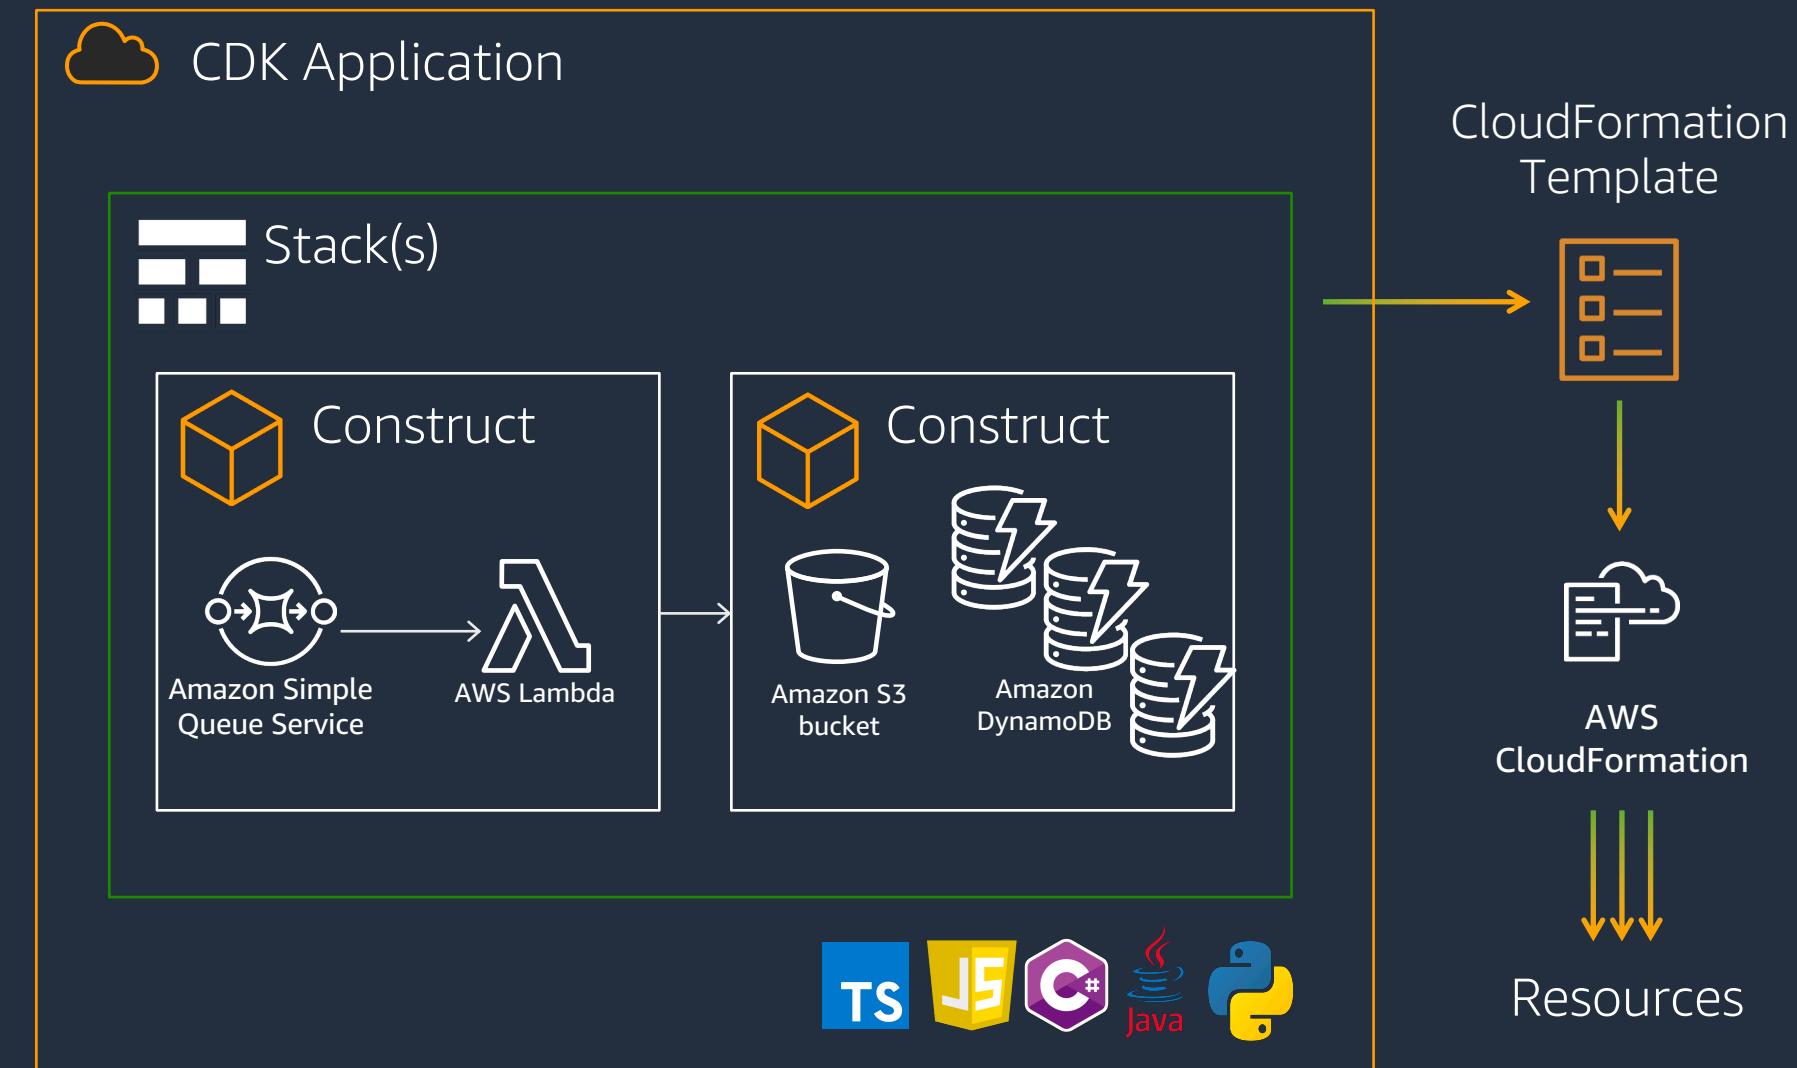
const AWS = require('aws-sdk');
const DynamoDB = new AWS.DynamoDB.DocumentClient();
const SSM = new AWS.SSM();

const readSessionFromSSM = function (callback) {
  let param = {
    "Name": process.env.SESSION_PARAMETER
  };
  SSM.getParameter(param,
    function (error, sessionParamResponse) {
      if (error) {
        let errorMessage = "Error reading from SSM";
        console.log(errorMessage);
        console.log(error);
        let responseError = new Error(errorMessage);
        responseError.code = "ErrorReadingSSM";
        responseError.details = error;
        callback(responseError,500);
      } else {
        let sessionData = null;
        try {
          sessionData = JSON.parse(sessionParamResponse.Parameter.Value);
          callback(null, sessionData);
        } catch (error) {
          let errorMessage = "Error parsing session data from SSM";
          console.log(errorMessage);
          console.log(error);
          let responseError = new Error(errorMessage);
          responseError.code = "ErrorReadingFromSSM";
          responseError.details = error;
          console.log(sessionData);
        }
      }
    }
  );
}

1:1 JavaScript Spaces: 4
```

```
bash - "ip-172-31-35-50" ×
Now using node v10.15.3 (npm v6.4.1)
Admin:~/environment $
```

AWS Cloud Development Kit (CDK)



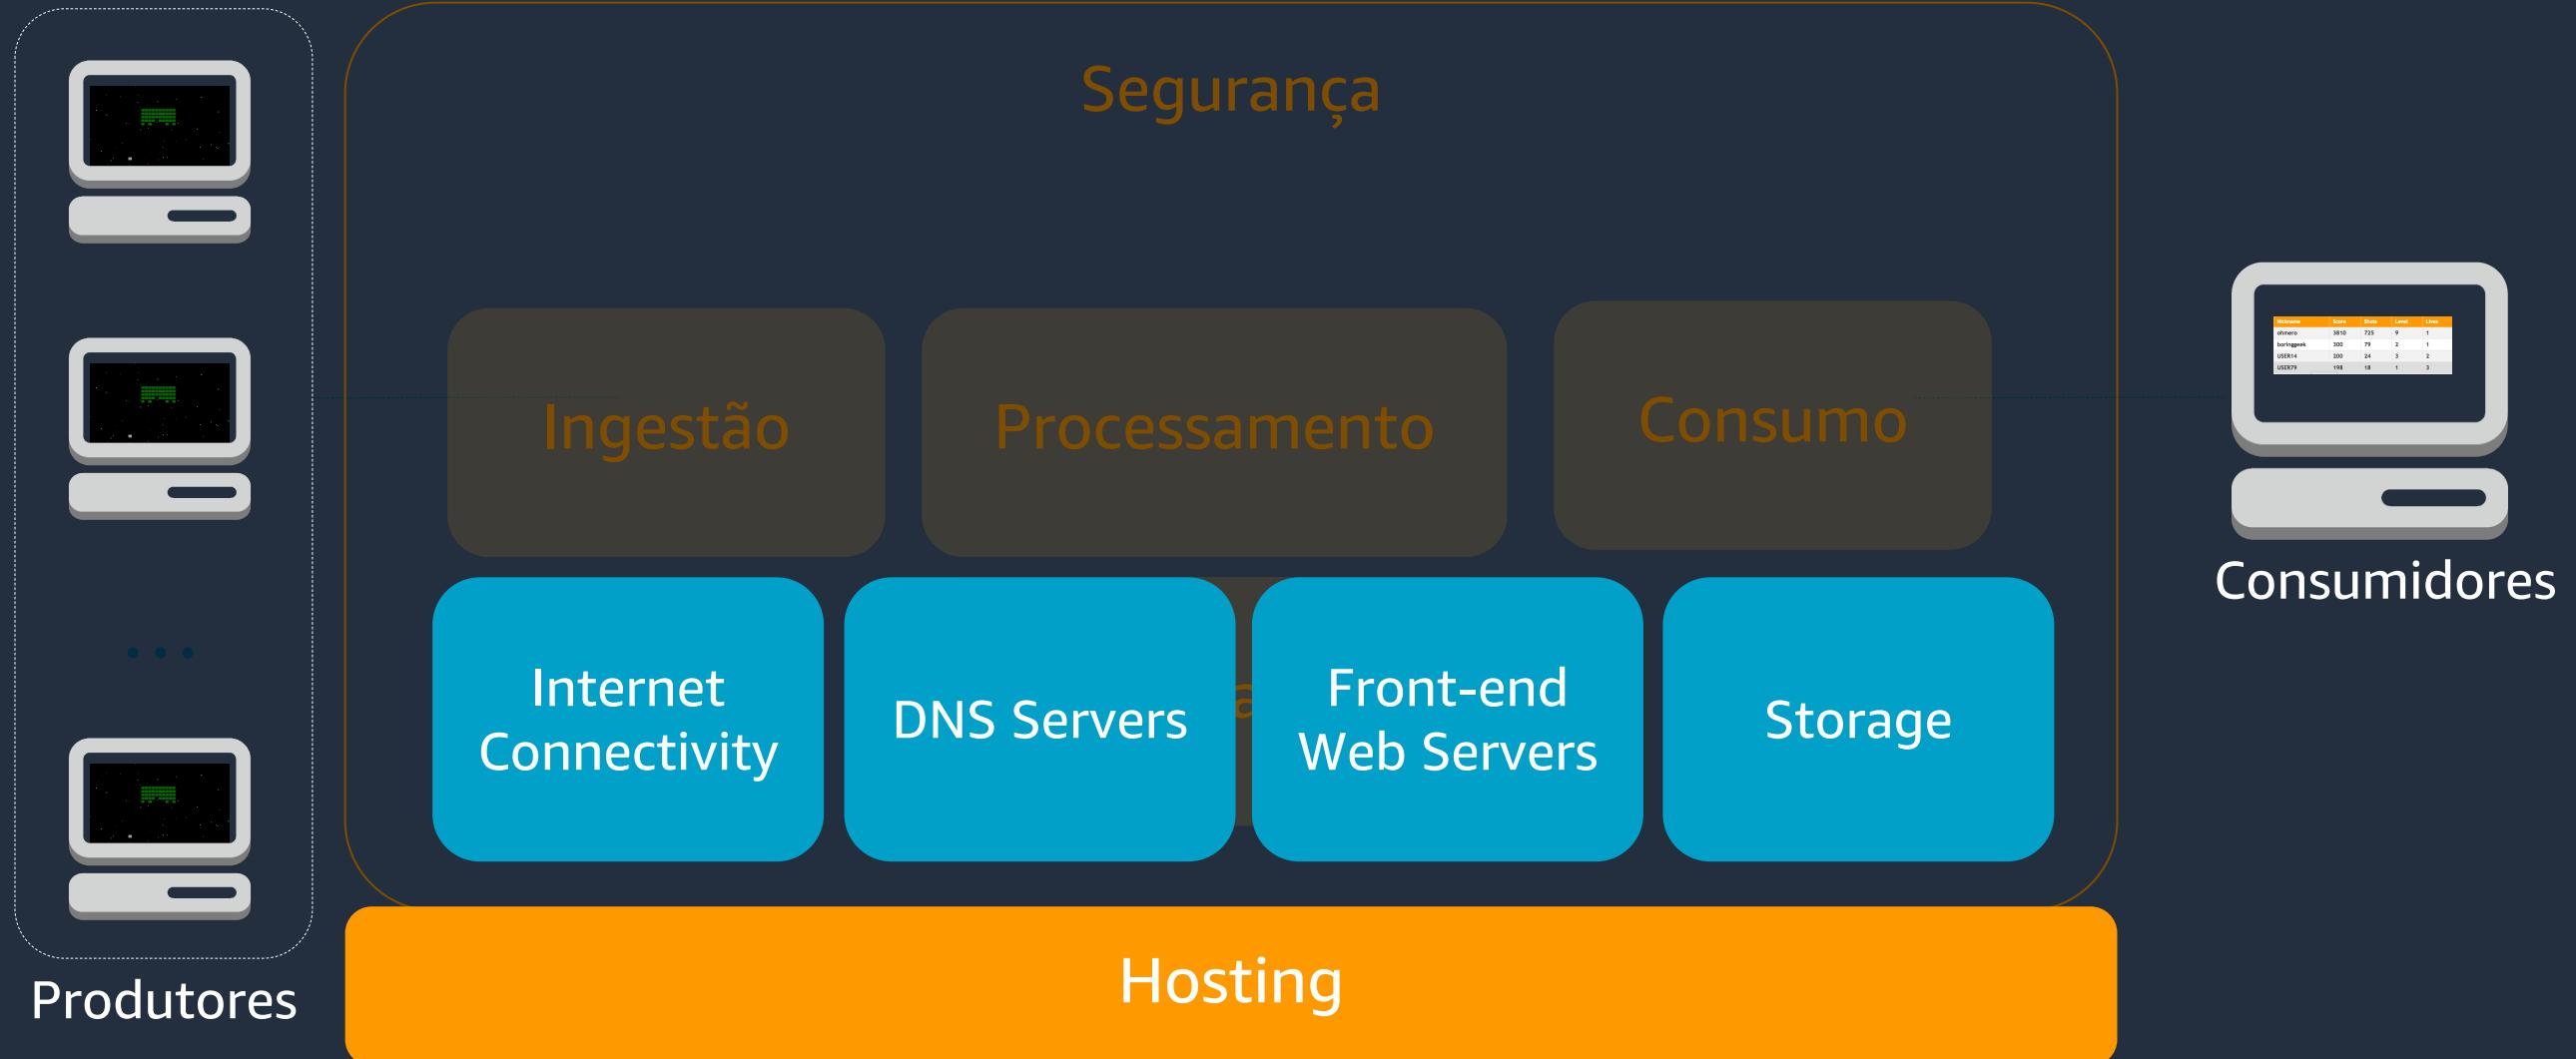
CDK Sample Code

```
export class ECSStack extends cdk.Stack {  
  
    public readonly theAPI: apigV1.RestApi  
    public readonly cluster: ecs.Cluster  
    public readonly loadBalancer: elbv2.NetworkLoadBalancer  
    public readonly vpcLink: apigV1.VpcLink  
  
    constructor(scope: cdk.App,  
               id: string, props?: cdk.StackProps)  
    {  
        super(scope, id, props);  
  
        const vpc =  
            new ec2.Vpc(this, 'ecs_Vpc', { maxAzs: 2 });  
  
        this.cluster = new ecs.Cluster(this, 'ecsCluster', {  
            clusterName: "ecs-Cluster",  
            ContainerInsights: true,  
            vpc  
        });  
    }  
}
```

```
    this.loadBalancer  
        = new elbv2.NetworkLoadBalancer(this, 'ecsLB',  
            { vpc, internetFacing: false });  
  
    this.vpcLink =  
        new apigV1.VpcLink( this, 'ecsVpcLink',  
            { targets: [this.loadBalancer] });  
  
    this.theAPI =  
        new apigV1.RestApi(this, "RestApi",  
            { restApiName: `API ${this.stackName}`,  
            endpointTypes:  
                [apigV1.EndpointType.REGIONAL],  
            }  
        );  
  
    const configResource =  
        this.addResource(`config`, 8080);  
    const updateScoreResource =  
        this.addResource(`updatescore`, 8081);  
  
    . . .  
}
```

Definindo a Infra-Estrutura

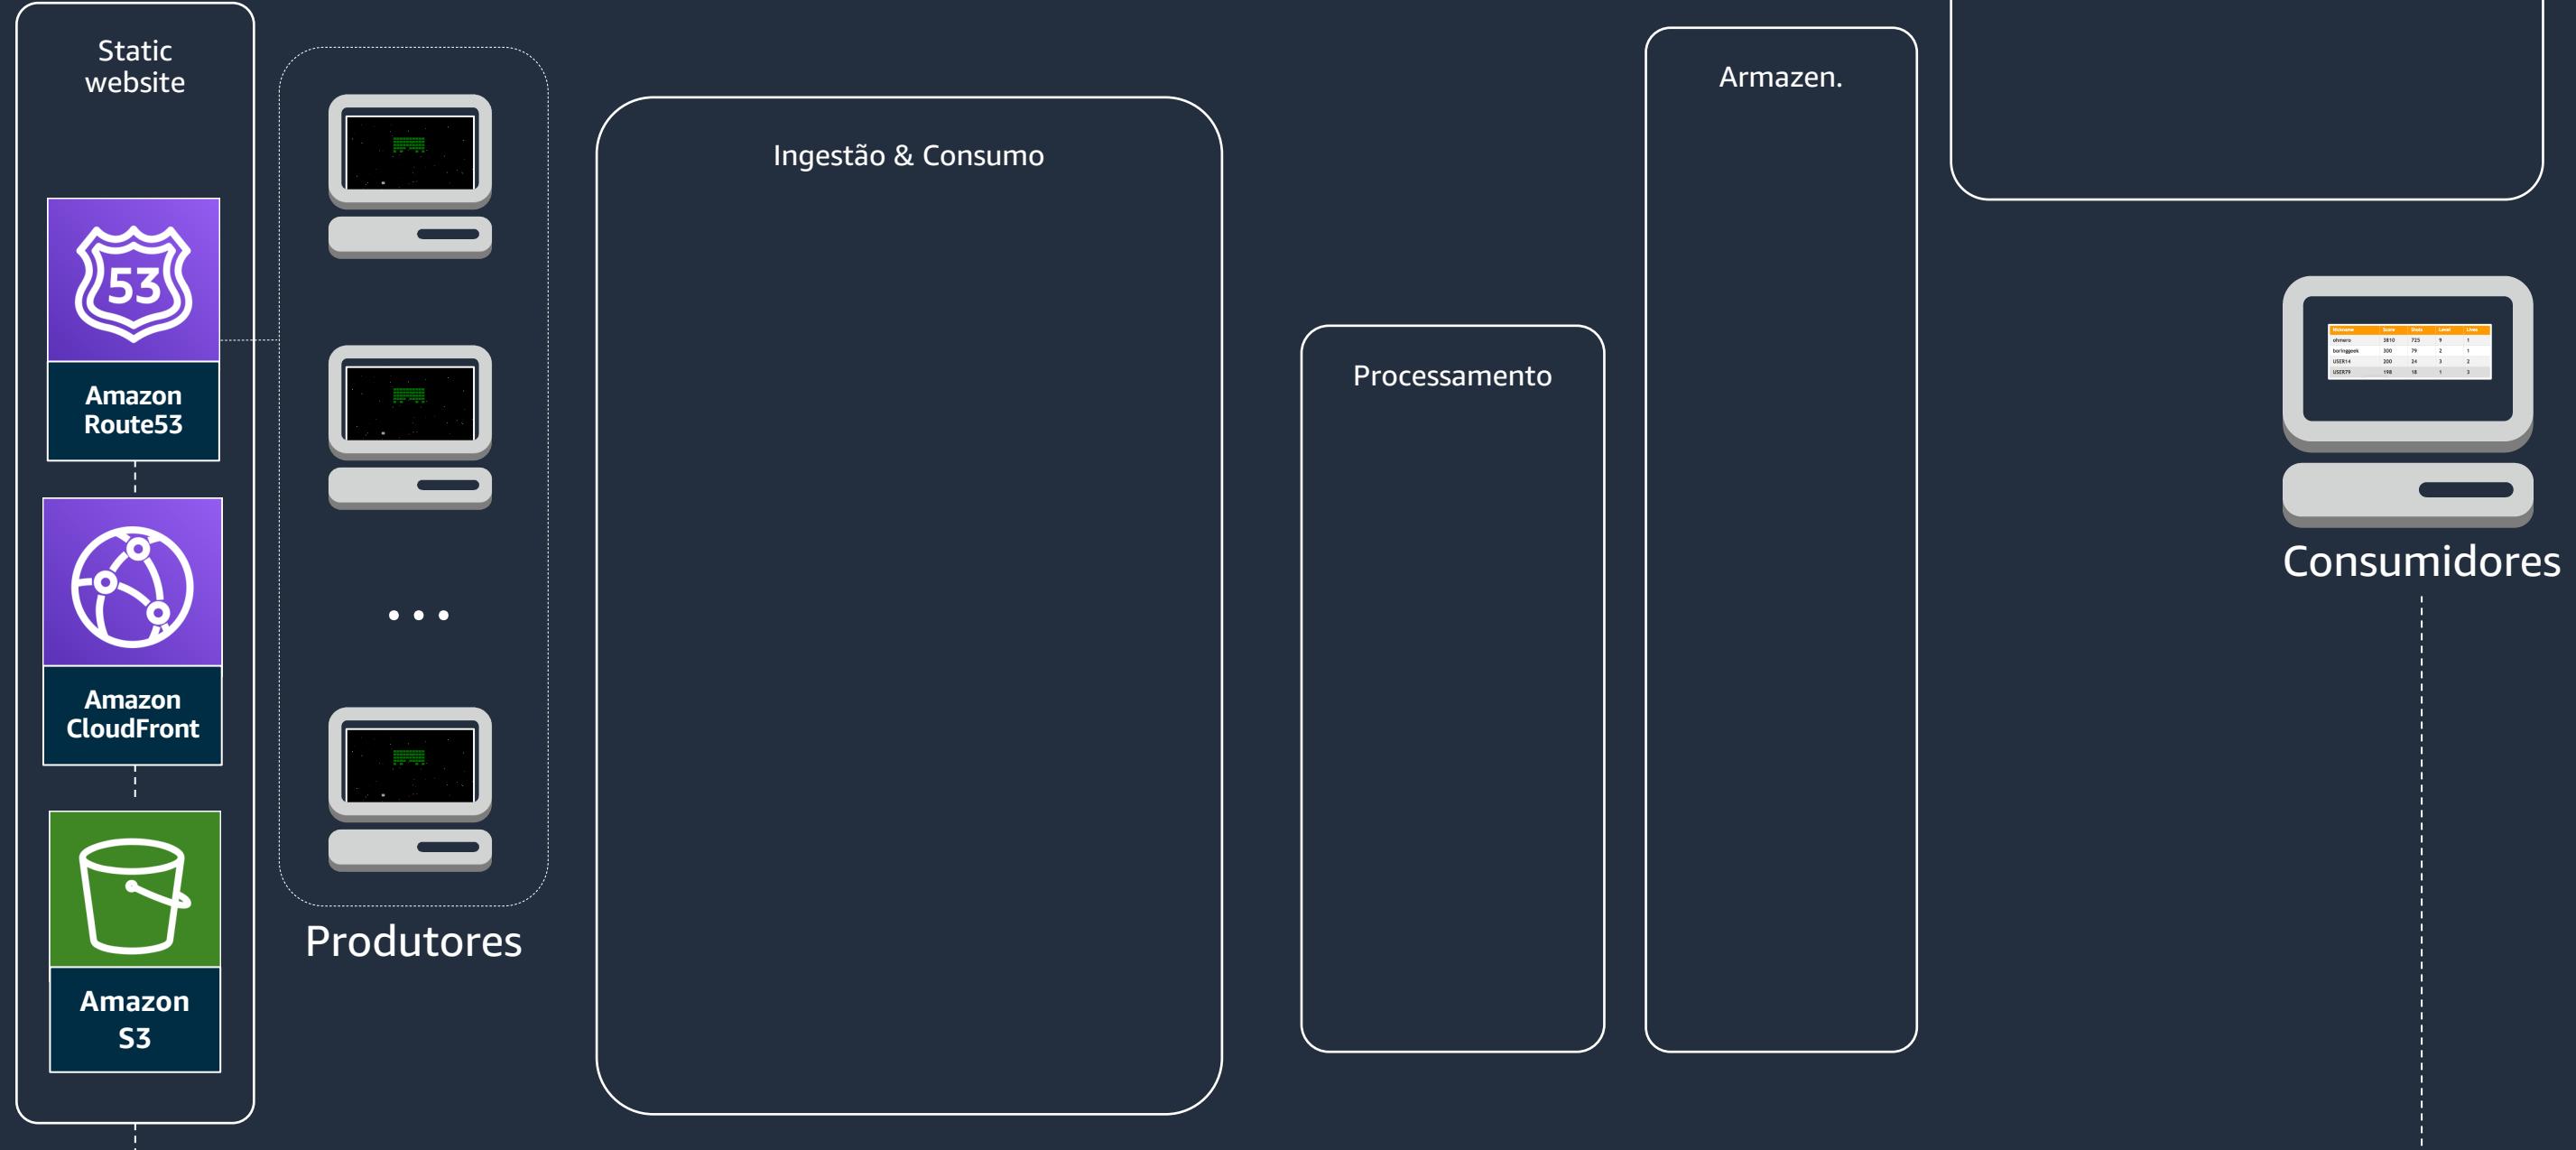
Website hosting



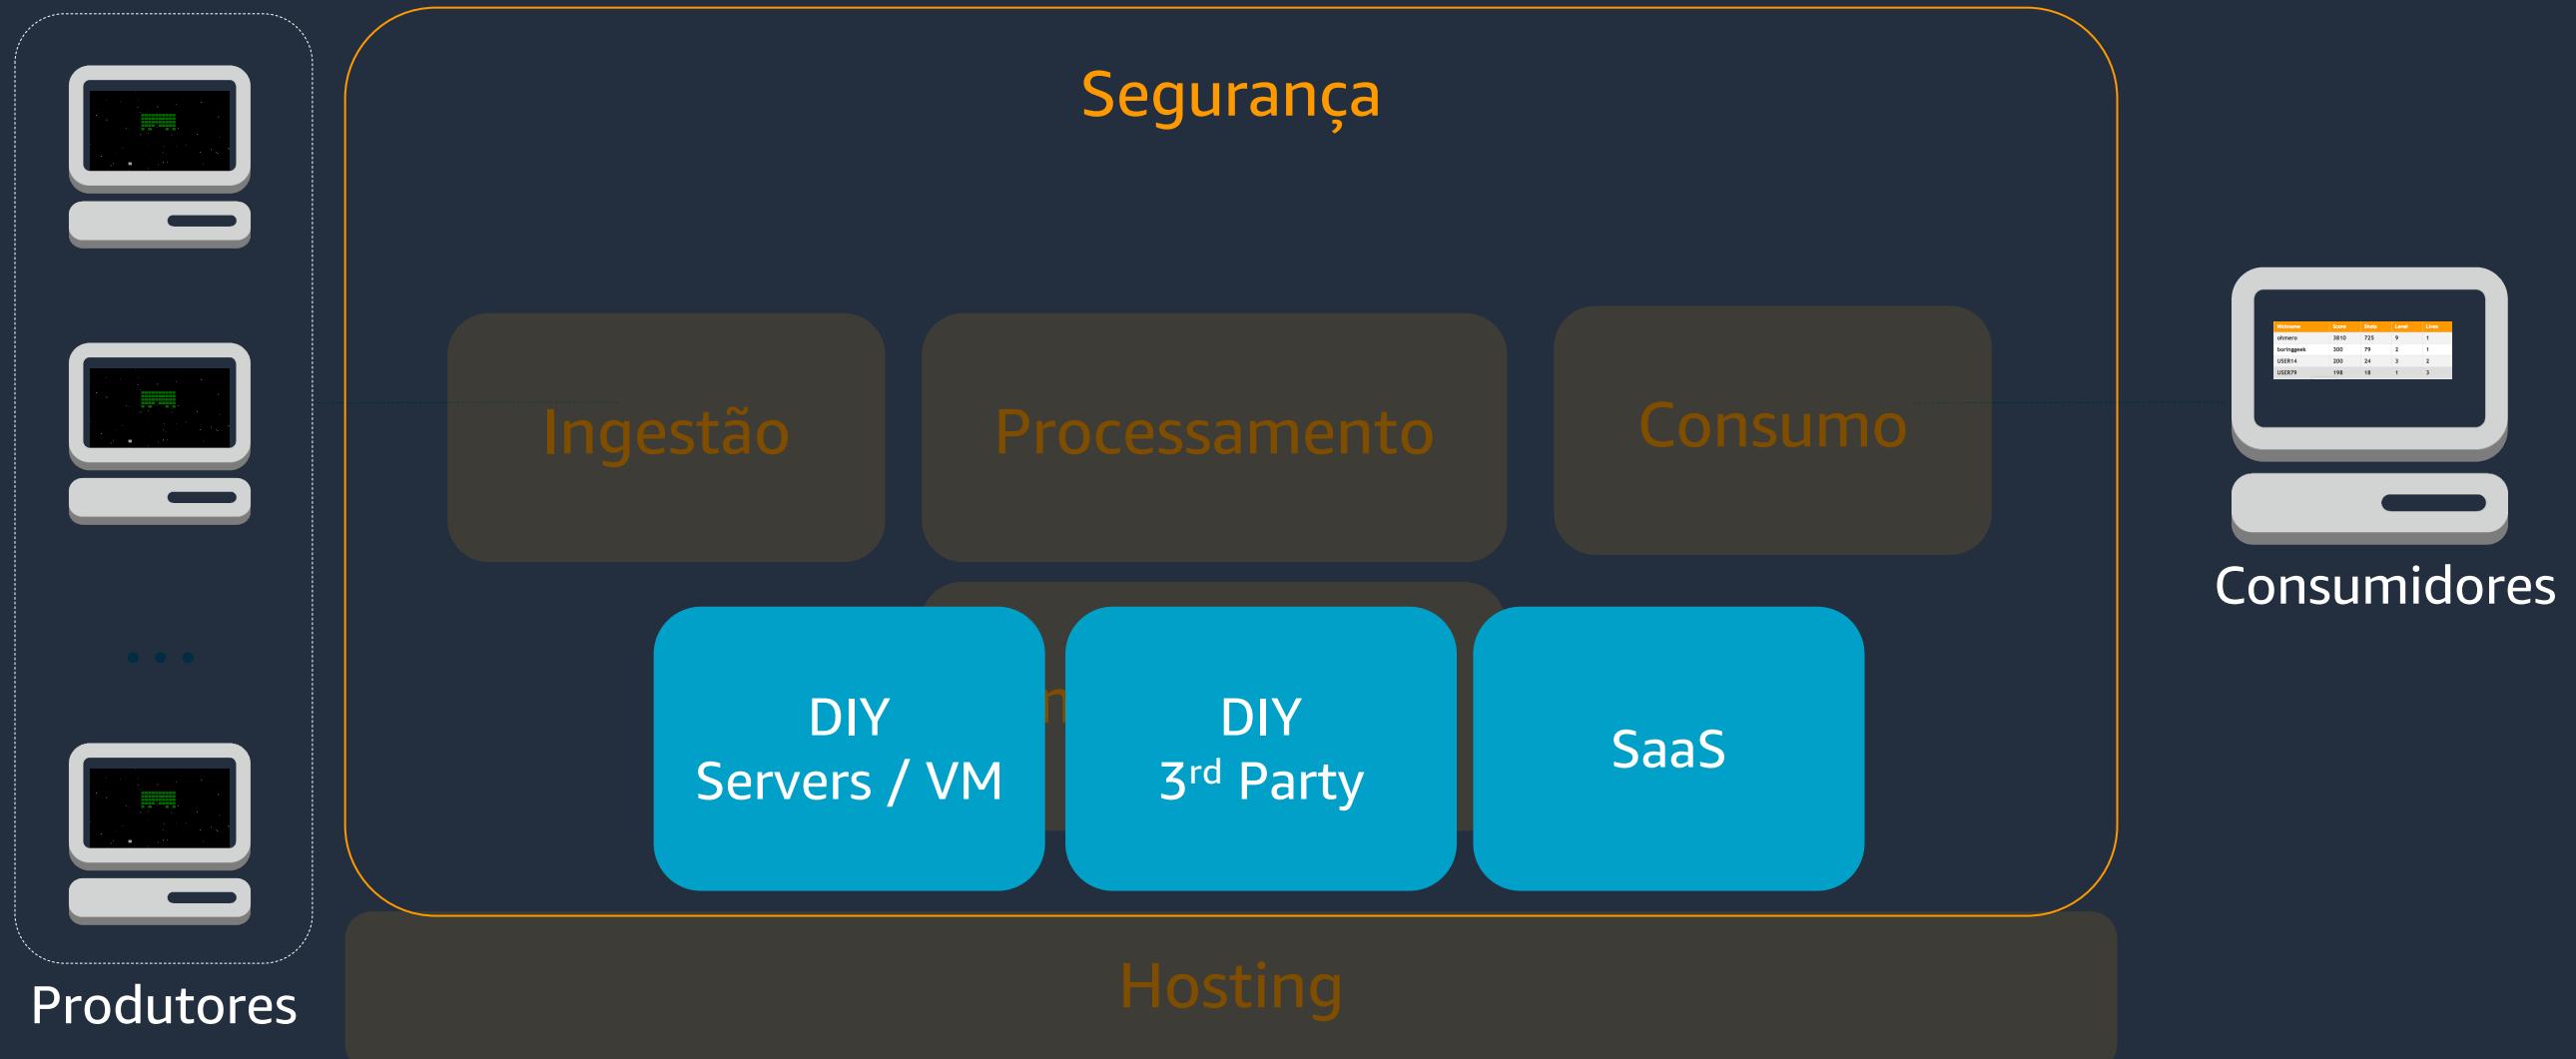
Website hosting



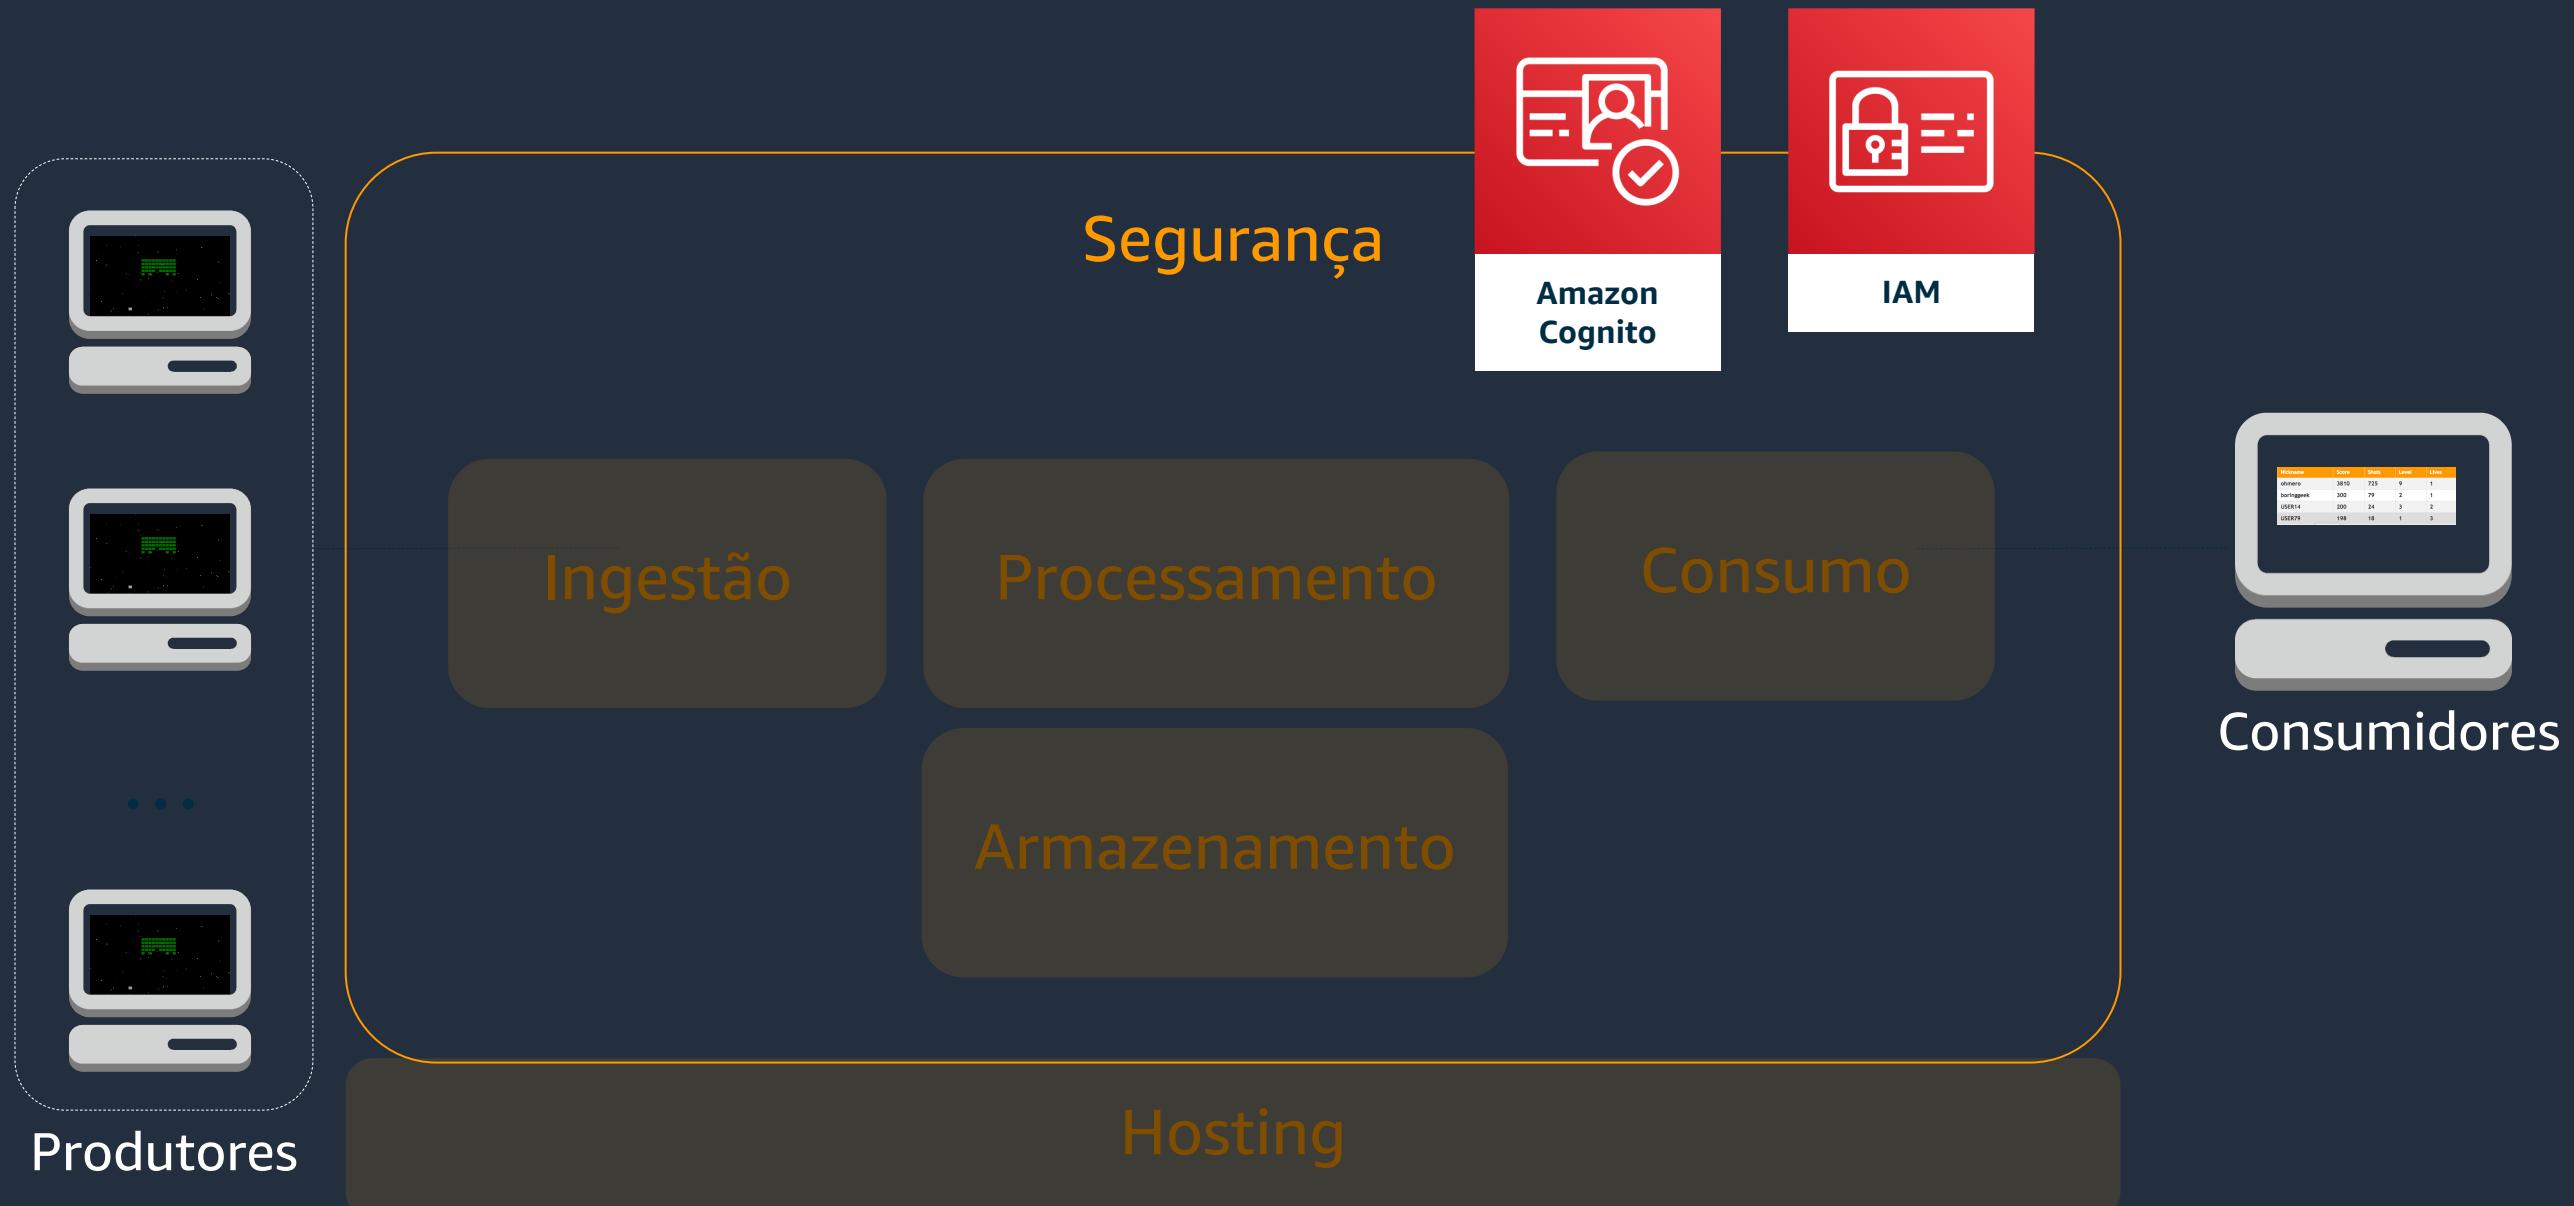
A arquitetura implementada



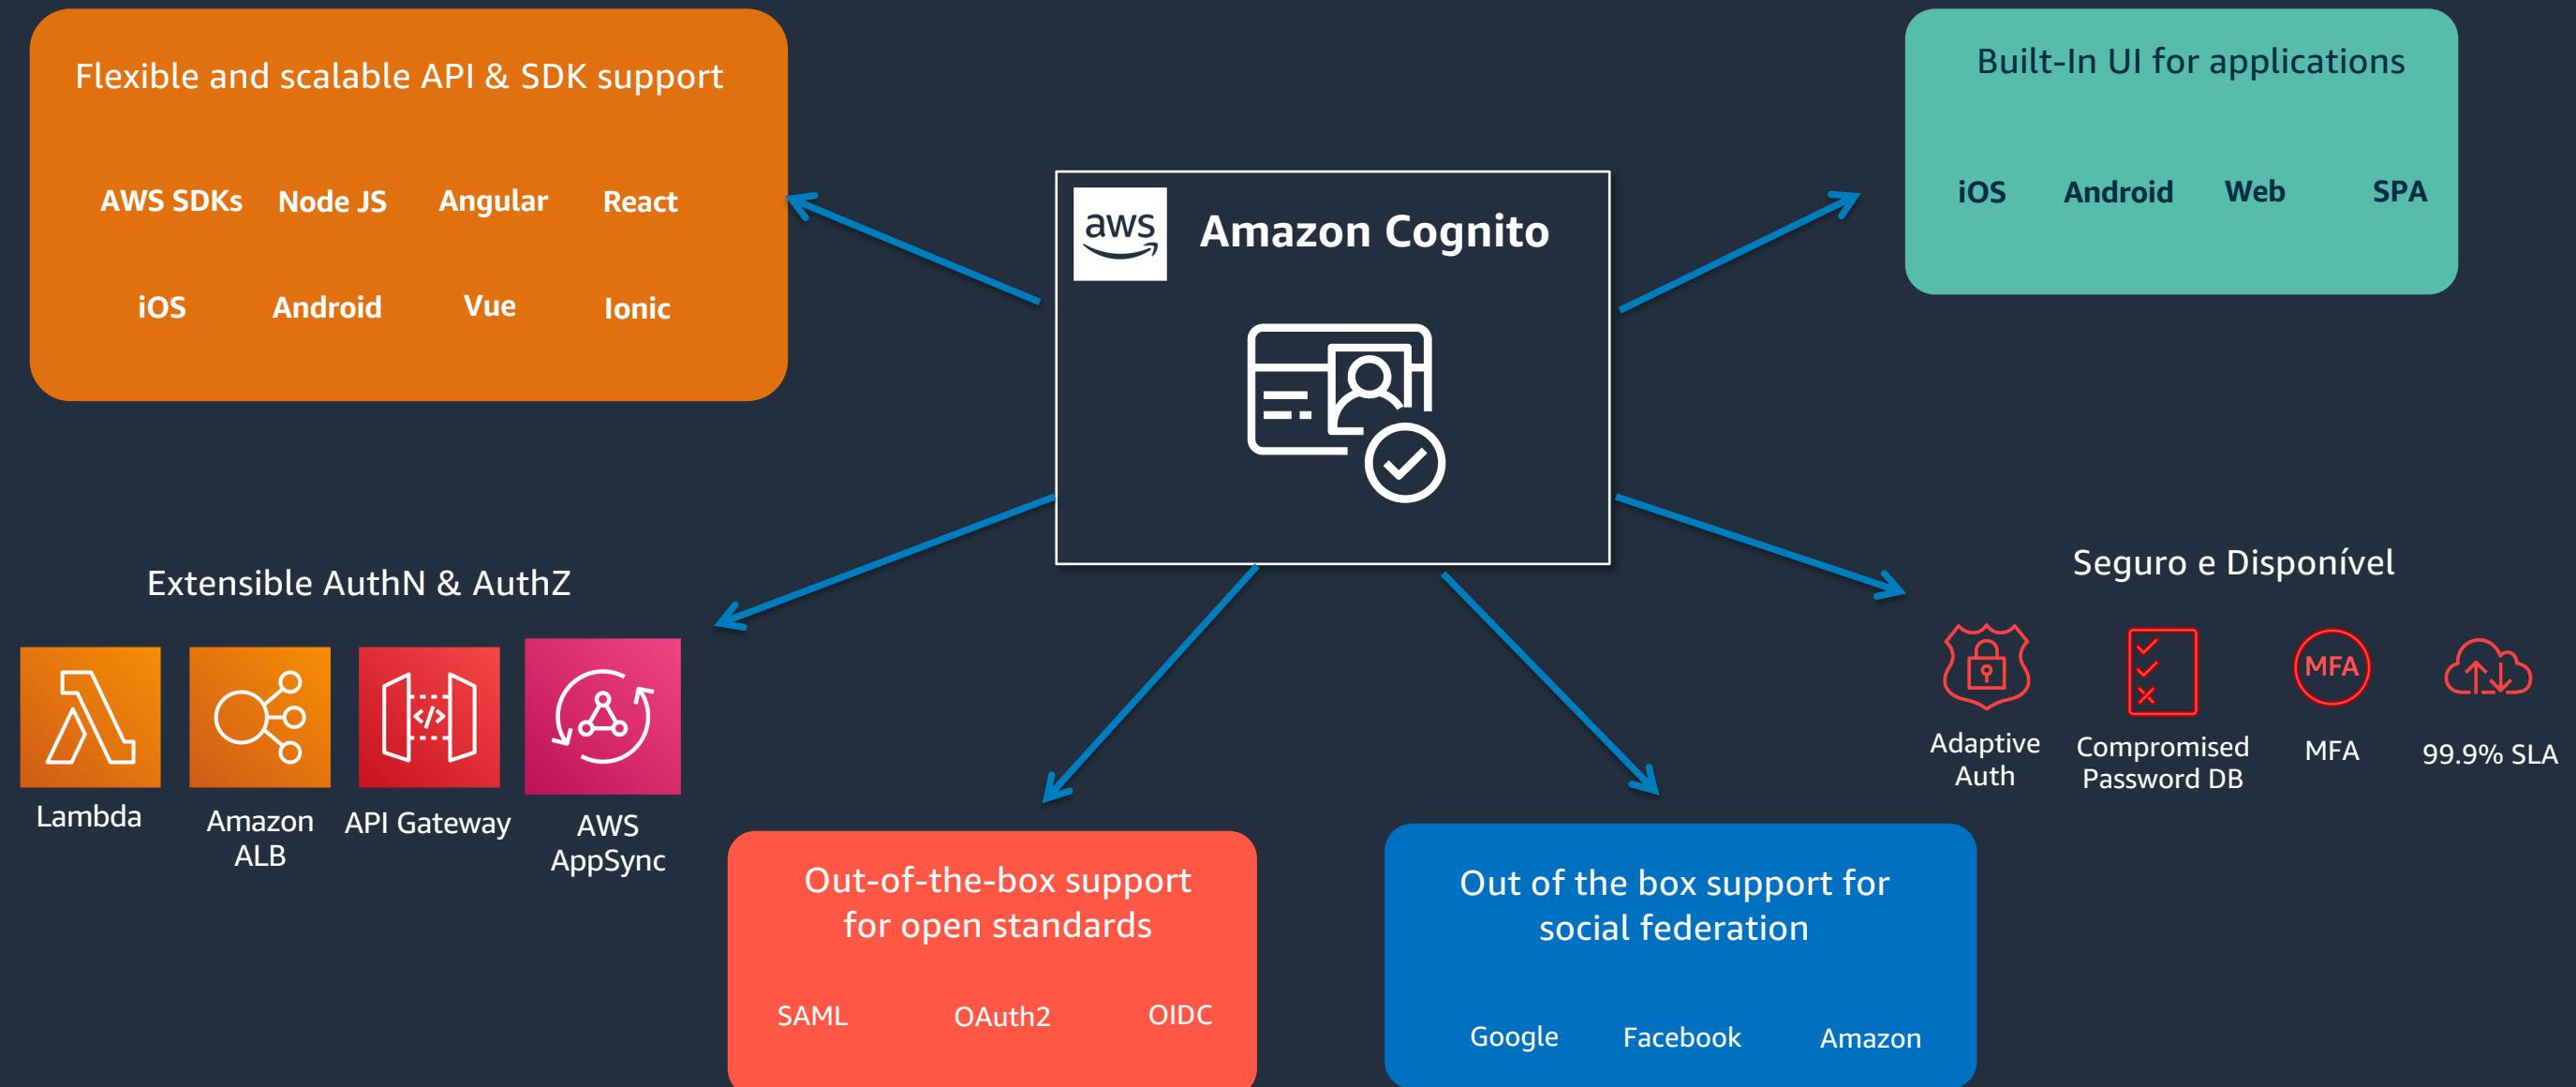
Segurança (RBAC – Role Based Access Control)



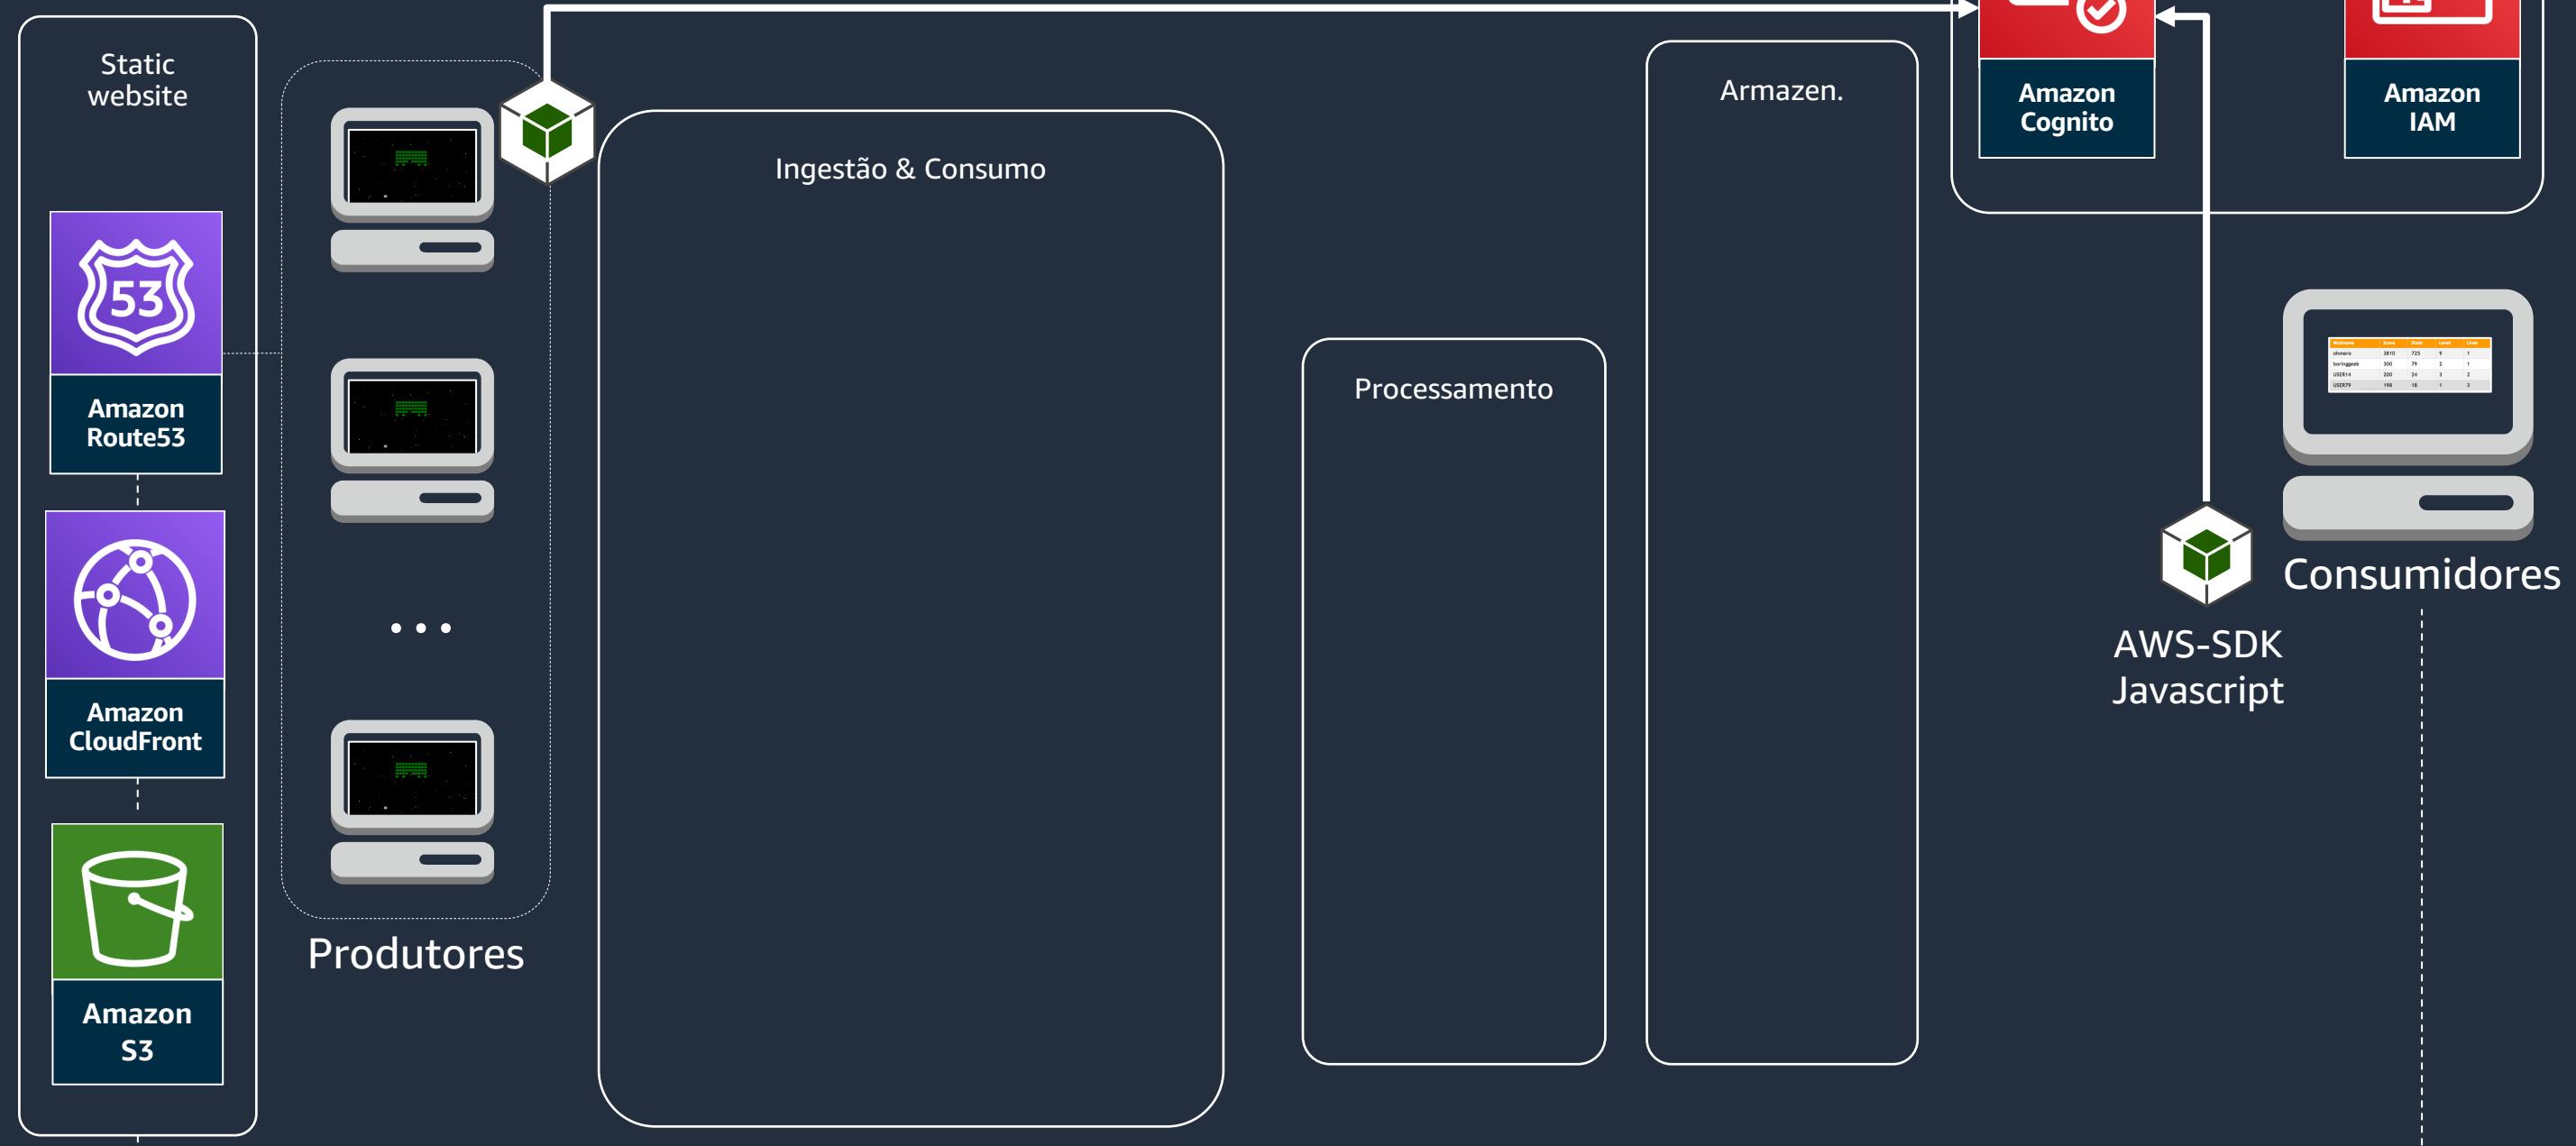
Segurança (RBAC – Role Based Access Control)



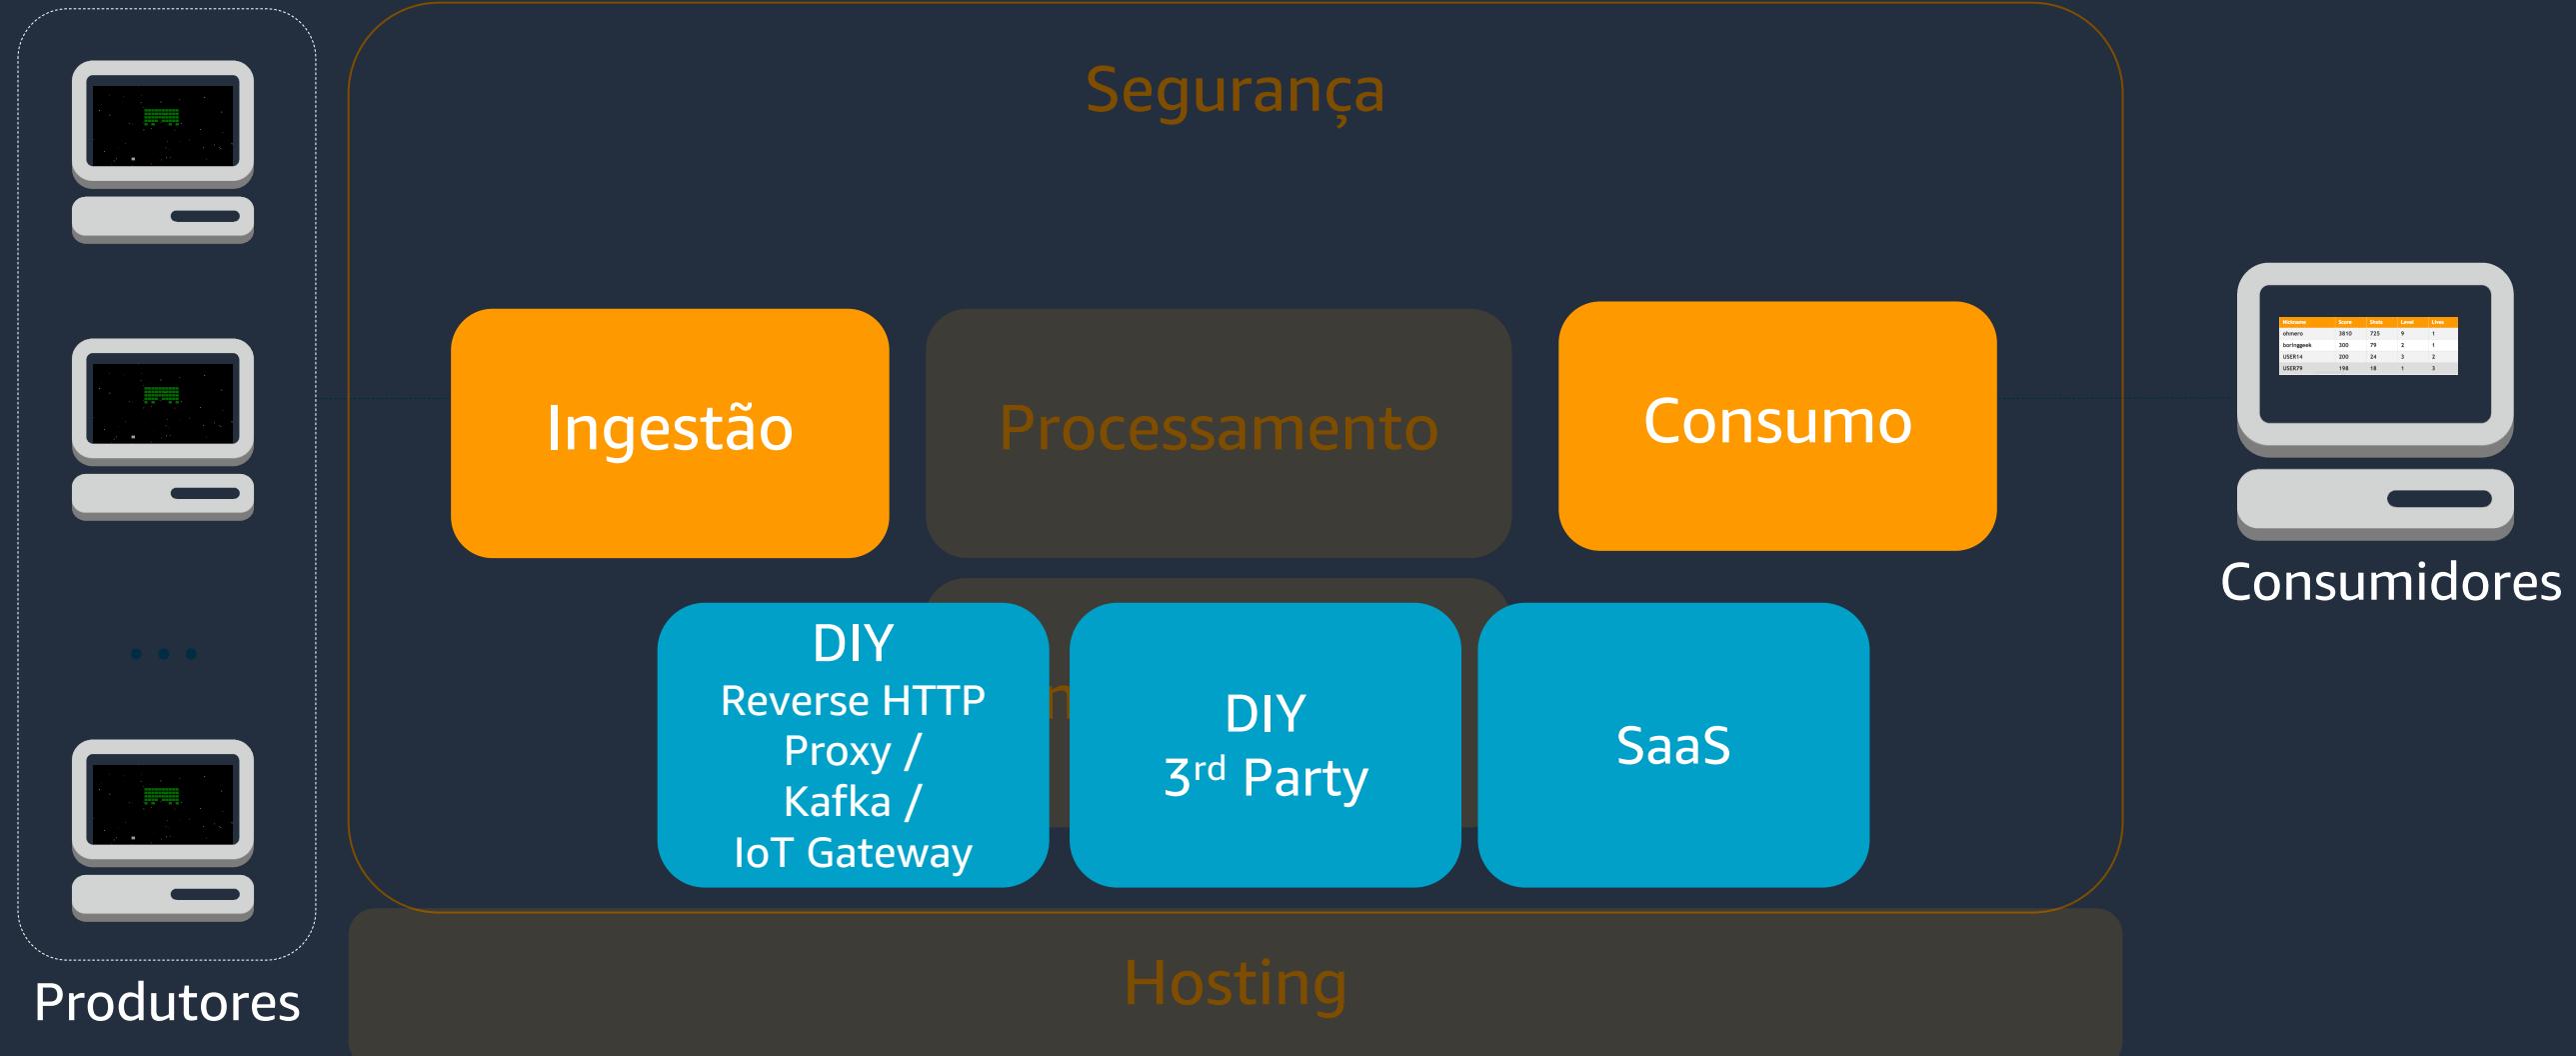
Amazon Cognito



A arquitetura implementada



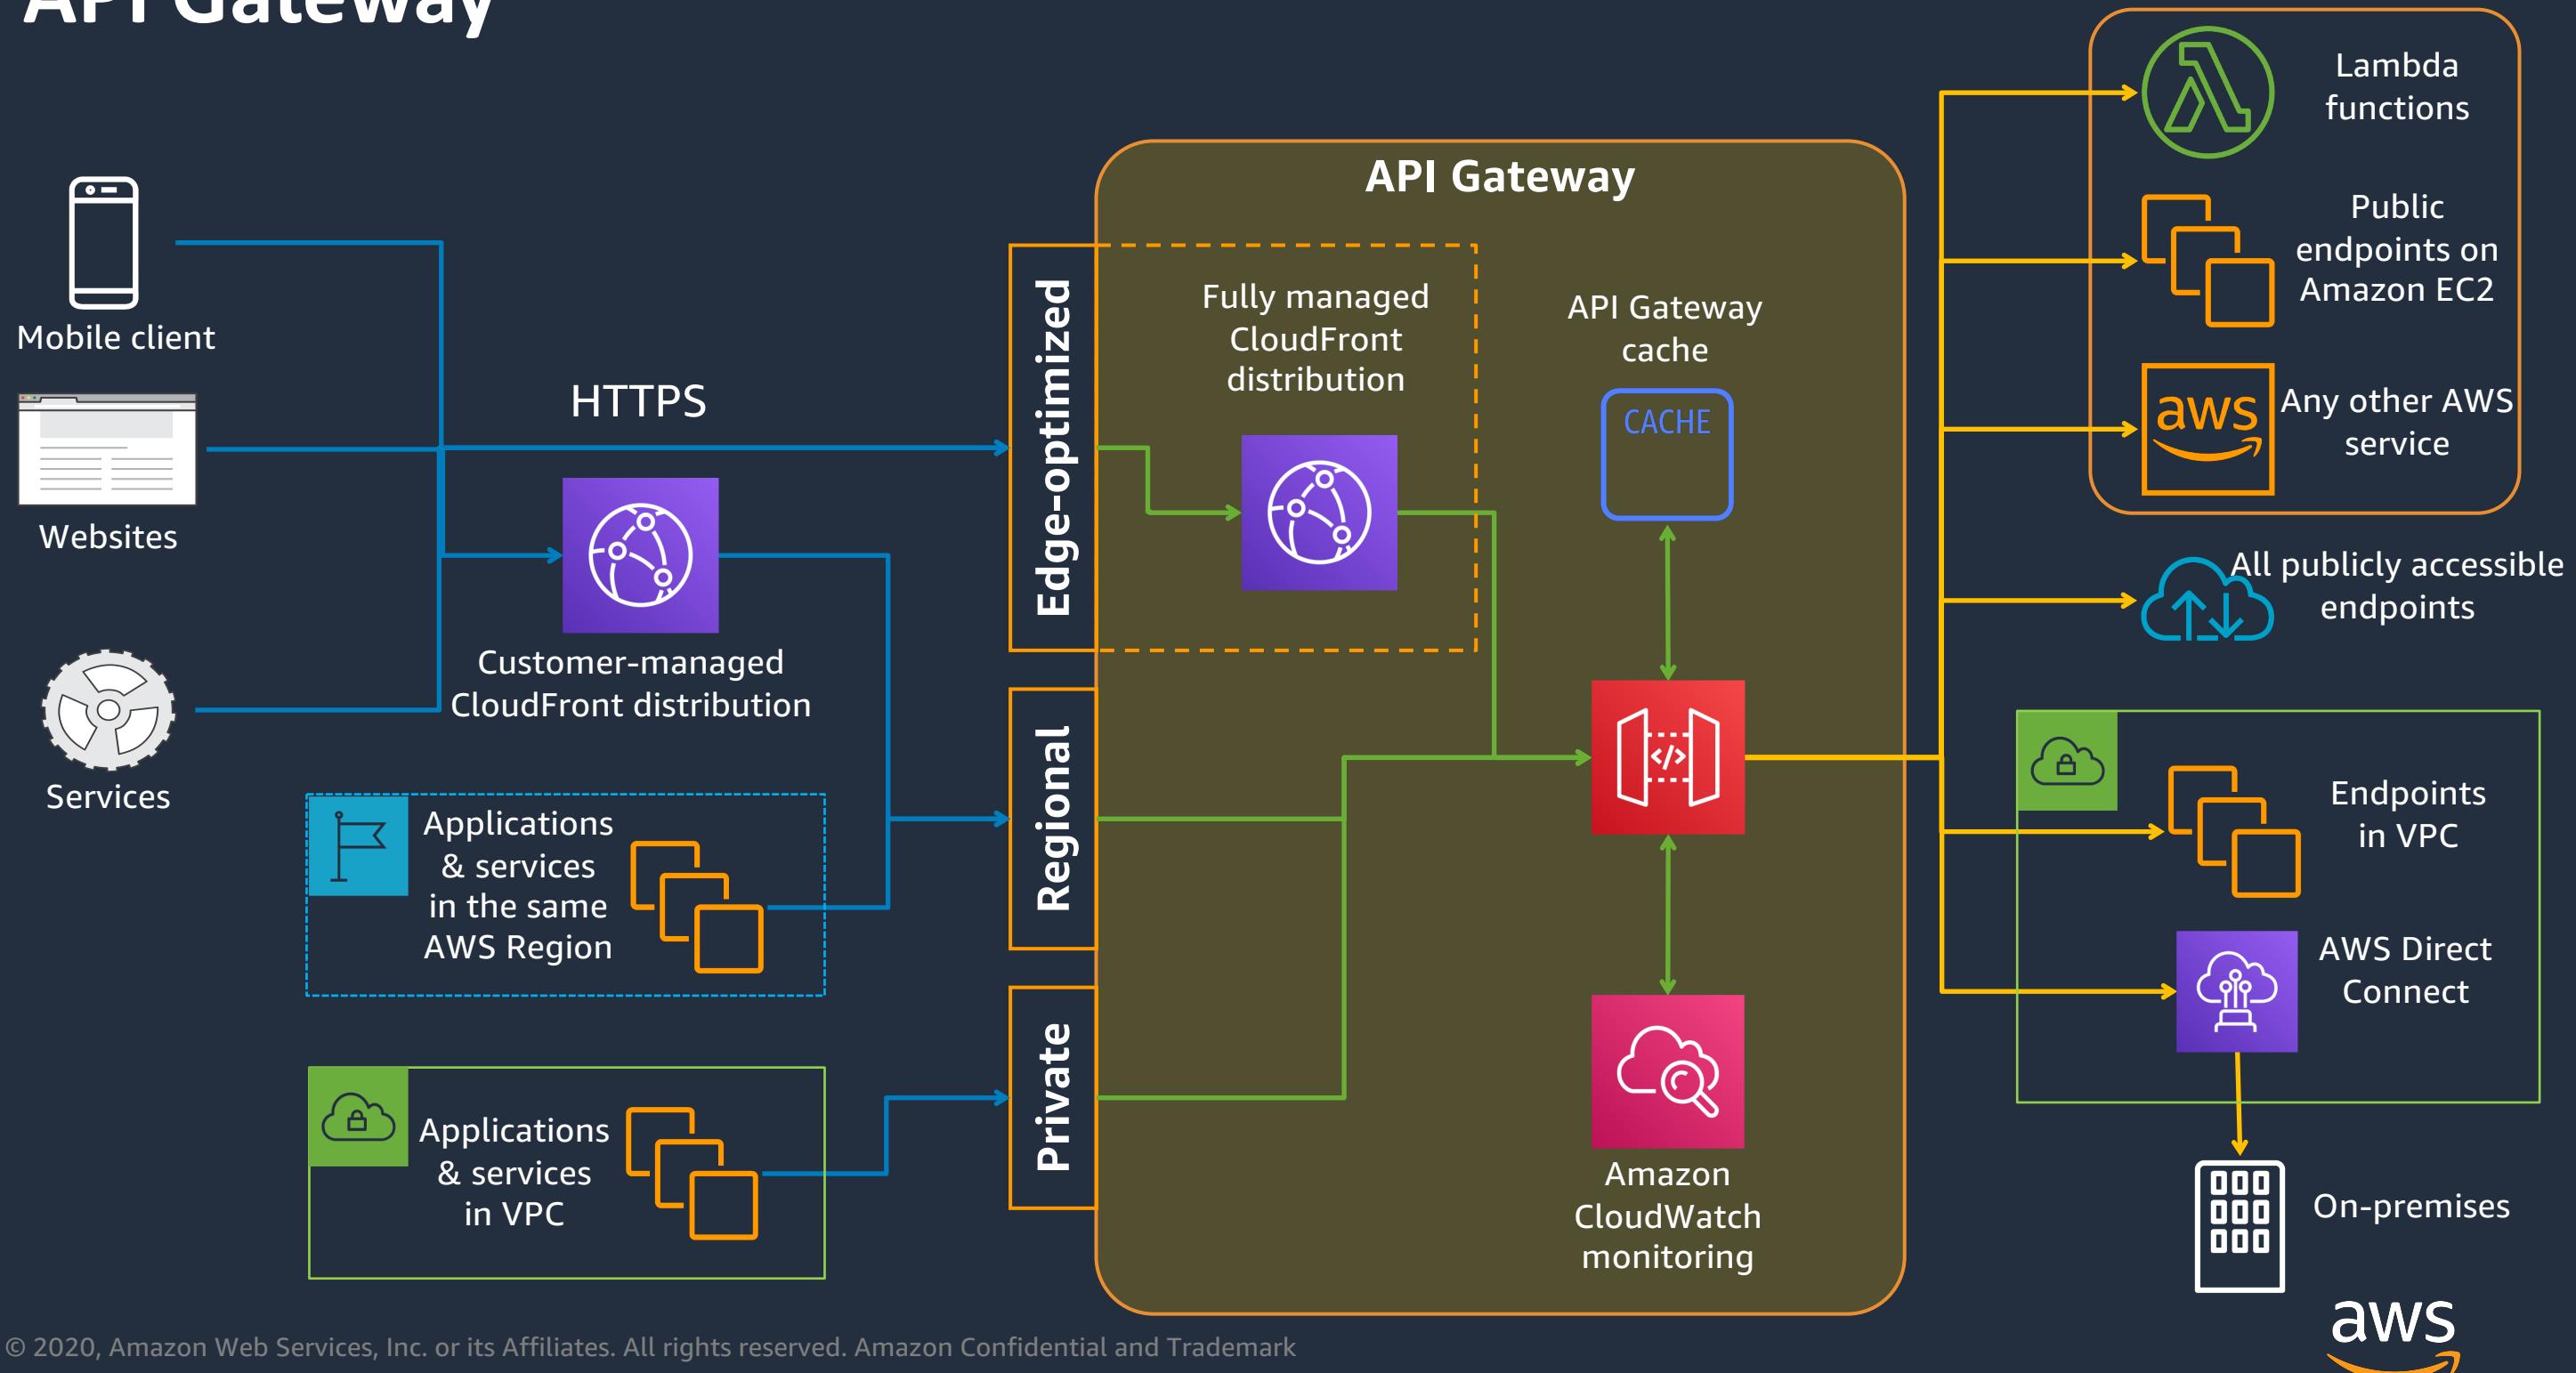
Ingestão e Consumo de Dados



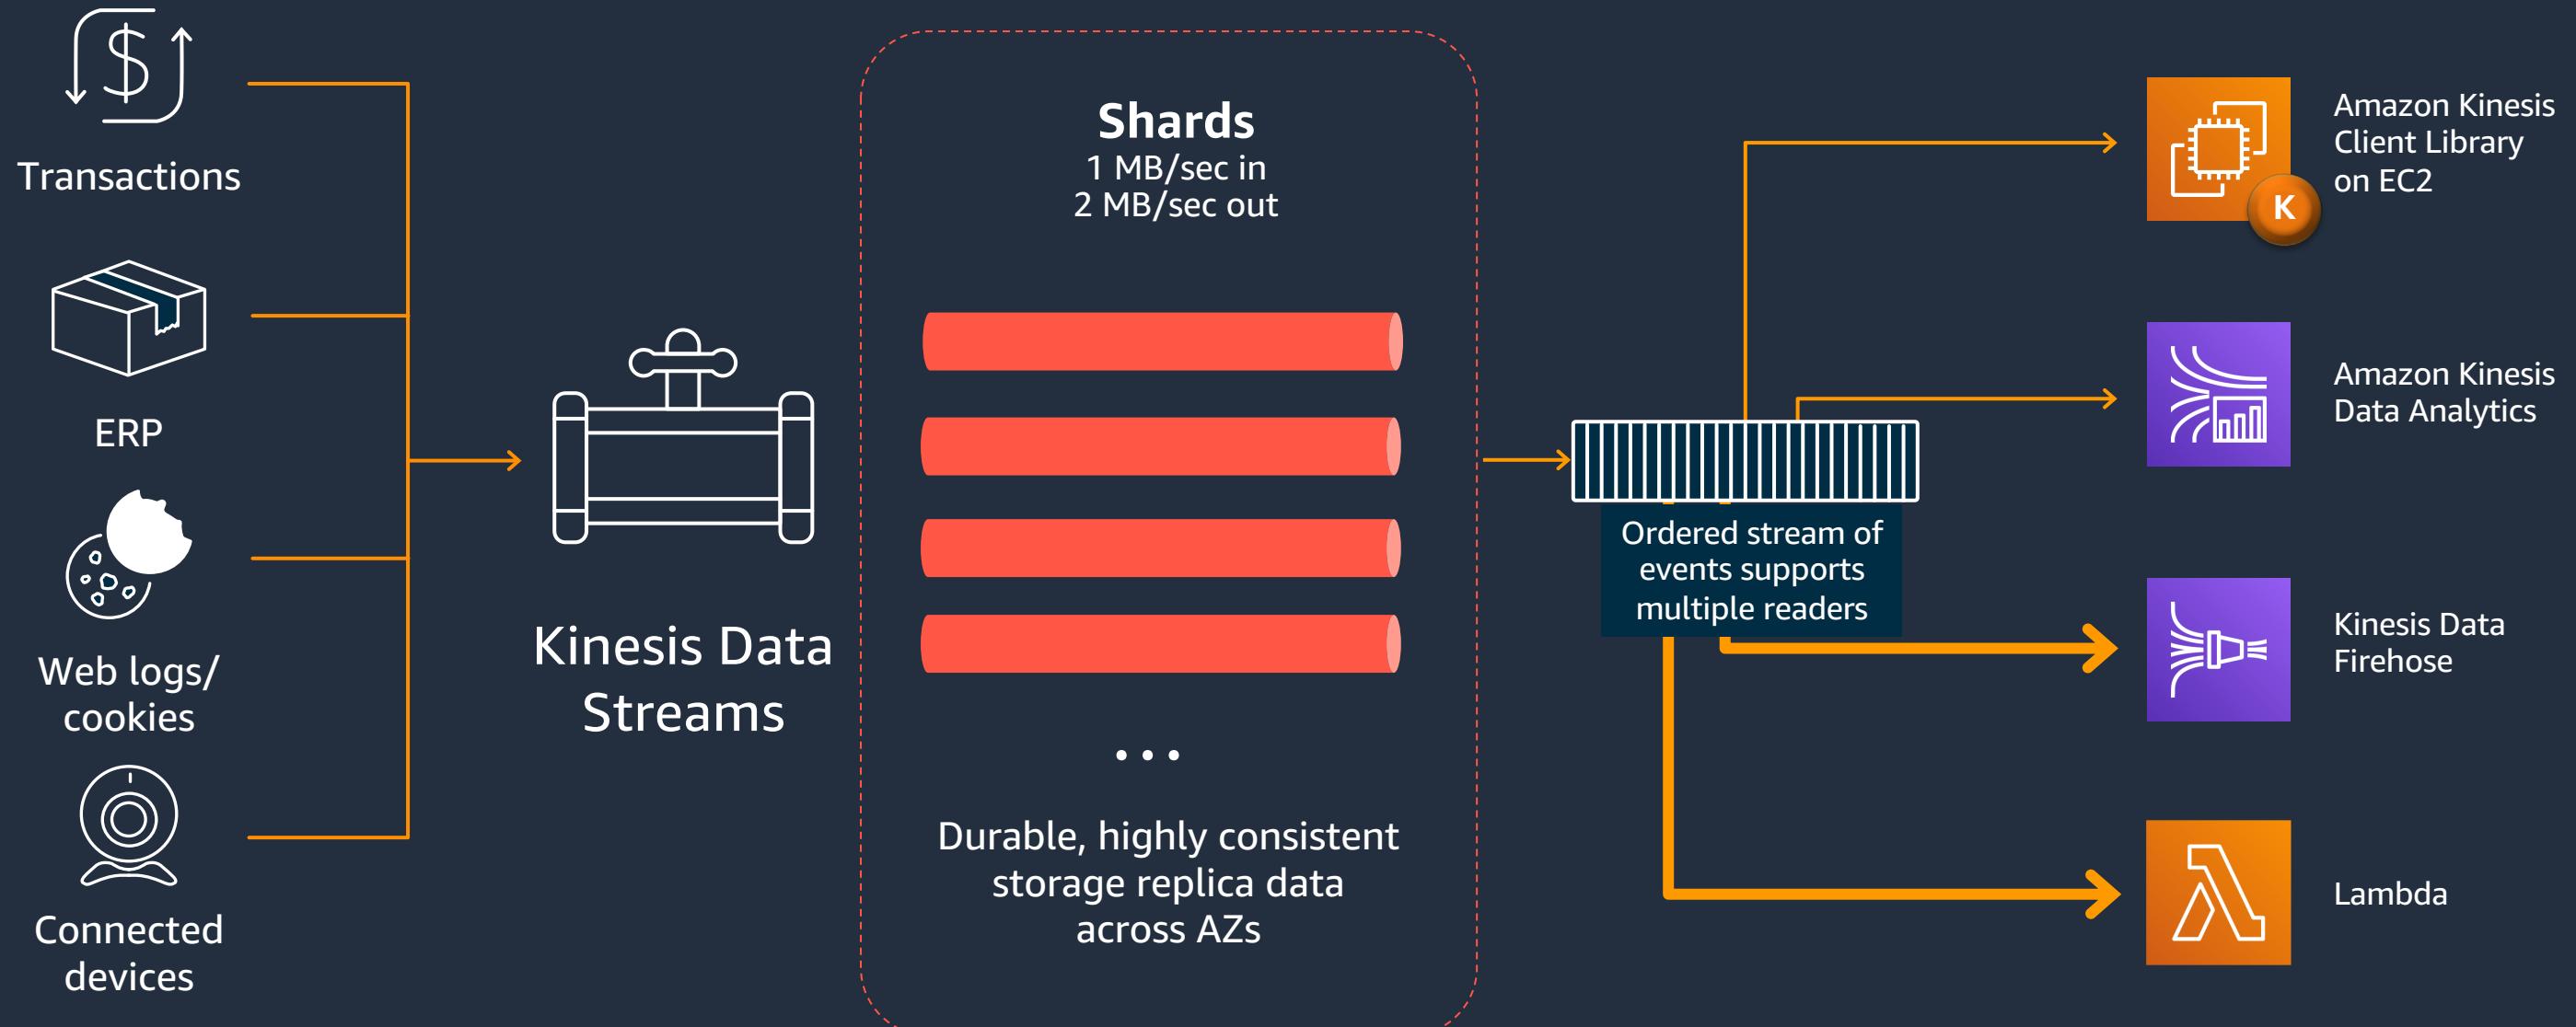
Ingestão e Consumo de Dados



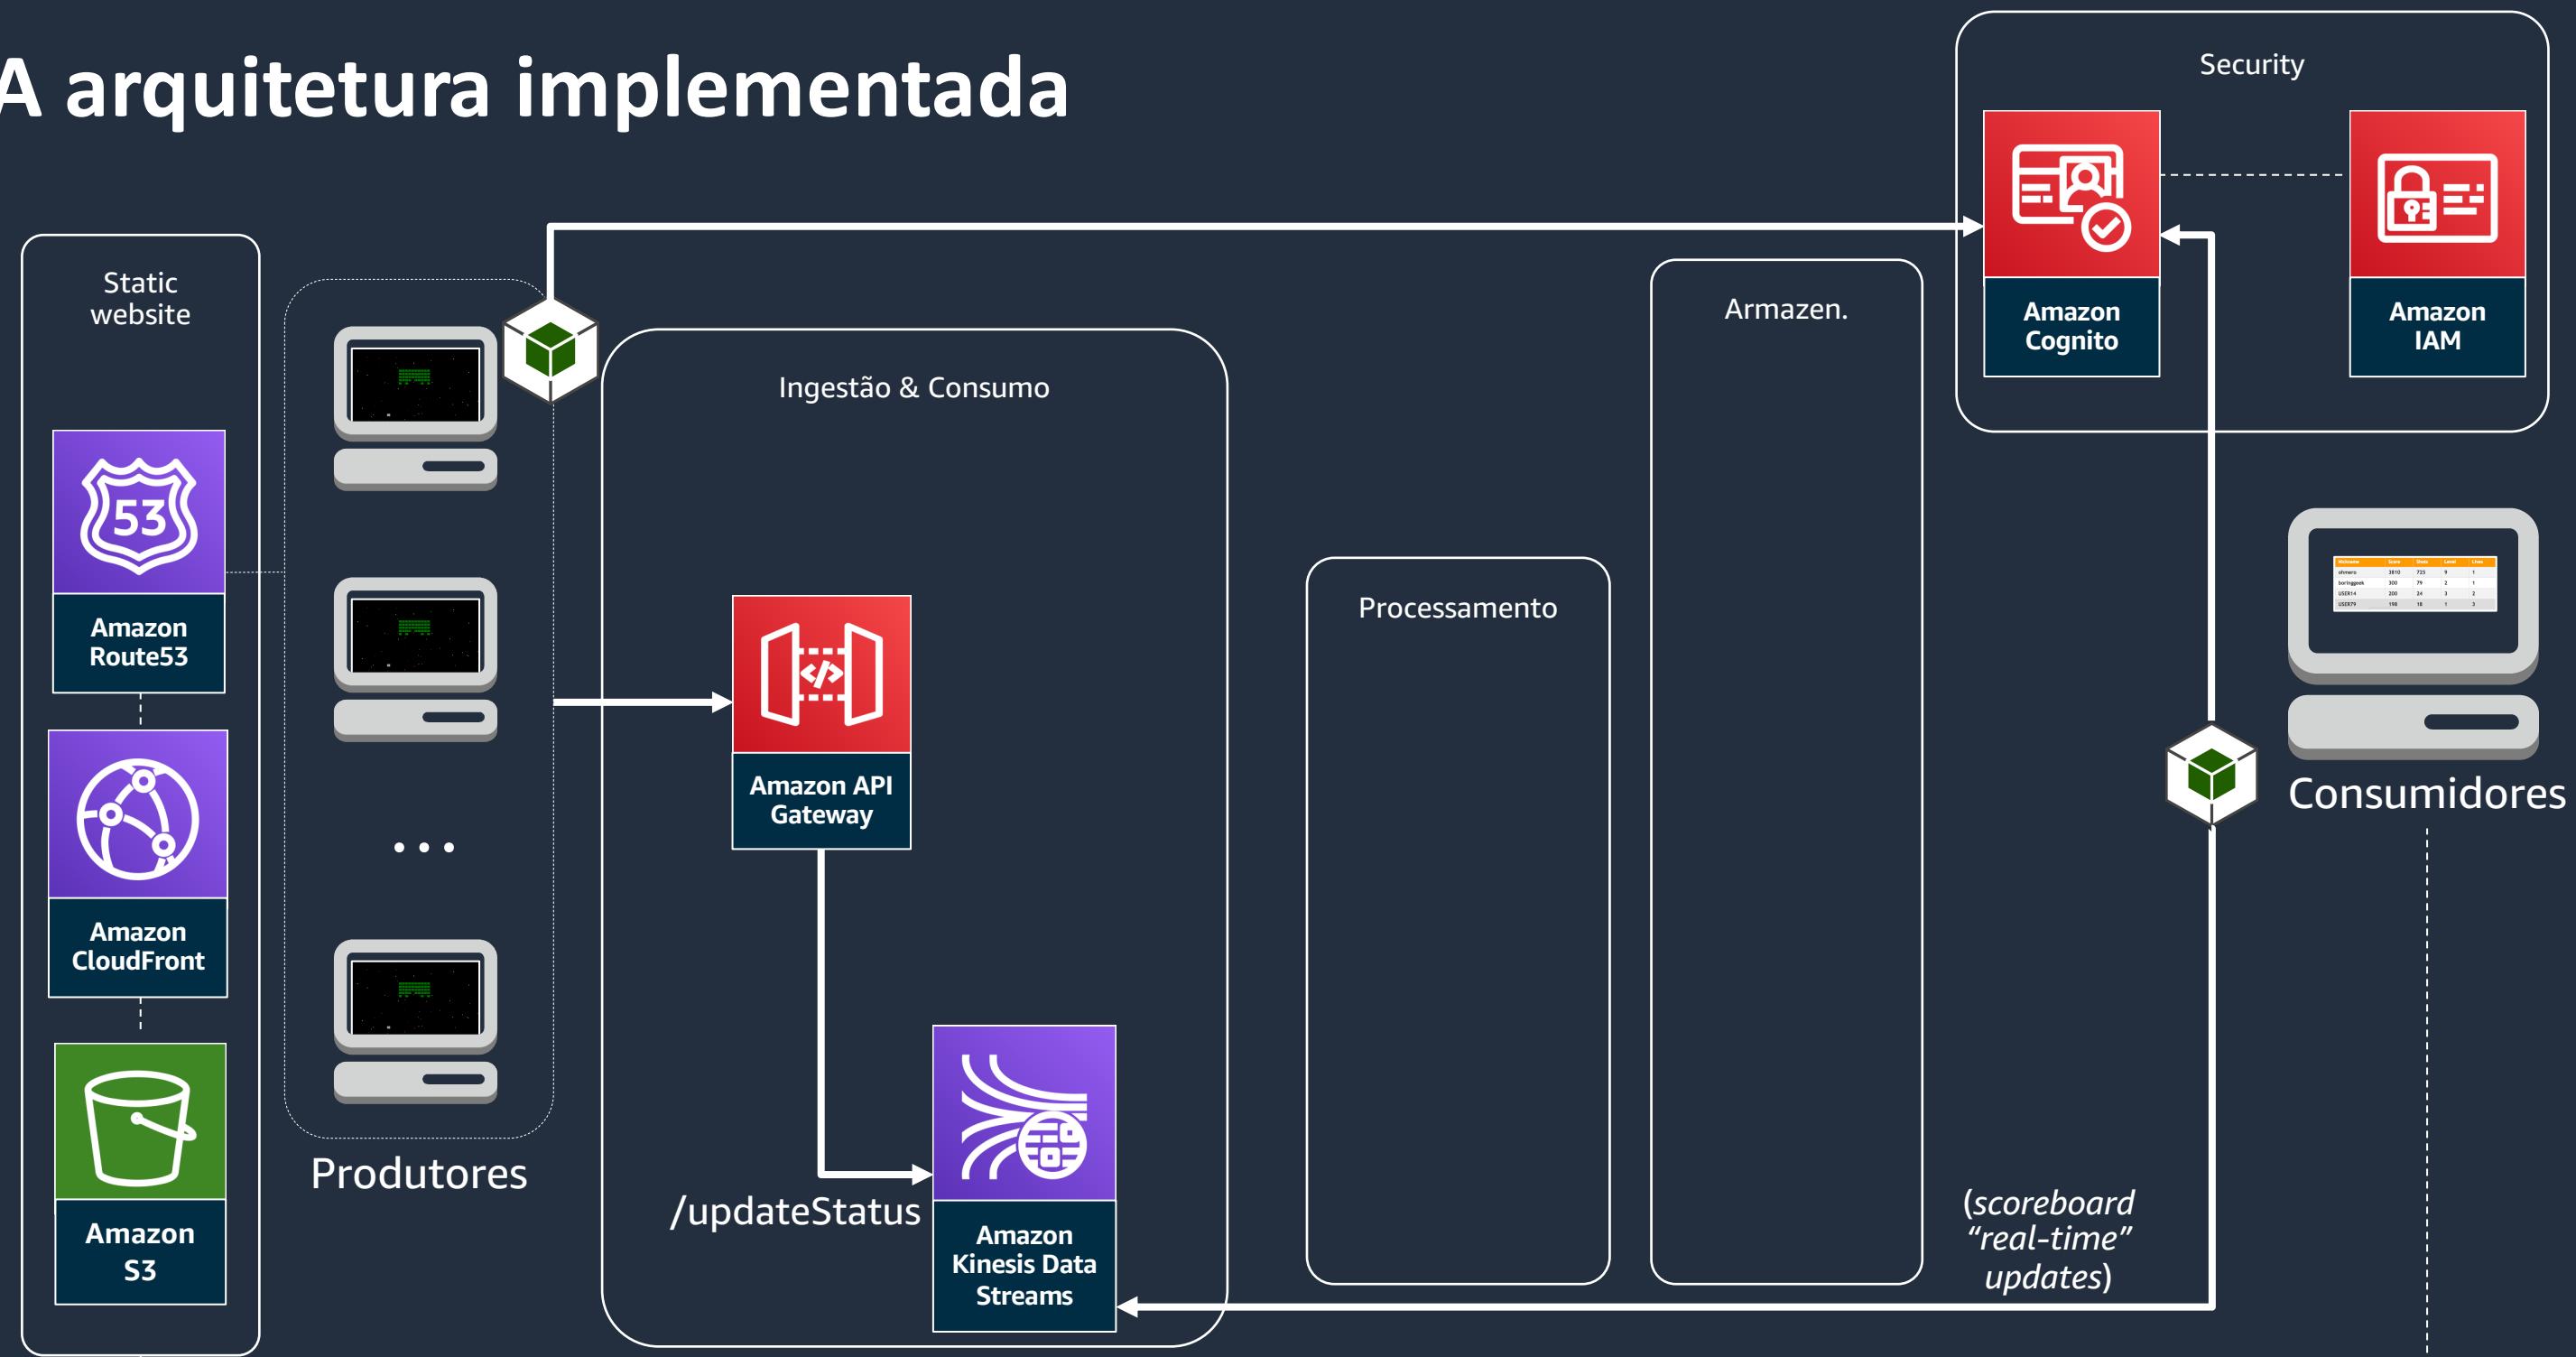
API Gateway



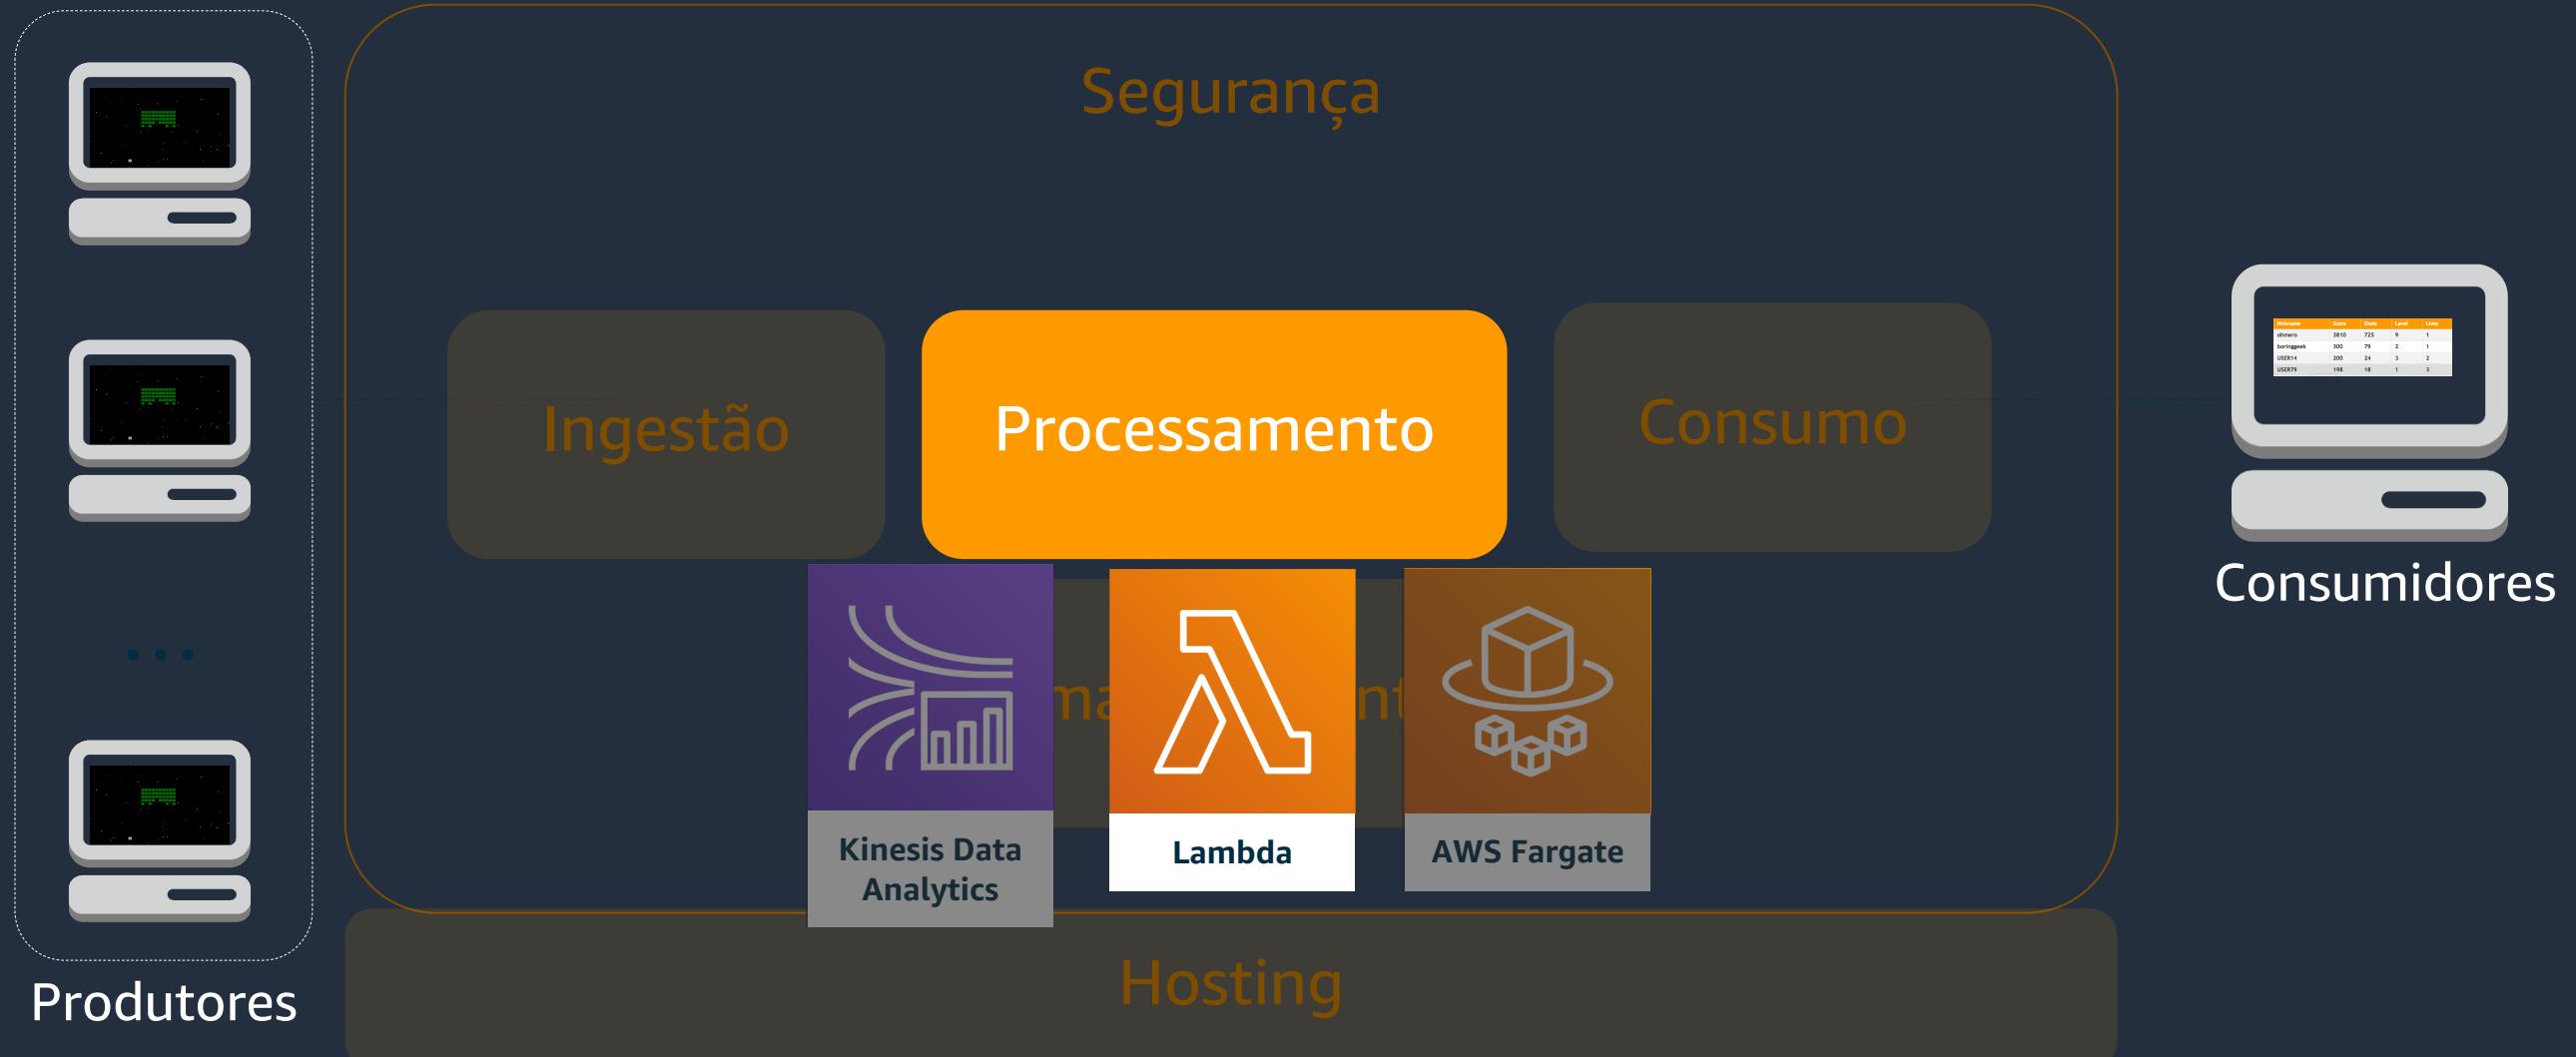
Kinesis Data Streams



A arquitetura implementada



Processamento



AWS Lambda

Event source



Changes in
data state



Requests to
endpoints



Changes in
resource state



Function



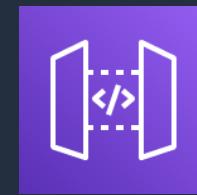
Node.js
Python
Java
.NET Core
Go
Ruby
BYOL

Services (anything)

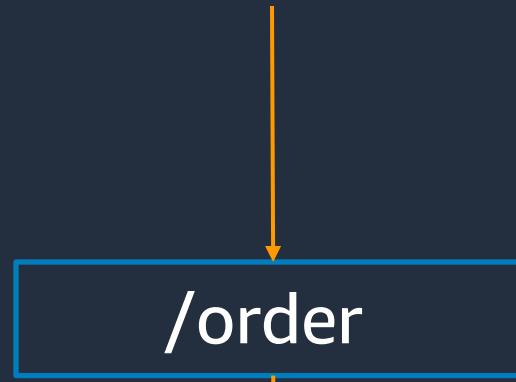


AWS Lambda Execution Model

Synchronous (push)



API Gateway

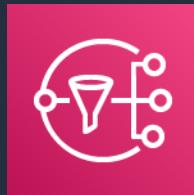


/order

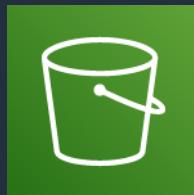


Lambda
function

Asynchronous (event)



Amazon
SNS



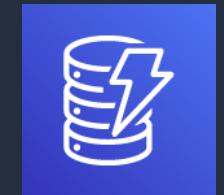
Amazon
S3

reqs



Lambda
function

Poll-based



DynamoDB



Amazon
Kinesis

changes

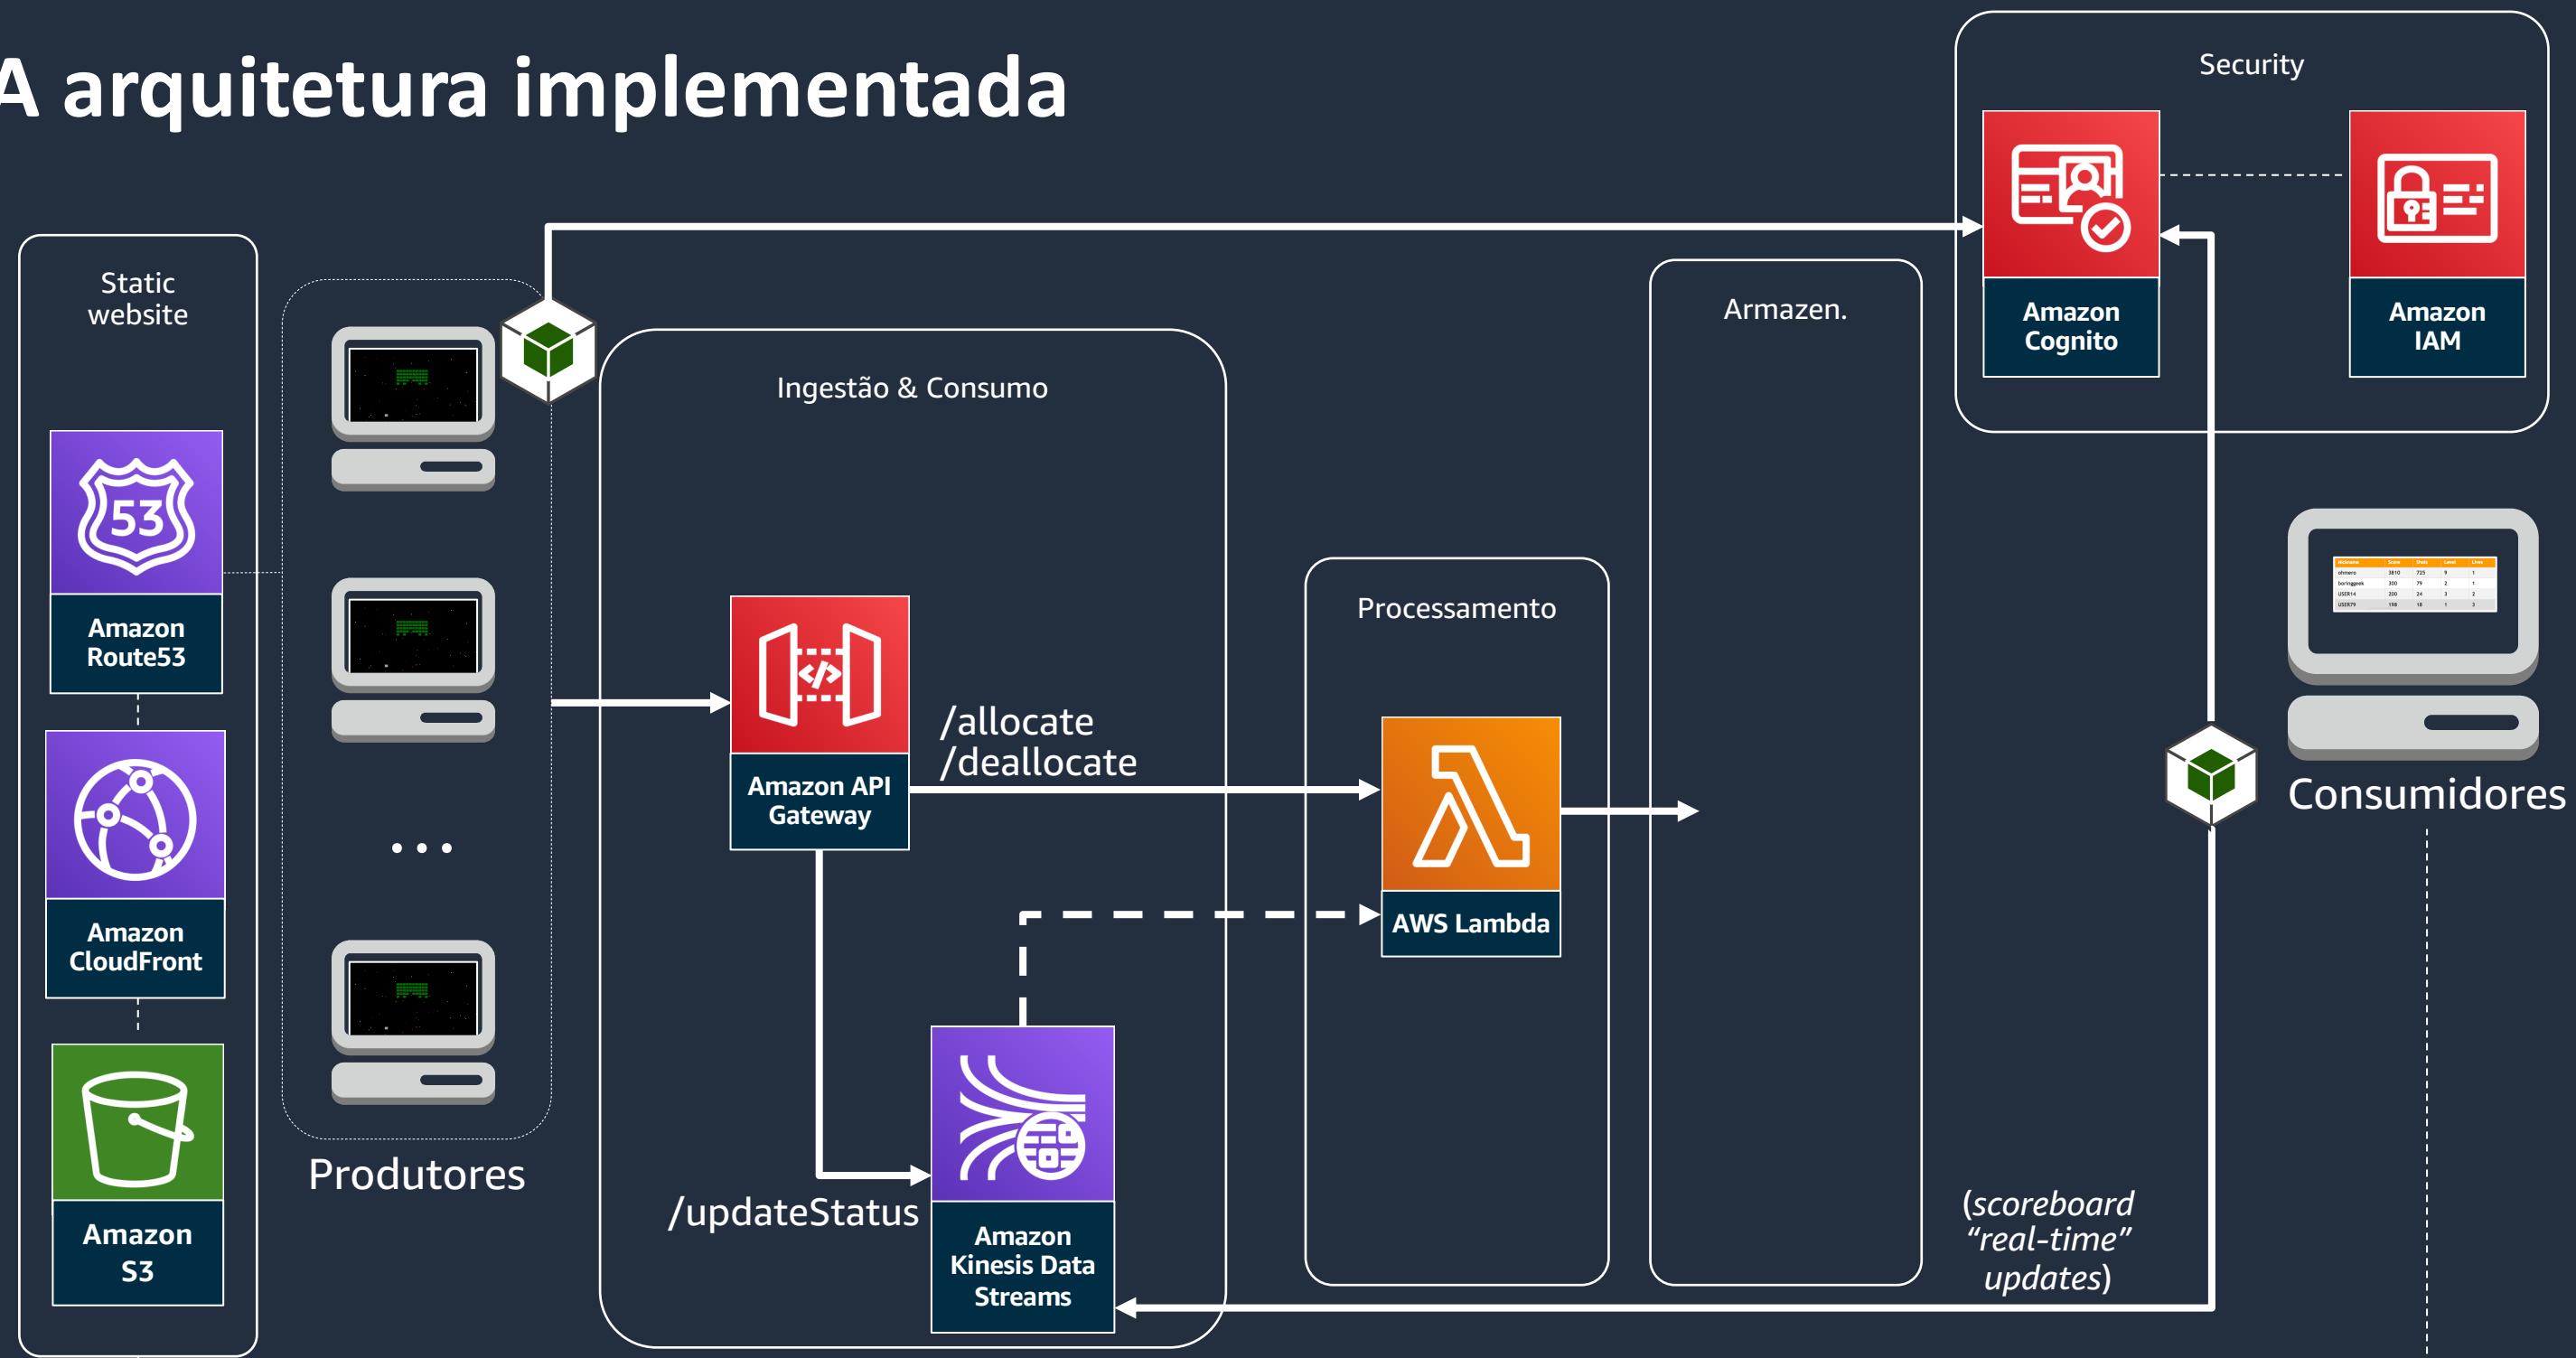


Lambda service

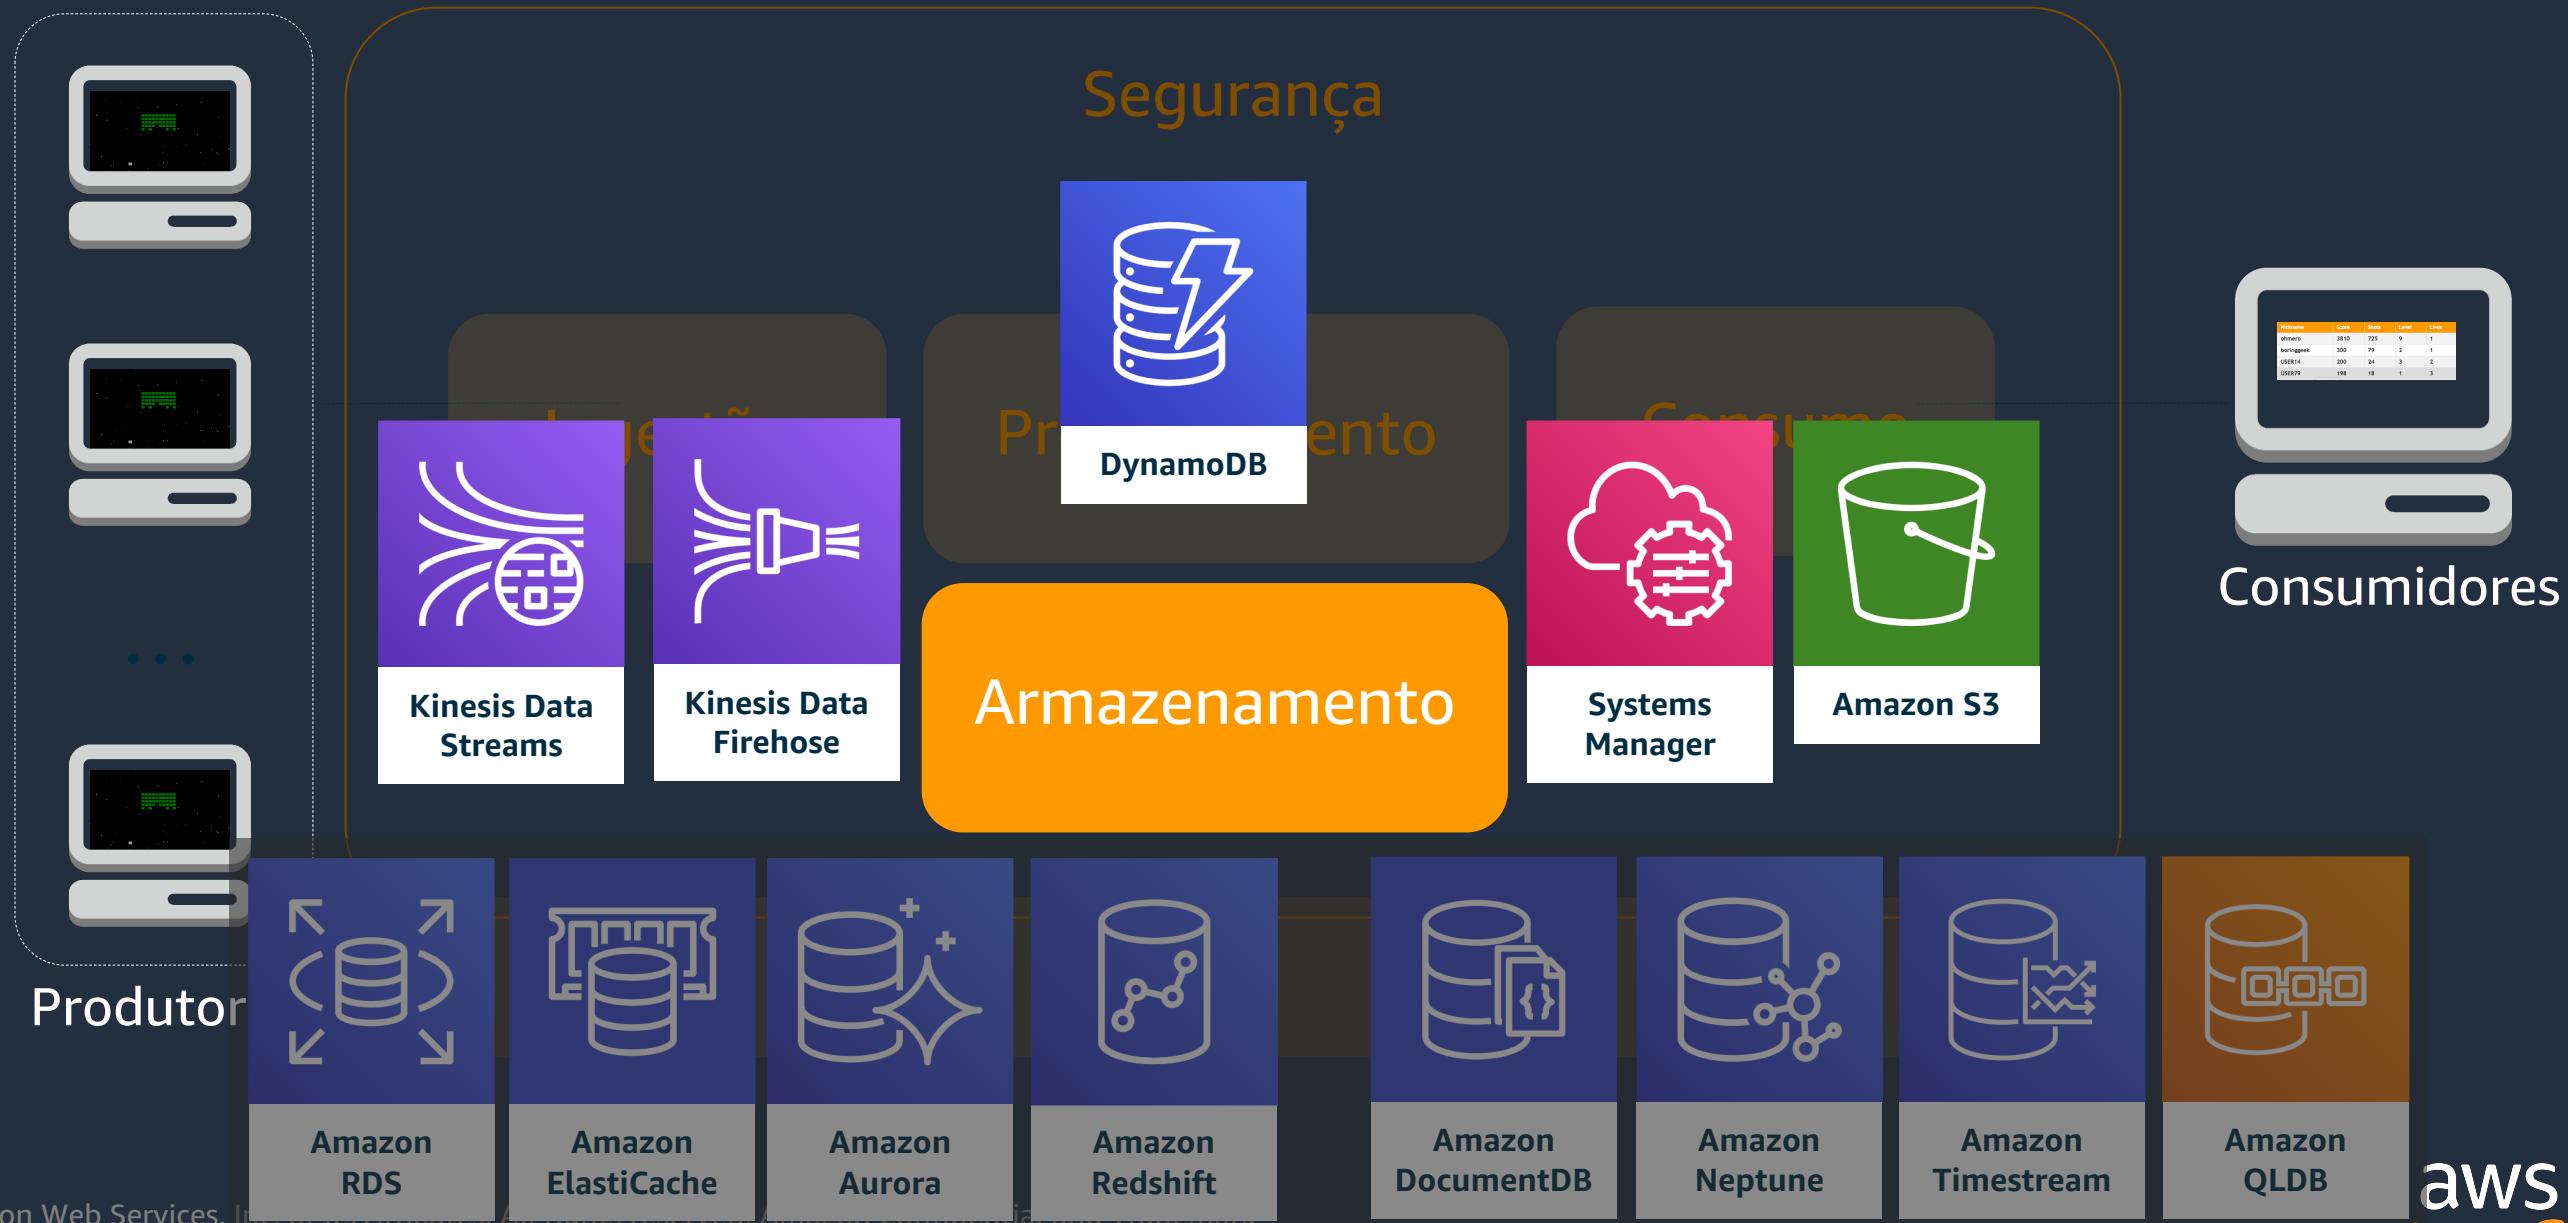


Lambda
function
aws

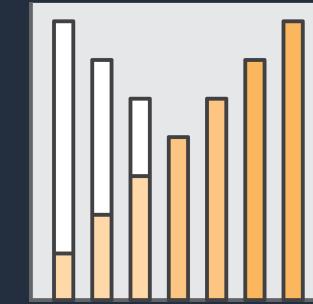
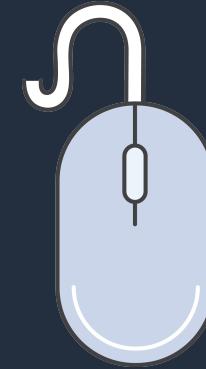
A arquitetura implementada



Armazenamento



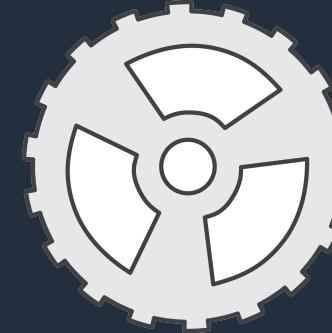
Amazon DynamoDB



Fully managed NoSQL

Document or key-value

Scales to any workload

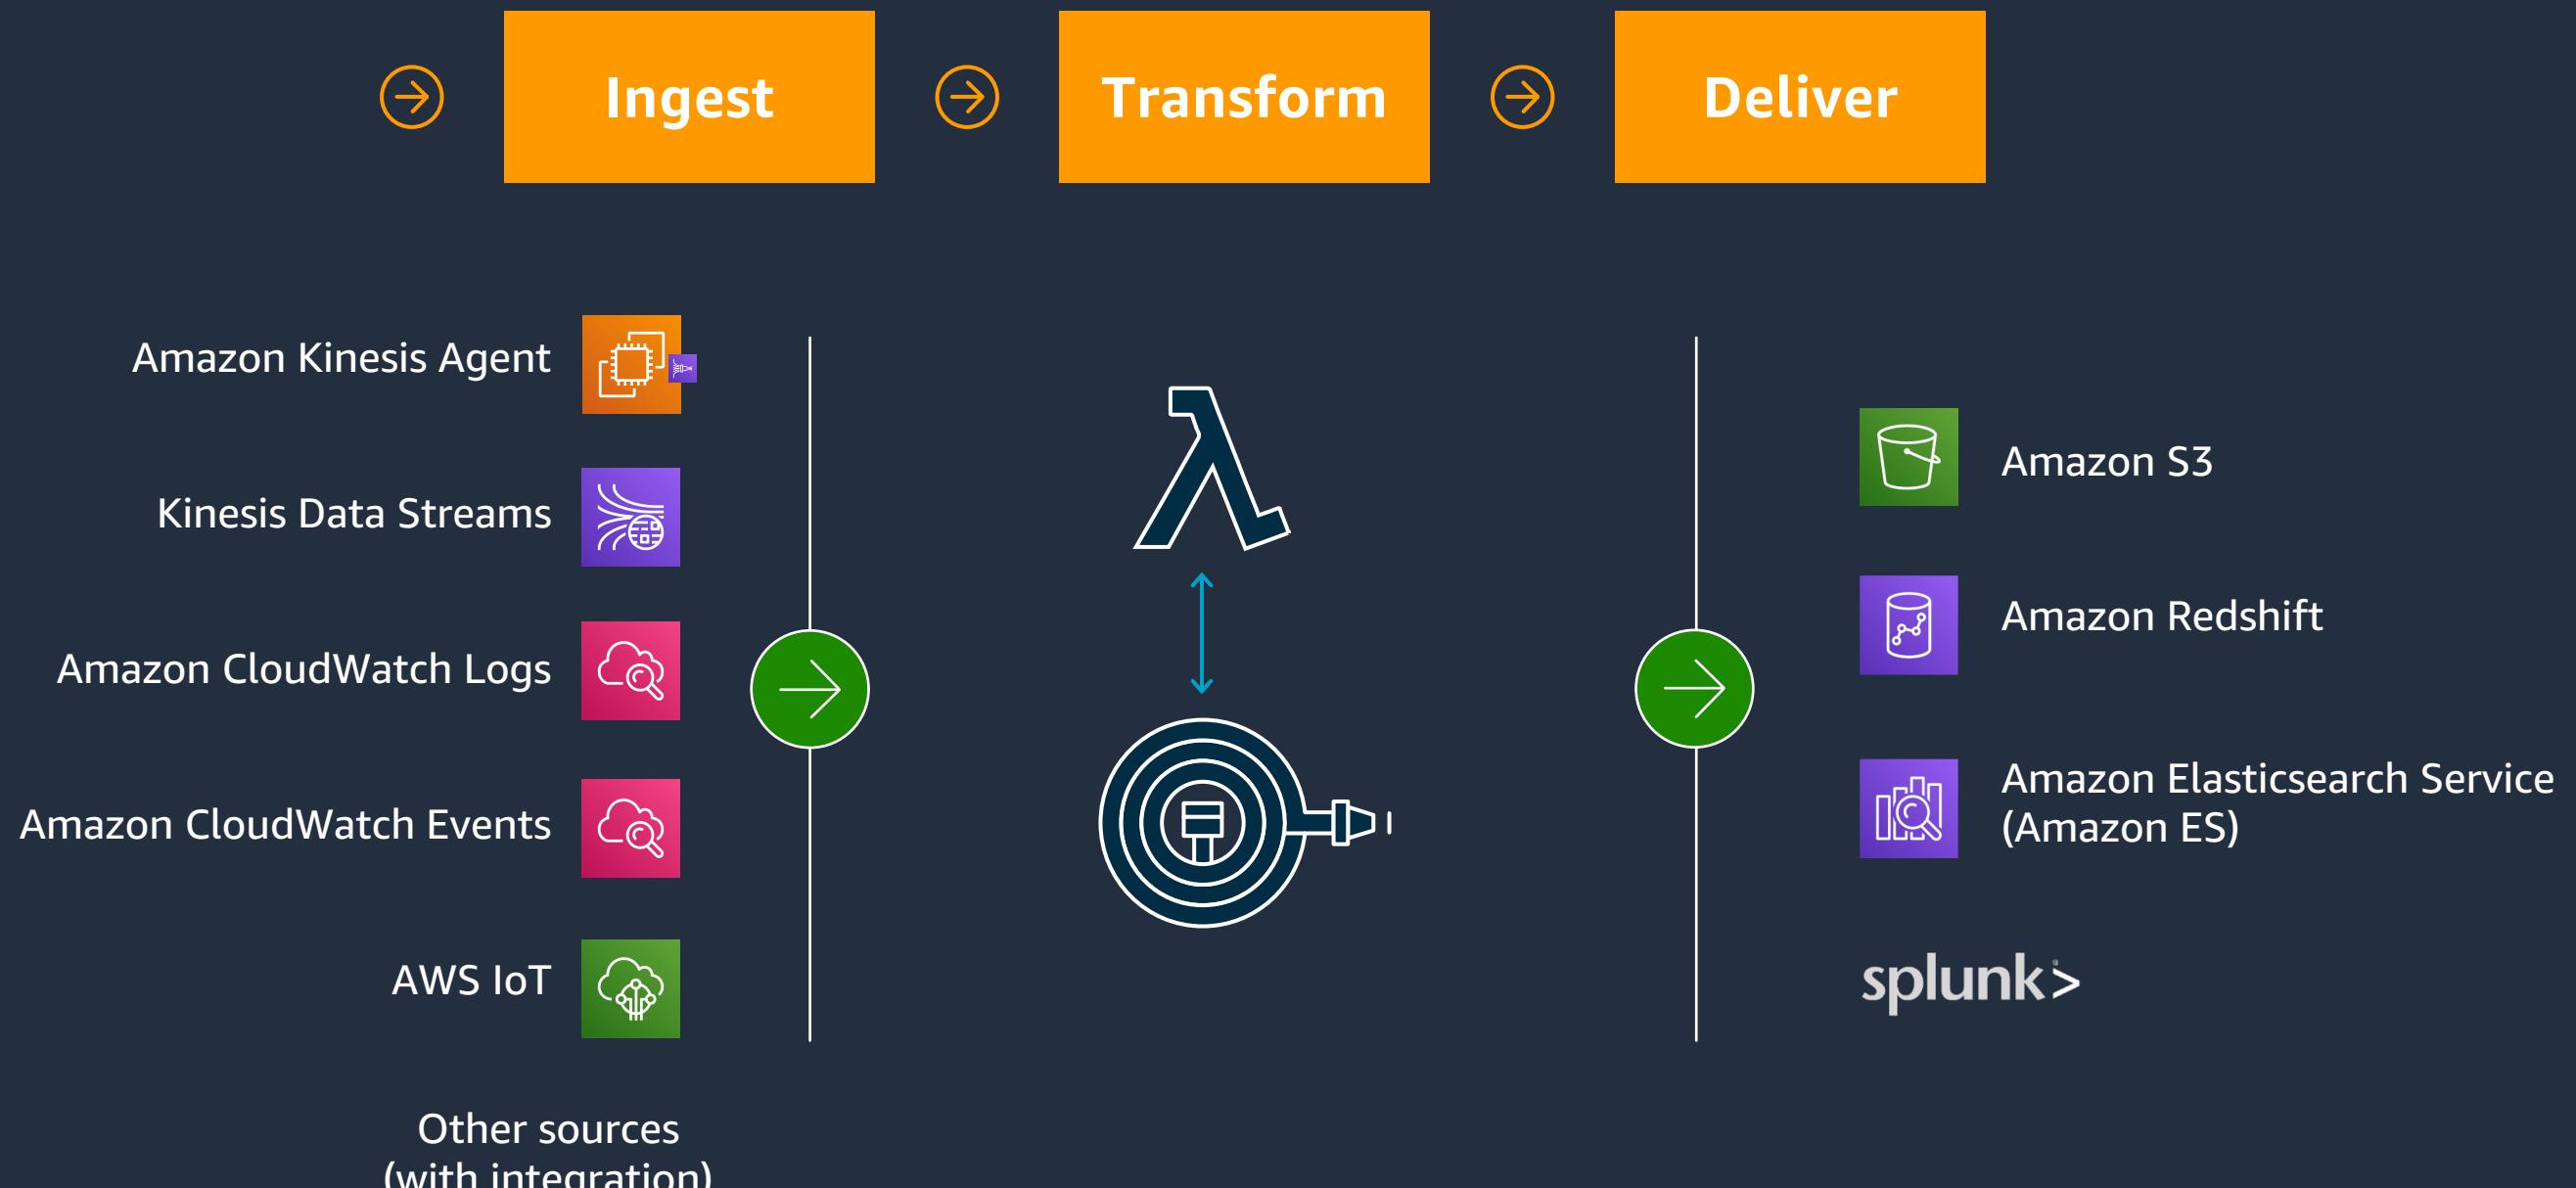


Fast and consistent

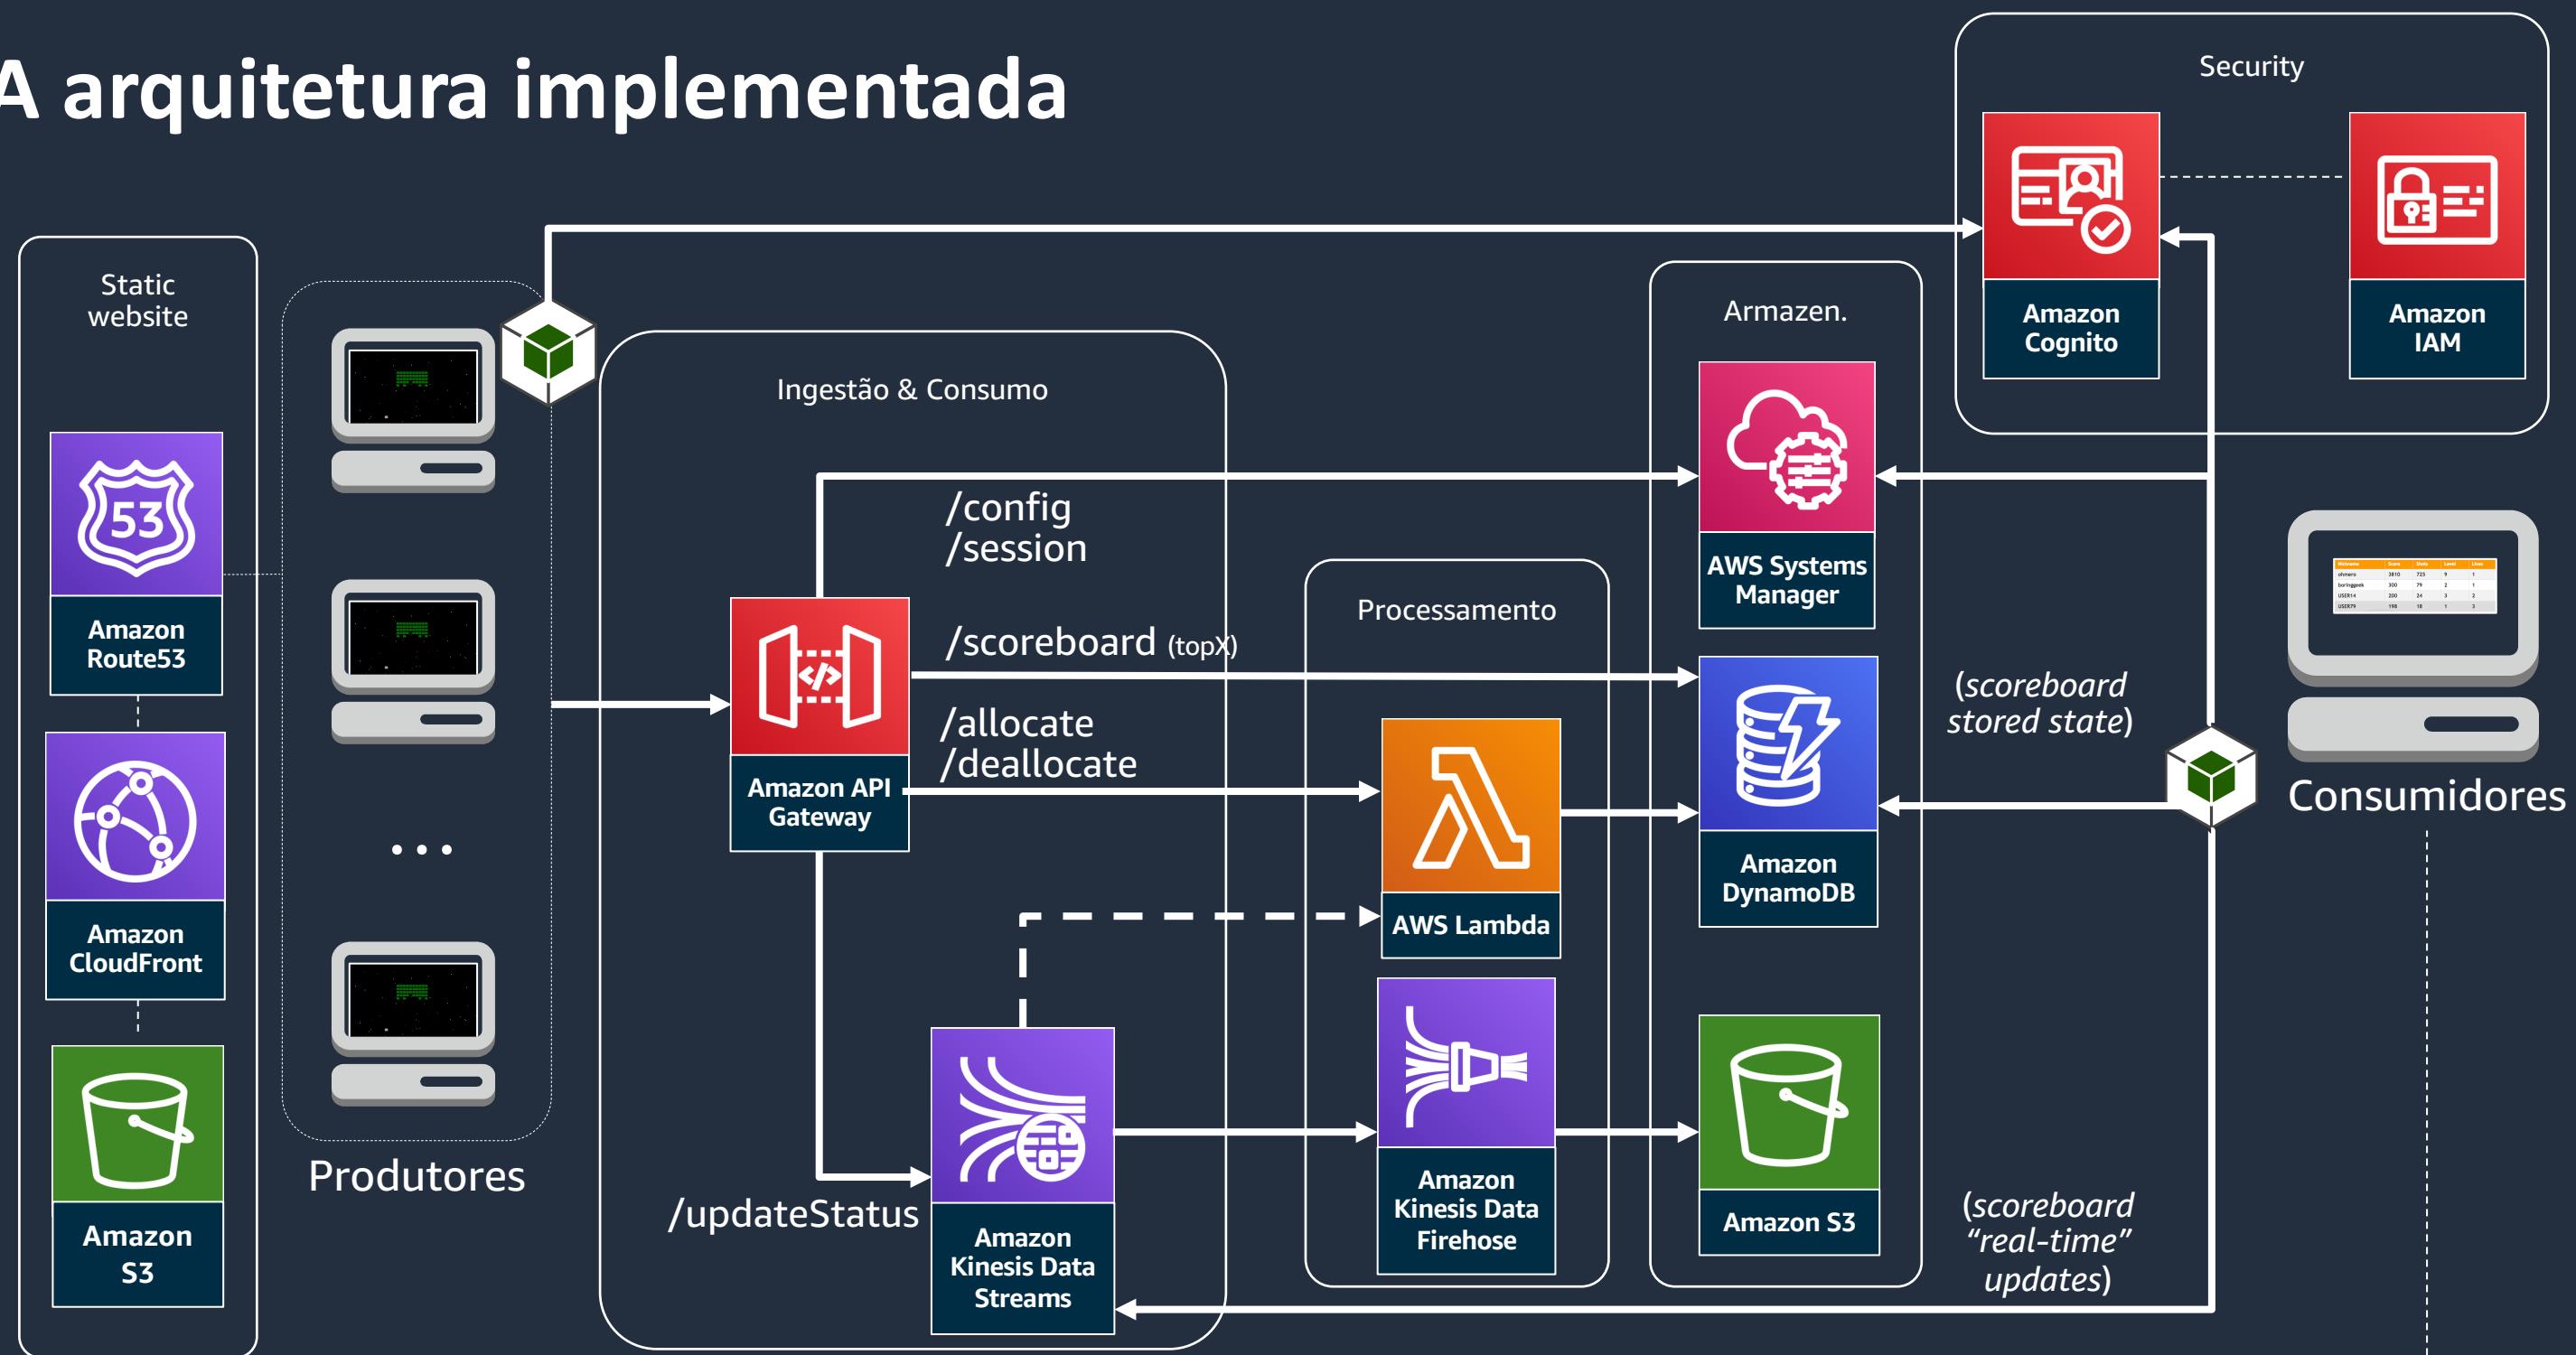
Access control

Event-driven programming

Amazon Kinesis Data Firehose

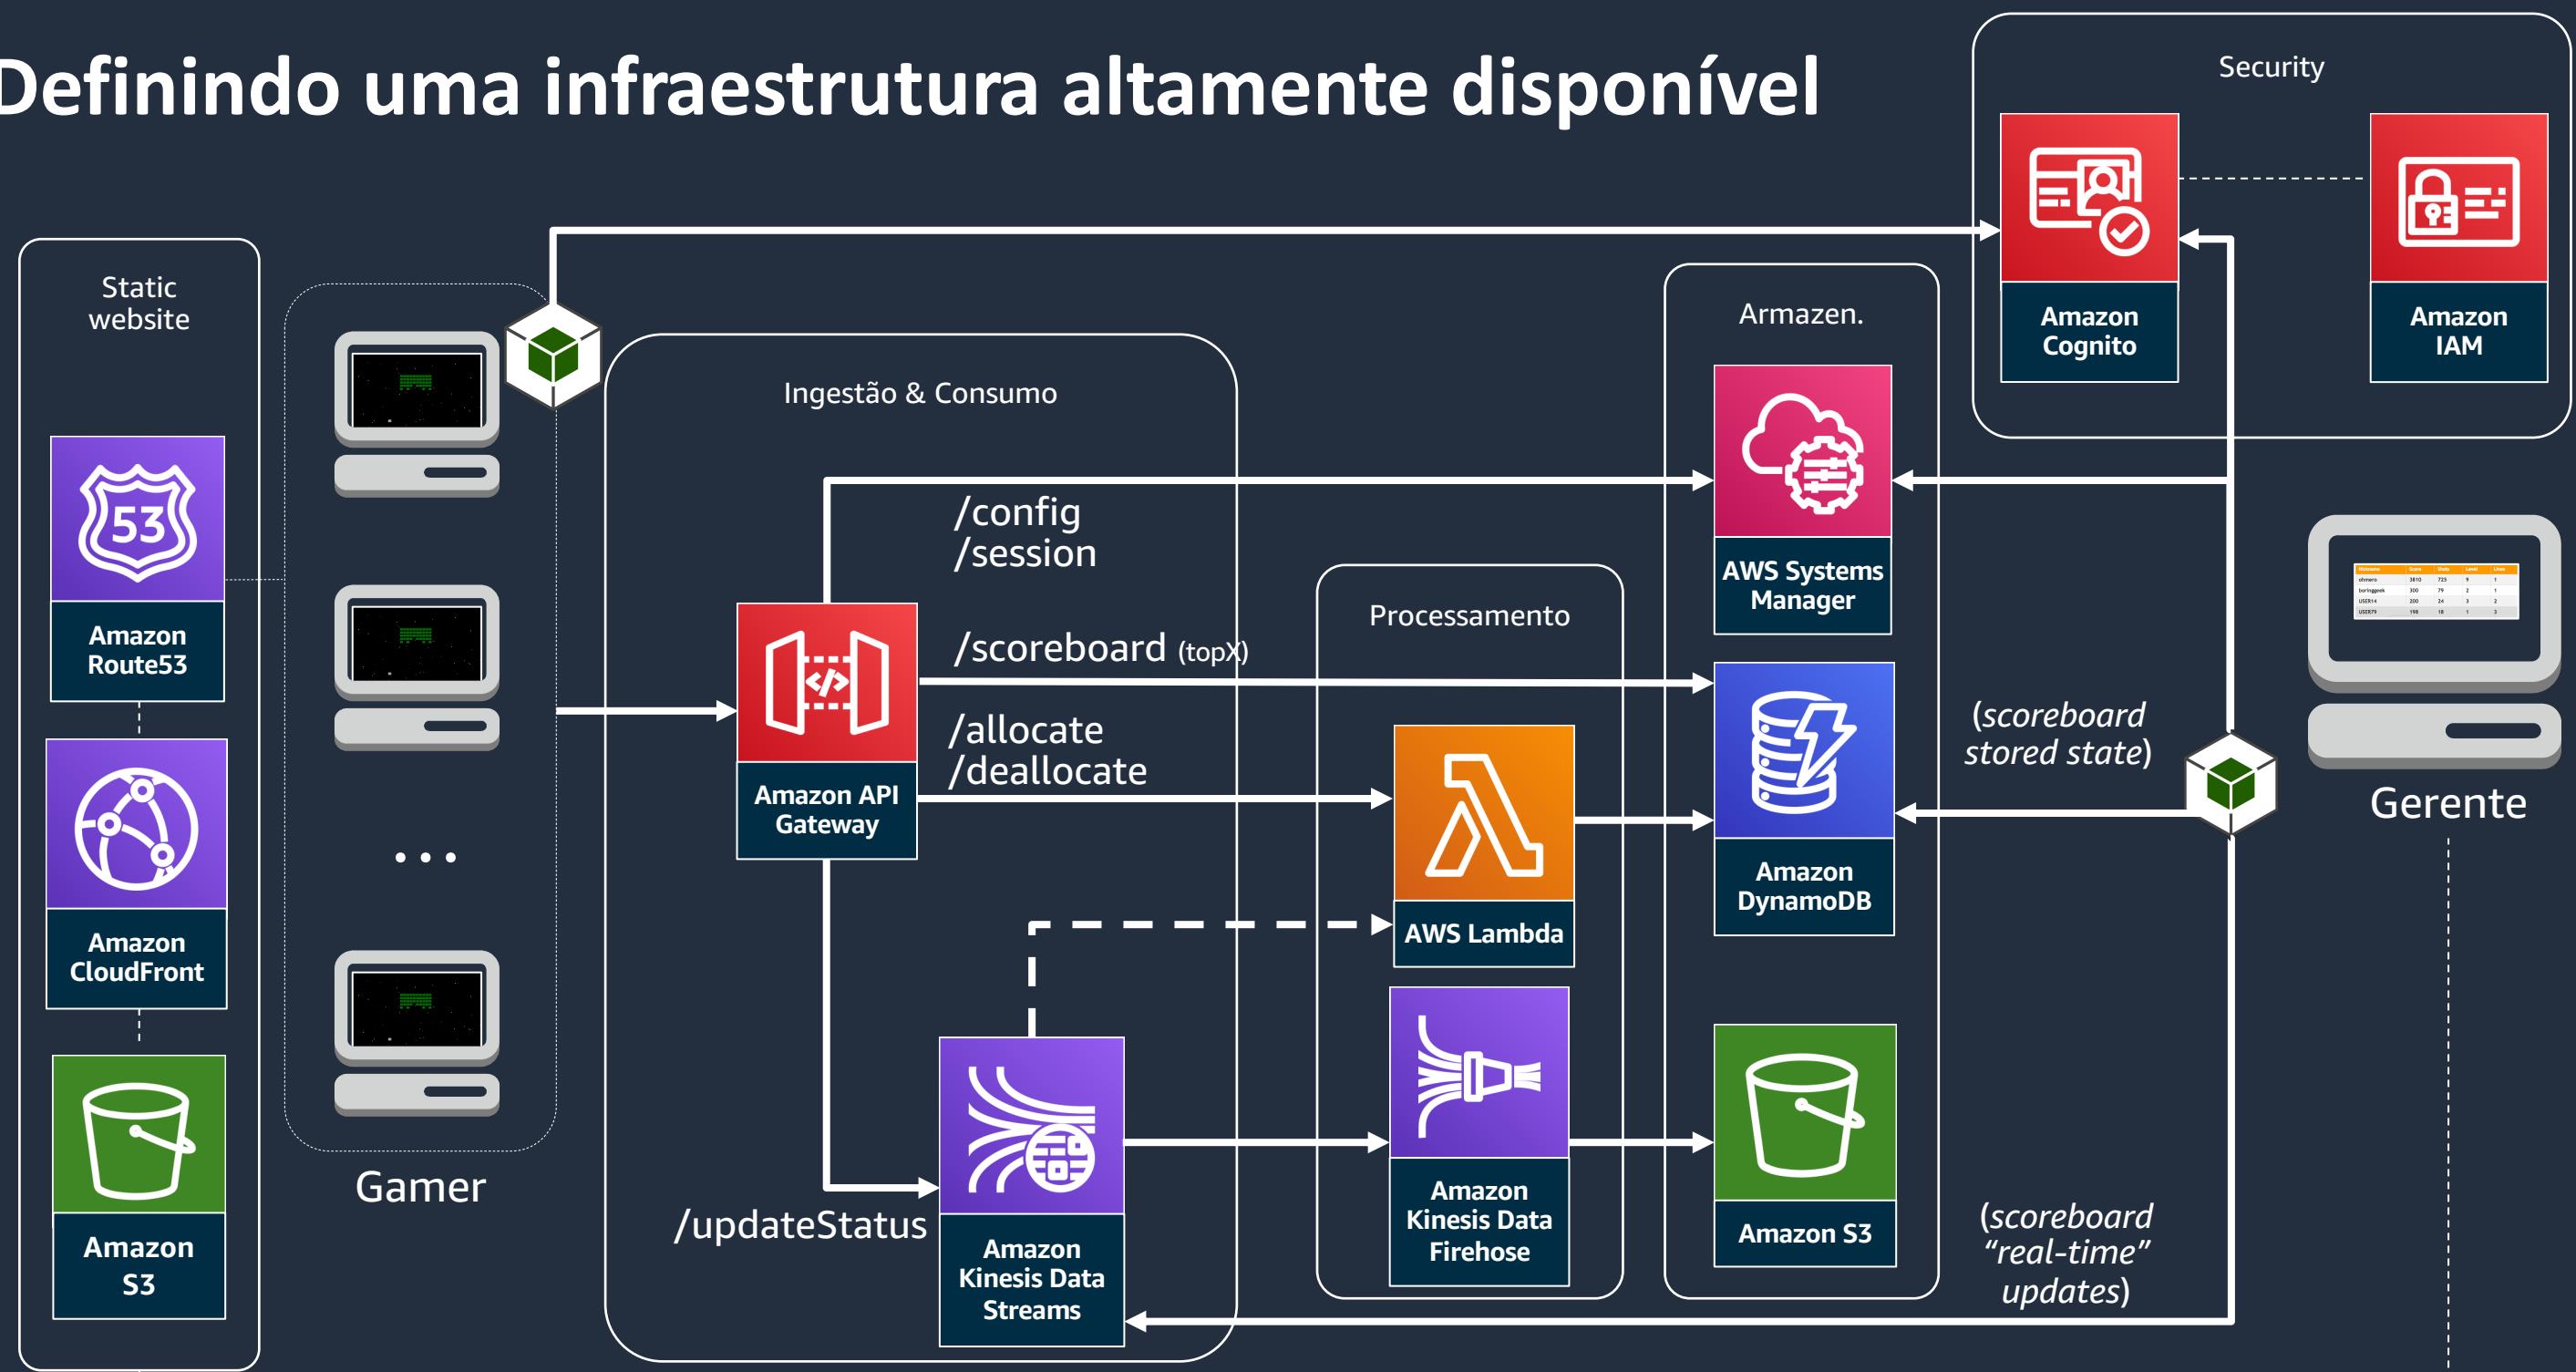


A arquitetura implementada



Revisando a arquitetura passo-a-passo

Definindo uma infraestrutura altamente disponível

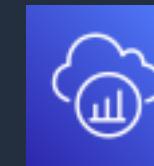
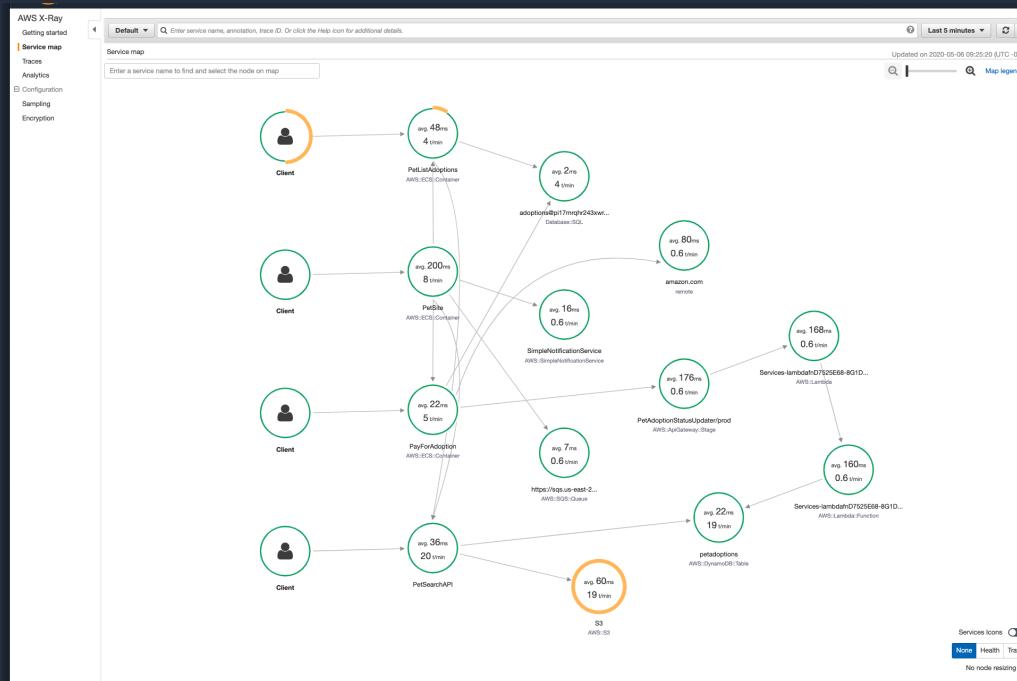


Operação

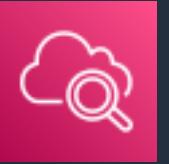
Responsabilidades da Operação

- Planejamento de Capacidade
- Updating/Upgrading/Patching
- Gerenciar Licenças
- Help/Service Desk
 - Atender a chamados de suporte (end-user e back-office)
 - Break&fix
 - Atender a pedidos de mudanças de configuração
 - Capacity update
- ...

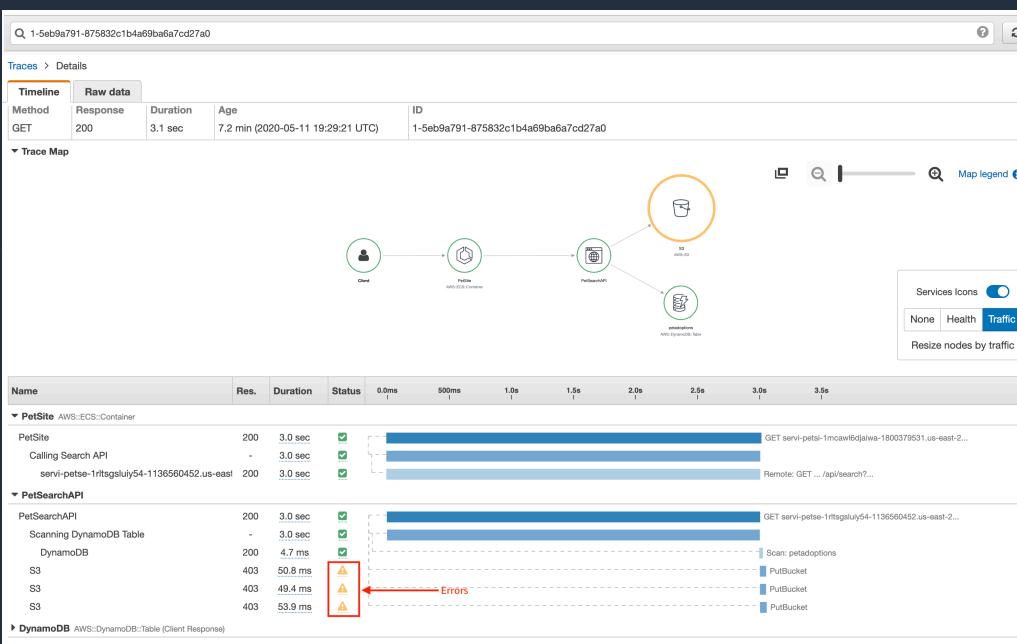
Observability



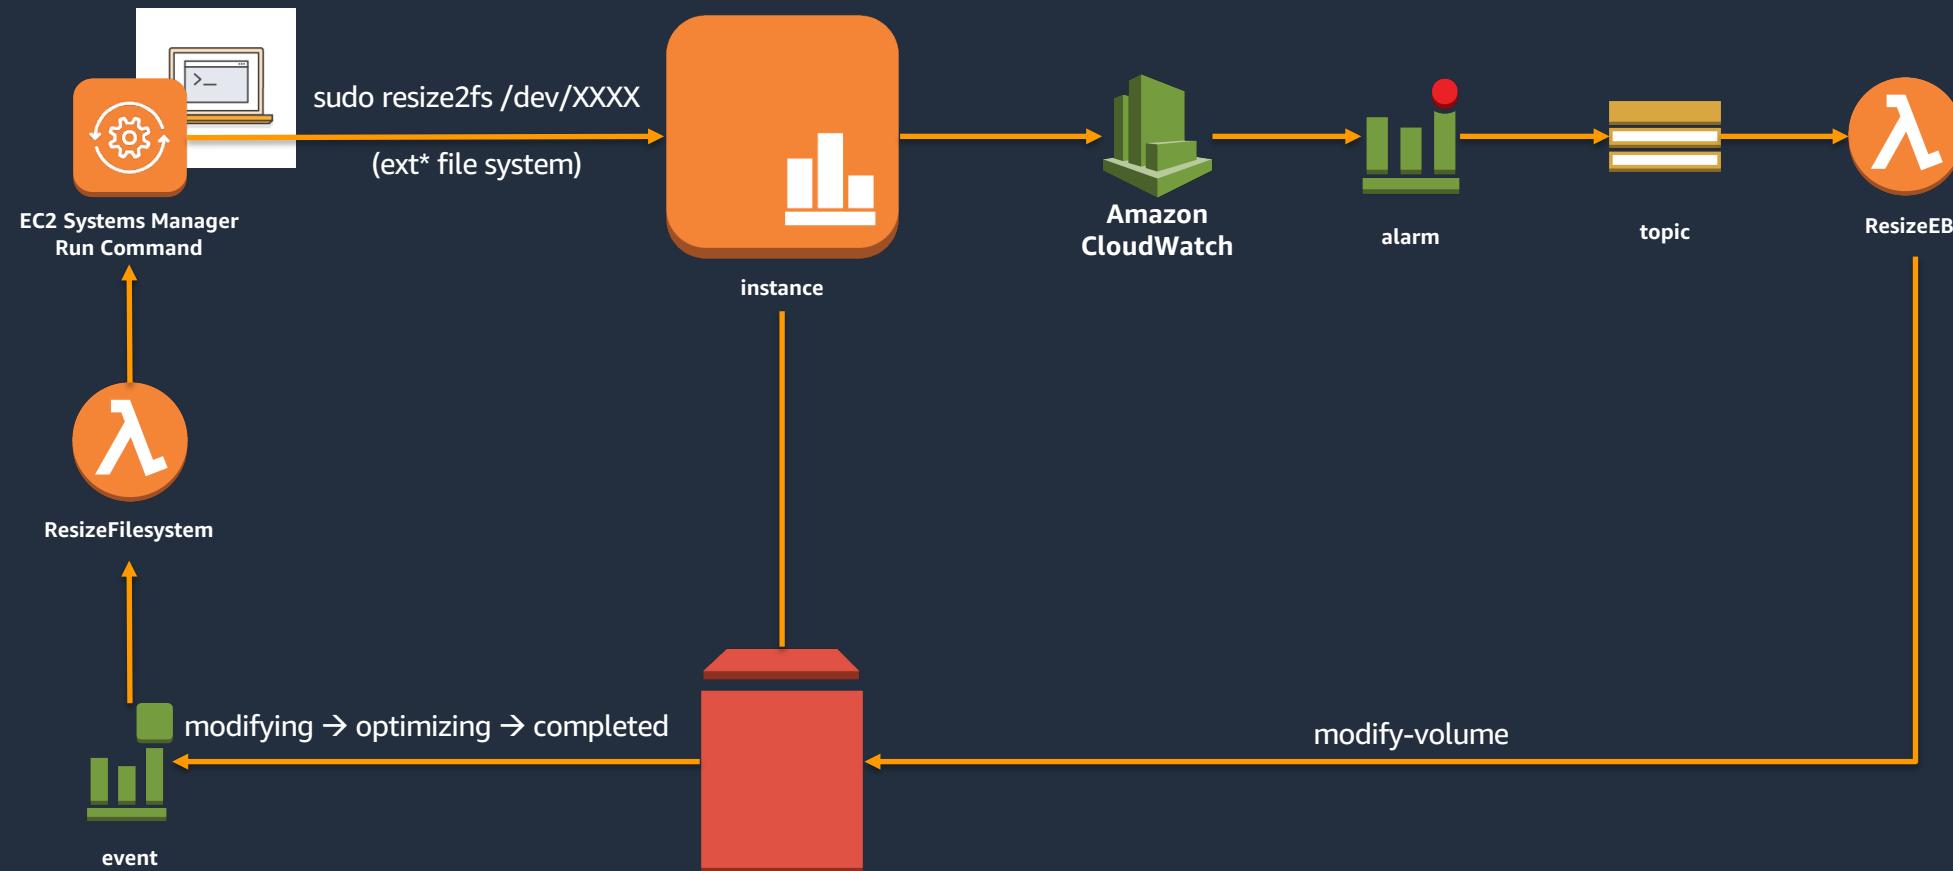
AWS X-Ray

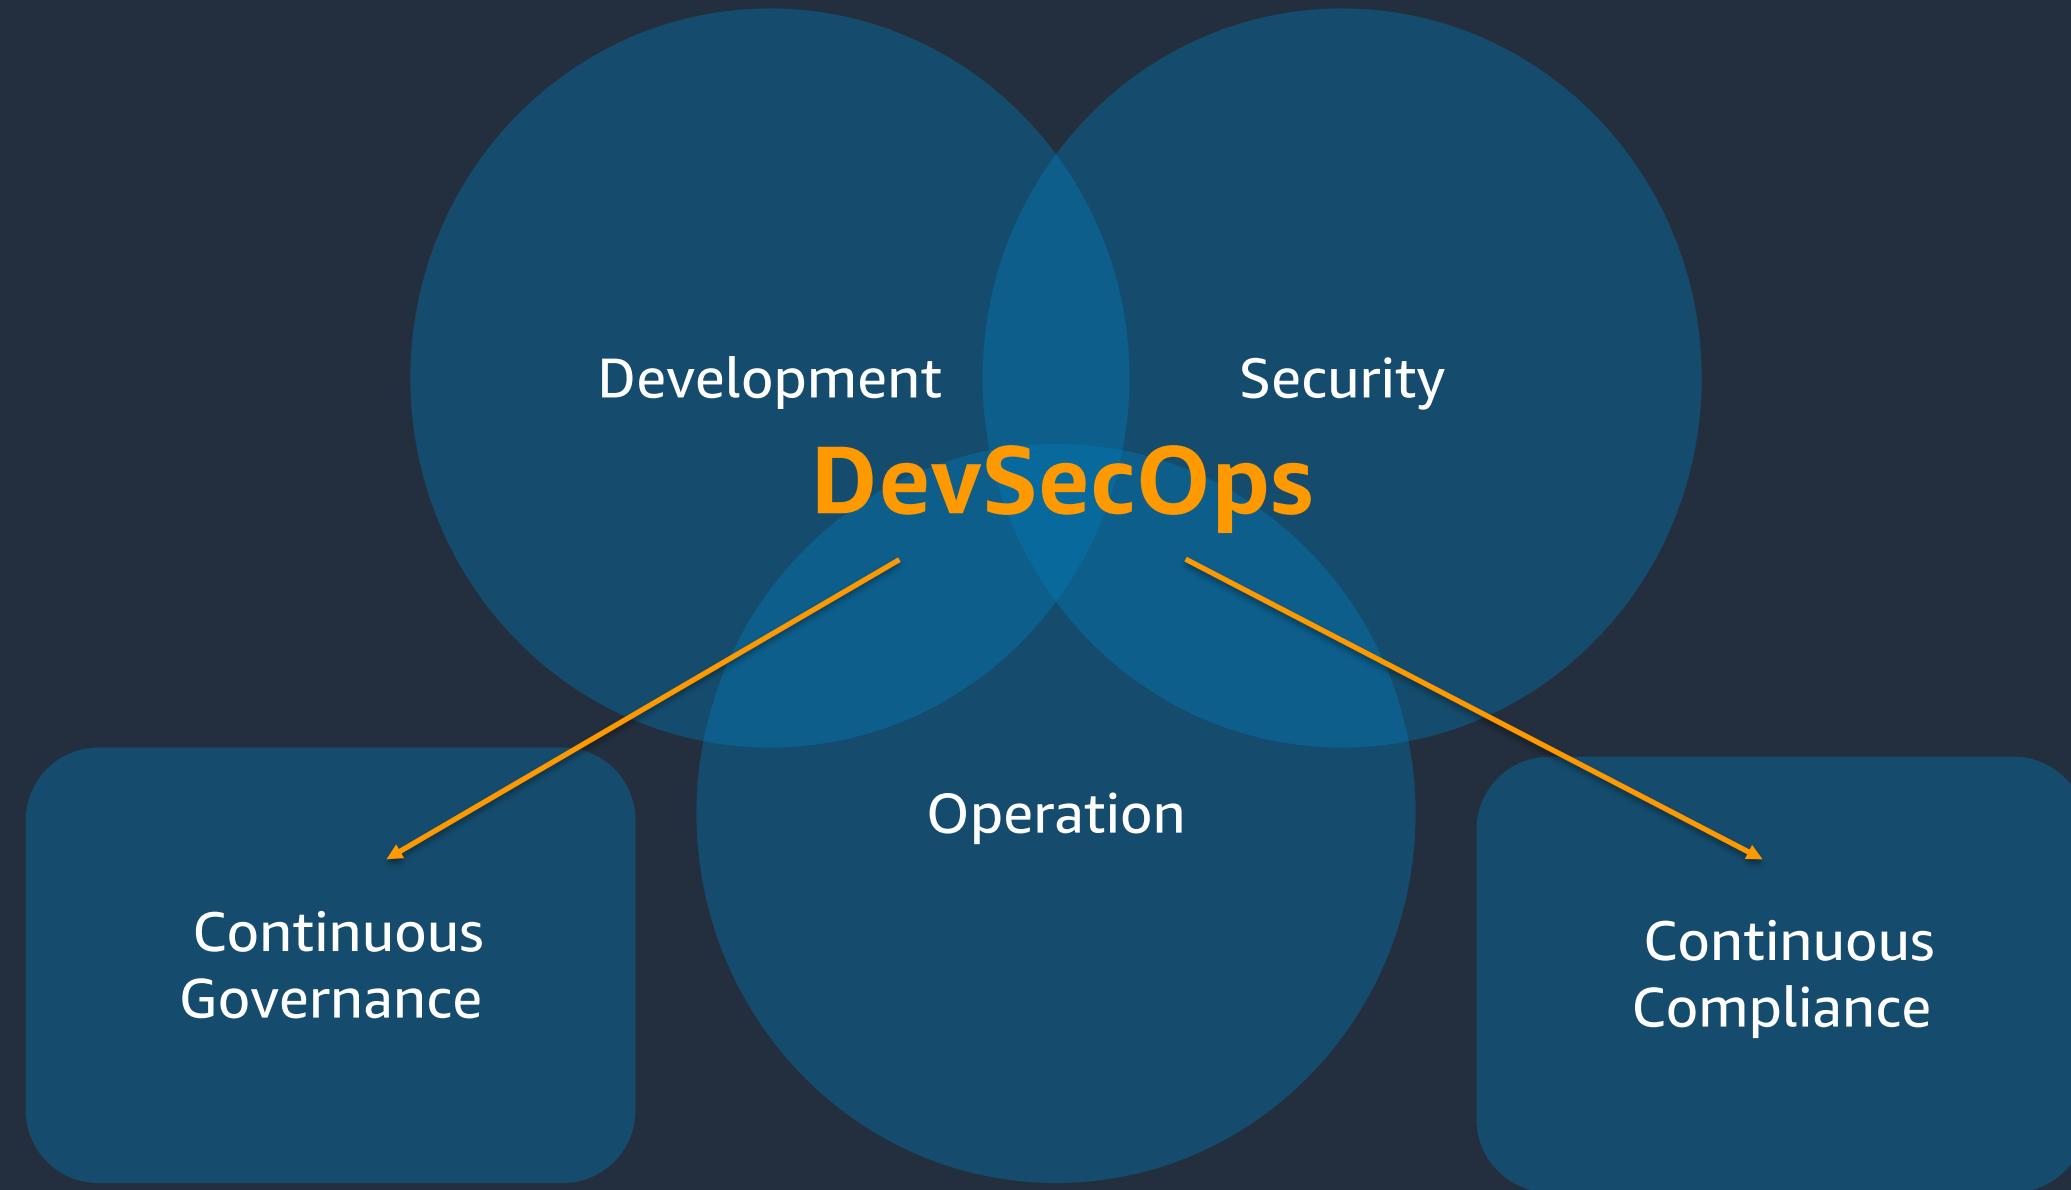


Amazon
CloudWatch



Operação reativa para Operação Proativa





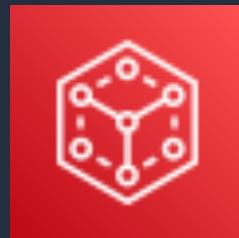
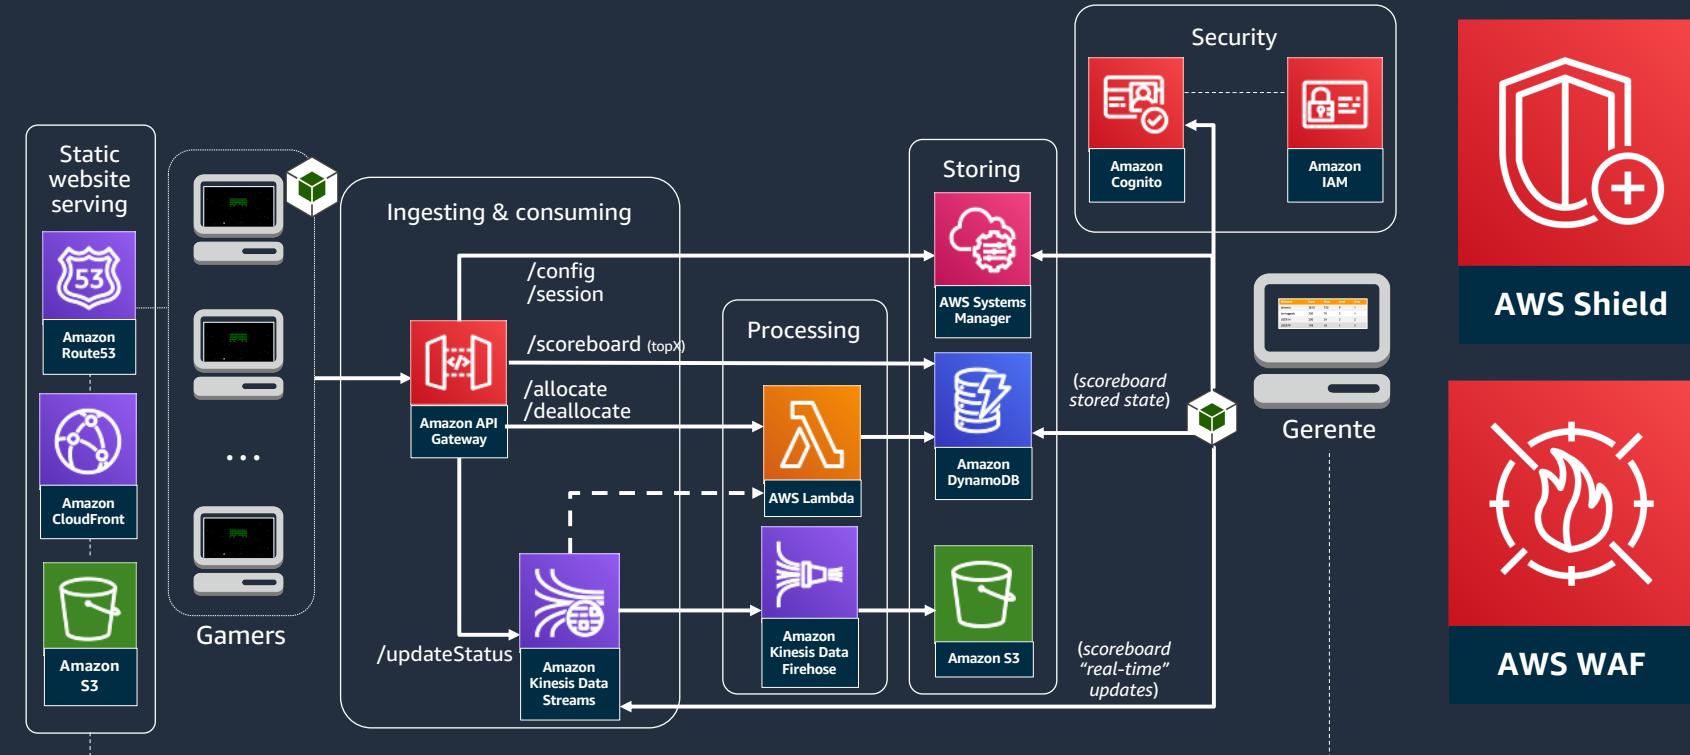
Teams of 8 to 10

-  Applications
-  Infrastructure
-  Leadership
-  Security
-  Operations



Achieve speed and agility
with **two-pizza teams**, while retaining
reporting structures

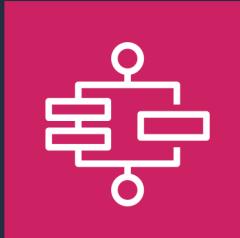
E a partir daqui?



AWS
Device Farm



AWS AppSync



AWS Step
Functions



Amazon
Simple Queue
Service



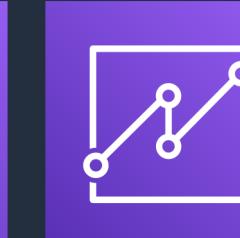
AWS IoT Core



AWS Glue



Amazon
Athena



Amazon
Quicksight



Amazon
Sagemaker

... e muito mais

Clientes usando e construindo soluções Serverless



StreamAlert by Airbnb

StreamAlert is a serverless, real-time data analysis framework which empowers you to ingest, analyze, and alert on data from different sources and alerting logic you define.

[Visit StreamAlert GitHub >>](#)

[Visit Airbnb.io >>](#)



Cloud Custodian by Capital One

Cloud Custodian is a rules engine for managing public cloud accounts and resources. It allows users to define policies to enable a well managed cloud infrastructure and consolidates adhoc scripts into a lightweight and flexible tool.

[Visit Cloud Custodian GitHub >>](#)



Odin & Fenrir by Coinbase

Odin allows you to deploy 12-factor applications to AWS. Fenrir is a secure AWS SAM deployer that helps manage serverless projects and scale serverless across your teams and orgs.

[Visit Odin GitHub >>](#)

[Visit Fenrir GitHub >>](#)



ShadowReader by Edmunds

ShadowReader is a load-testing tool that replays production traffic to a destination of your choice by collecting traffic patterns from access logs.

[Visit ShadowReader GitHub >>](#)



Bender by Nextdoor

Bender provides an extendable Java framework for creating serverless ETL functions on AWS Lambda. Bender provides the interfaces necessary to build modules for all aspects of the ETL process.

[Visit Bender GitHub >>](#)



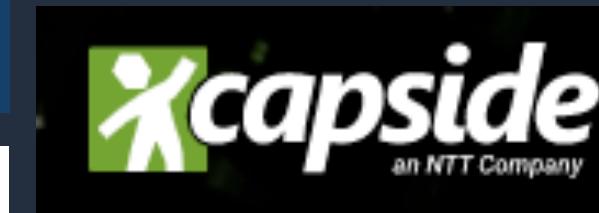
Jazz by T-Mobile

Jazz is a serverless development platform that accelerates the adoption of serverless technology within your enterprise. Jazz can help build functions, APIs, and static websites, and comes with CI/CD by default.

[Visit Jazz GitHub >>](#)



Partners no Brasil



#GoBuild!