



Convince your boss to go Serverless

by Vadym Kazulkin, ip.labs GmbH

Contact



Vadym Kazulkin

ip.labs GmbH Bonn, Germany

Co-Organizer of the Java User Group Bonn
and Serverless Bonn Meetup



v.kazulkin@gmail.com



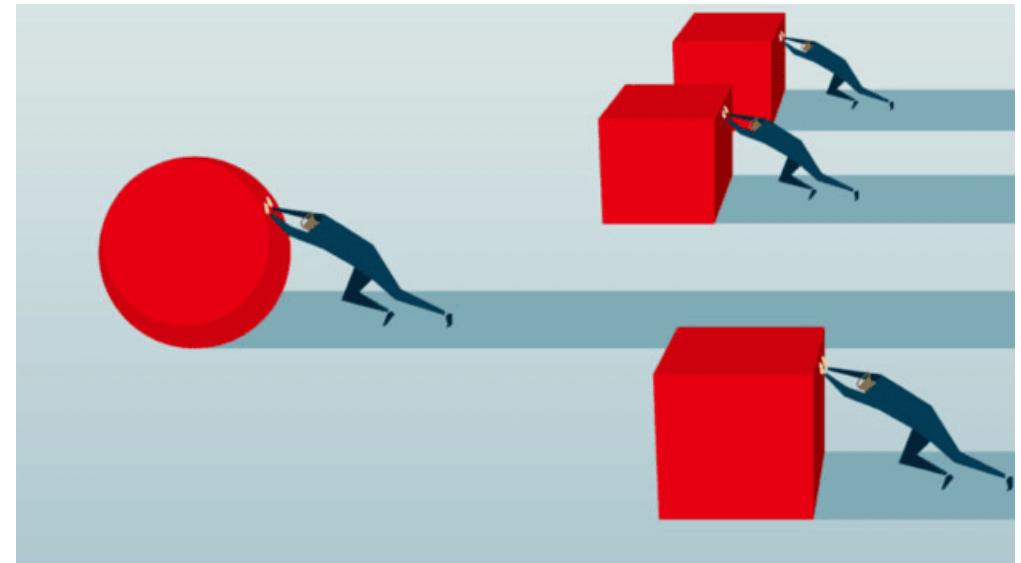
<https://www.linkedin.com/in/vadymkazulkin/>



@VKazulkin
@ServerlessBonn (Meetup)

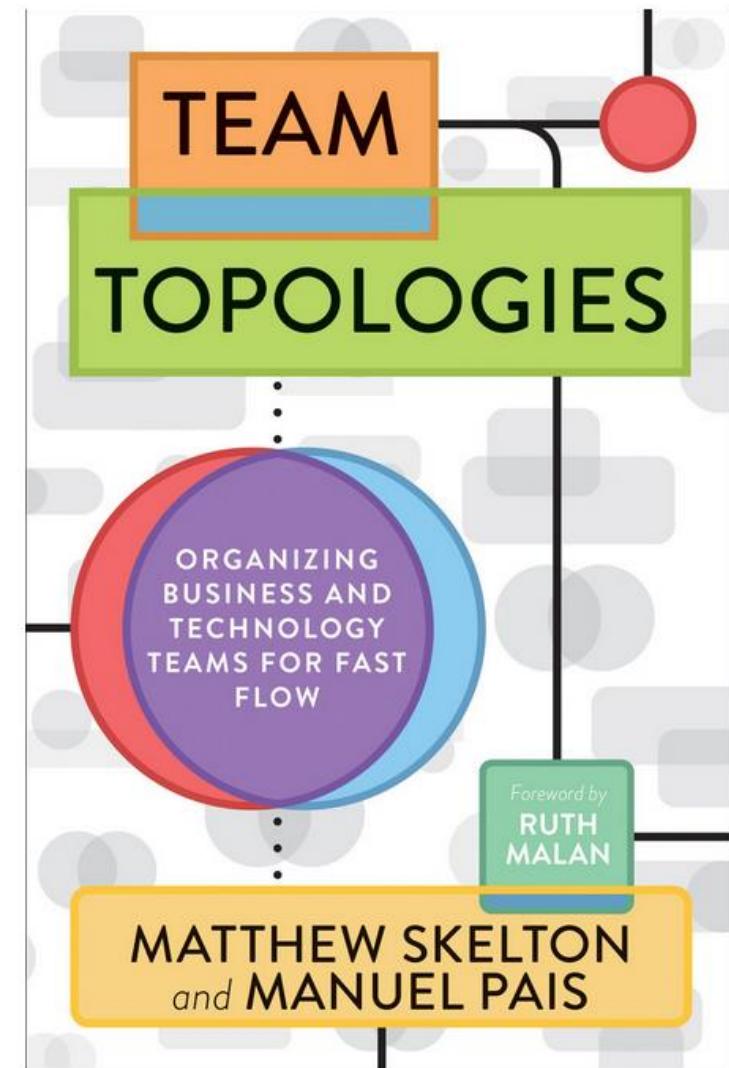


Let's talk about the
challenges of the
software development in
general first



Cognitive Load –
the total amount of mental effort
being used in the working memory

- Intrinsic
- Extraneous
- Germane



<https://teamtopologies.com/>



Cognitive Load

- **Intrinsic**
 - How to write a Java class or use a framework (Spring)
- Extraneous
- Germane

Cognitive Load

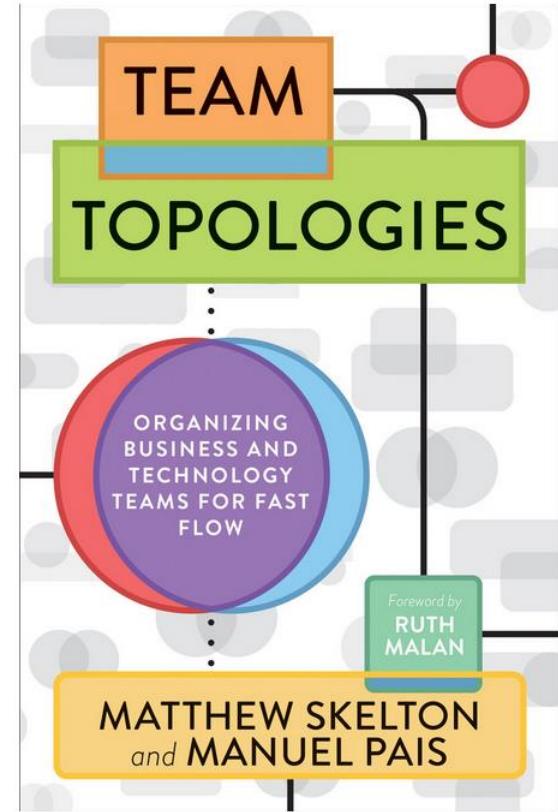
- Intrinsic
- **Extraneous**
 - How to automate tests (unit, integration, end-to-end, web, desktop, mobile)
 - How to build, package, deploy and run my application
 - How to configure monitoring, alerting, auto-scaling, logging and tracing
 - How to operate and maintain infrastructure
 - How to build-in fault-tolerance and resiliency
 - How to make the hardware, networking and application secure
- Germane

Cognitive Load

- Intrinsic
- Extraneous
- Germane
 - Domain Knowledge (payment, e-commerce)
 - Business processes and workflows

Cognitive Load

- Intrinsic ->
become fluent in it
- Extraneous ->
minimize amount of what we implement/operate/support/own by ourselves
- Germane ->
minimize amount of what we have to implement by ourselves



What our boss wants from us?

... that we are productive

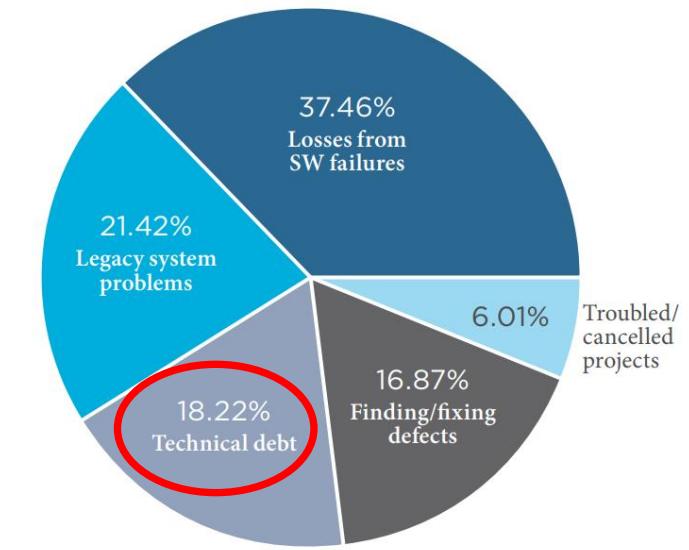
Productivity

We are productive if we regularly ship products,
which are successfully used by our customers

What is holding us back from being productive?

Technical Debt - reflects the implied cost of additional rework caused by choosing an easy (limited) solution now instead of using a better approach that would take longer

AREAS OF COST RELATING TO POOR IT/SOFTWARE QUALITY IN THE US



"The Cost of Poor Quality Software in the US: A 2018 Report"

<https://www.it-cisq.org/the-cost-of-poor-quality-software-in-the-us-a-2018-report/The-Cost-of-Poor-Quality-Software-in-the-US-2018-Report.pdf>

Technical Debt

- Even a perfect solution can become the technical debt over the time
 - Version of programming language comes out of support (Java 8)
 - Security considerations forces us to upgrade one of our dependencies (library or web application server version)
 - One of our dependencies (e.g. to open source project) is discontinued

Technical Debt

Think of what can happen to your software over
the entire life cycle of our product

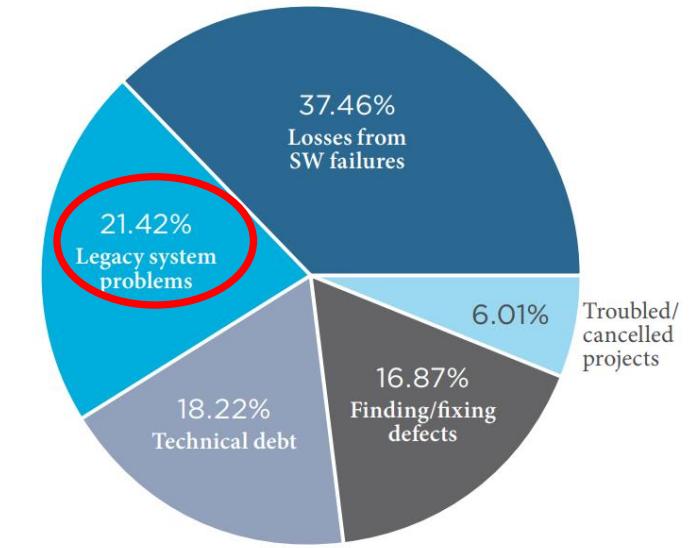
Technical Debt

- is related to amount of code written
- is related to amount of dependencies used
 - open source projects, programming languages, databases, (web) application servers

Legacy System

Legacy Systems are systems that can't evolve

AREAS OF COST RELATING TO POOR IT/SOFTWARE QUALITY IN THE US



"The Cost of Poor Quality Software in the US: A 2018 Report"

<https://www.it-cisq.org/the-cost-of-poor-quality-software-in-the-us-a-2018-report/The-Cost-of-Poor-Quality-Software-in-the-US-2018-Report.pdf>

Evolutionary Architecture – supports guided, incremental change across multiple dimensions

- Incremental change
- Appropriate architectural coupling
- Fitness functions

O'REILLY®

Building Evolutionary Architectures

SUPPORT CONSTANT CHANGE



Neal Ford, Rebecca Parsons & Patrick Kua



"Architectural Coupling" <https://learning.oreilly.com/library/view/building-evolutionary-architectures/9781491986356/ch04.html>

Evolutionary Architecture – Fitness functions

- **Source code metrics** (such as measuring cyclomatic complexity)
- **Unit tests** (% of coverage and % of success)
- **Performance metrics** (such as API latency or throughput)
- **Security** (encryption at rest, e.g. checking that all S3 buckets have encryption enabled, or automatic key rotation for all external APIs, with tools such as the AWS Secrets Manager)
- **ArchUnit, Sonar, CI/CD Tools**
 - CodeCommit,...CodeDeploy, Jenkins

O'REILLY®

Building Evolutionary Architectures

SUPPORT CONSTANT CHANGE



Neal Ford, Rebecca Parsons & Patrick Kua



The Value Proposition of Serverless

But let's talk about of Total Cost of Ownership of the Serverless paradigm

TCO Full Picture



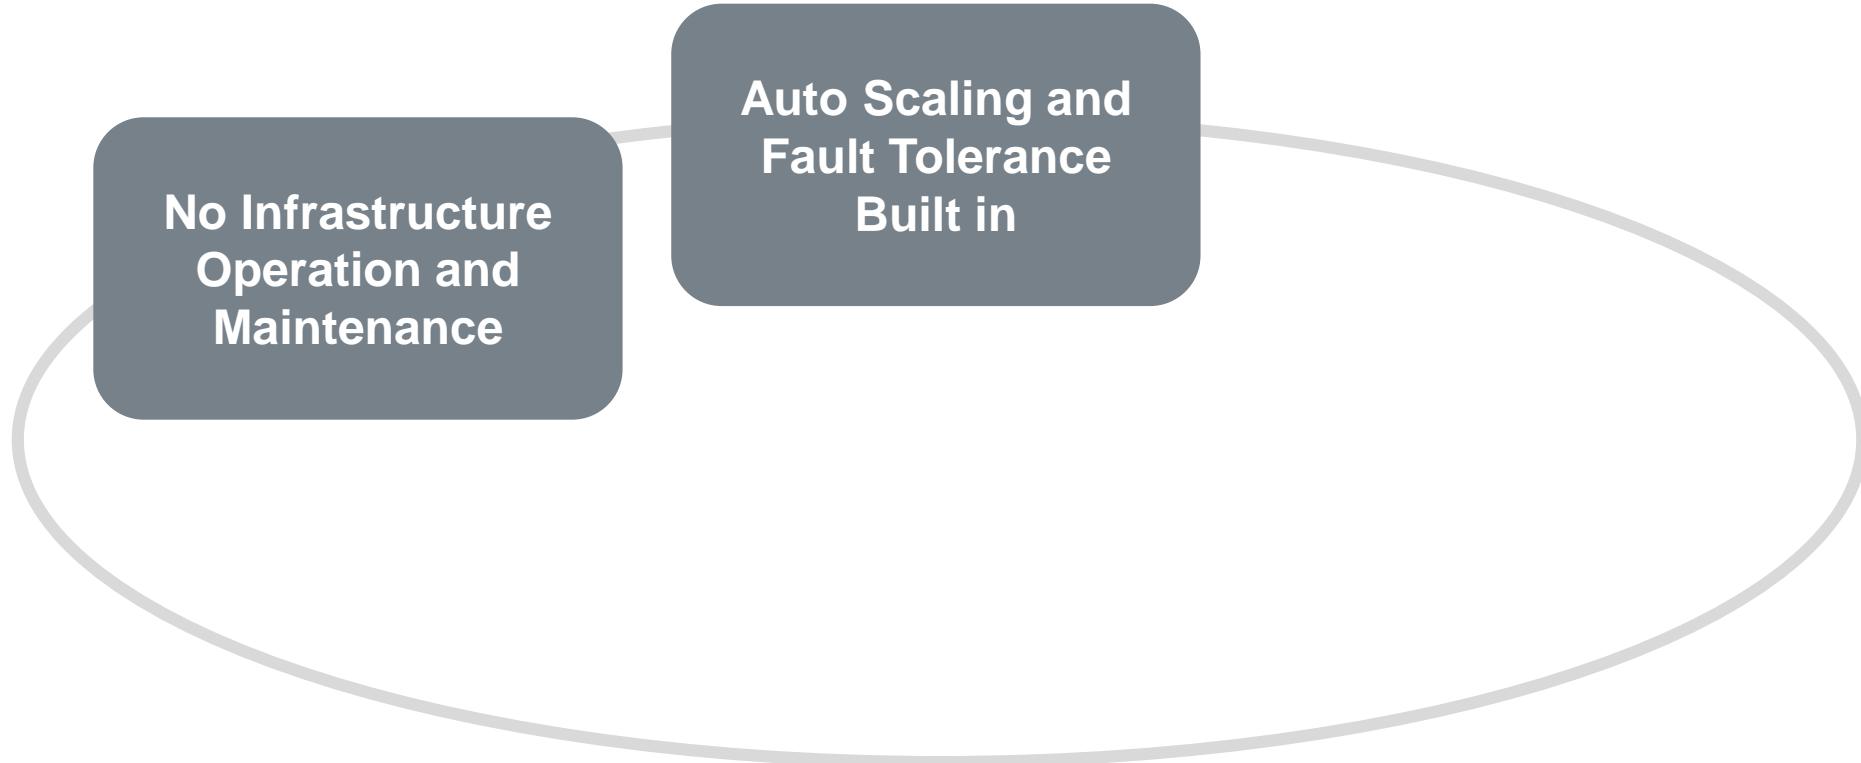
No Infrastructure
Operation and
Maintenance



No Infrastructure Maintenance

Is infrastructure maintenance and operation your core competency ?

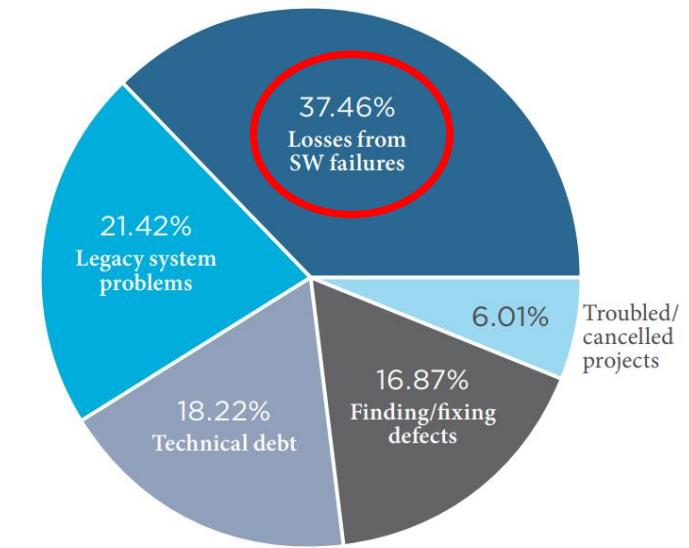
TCO Full Picture



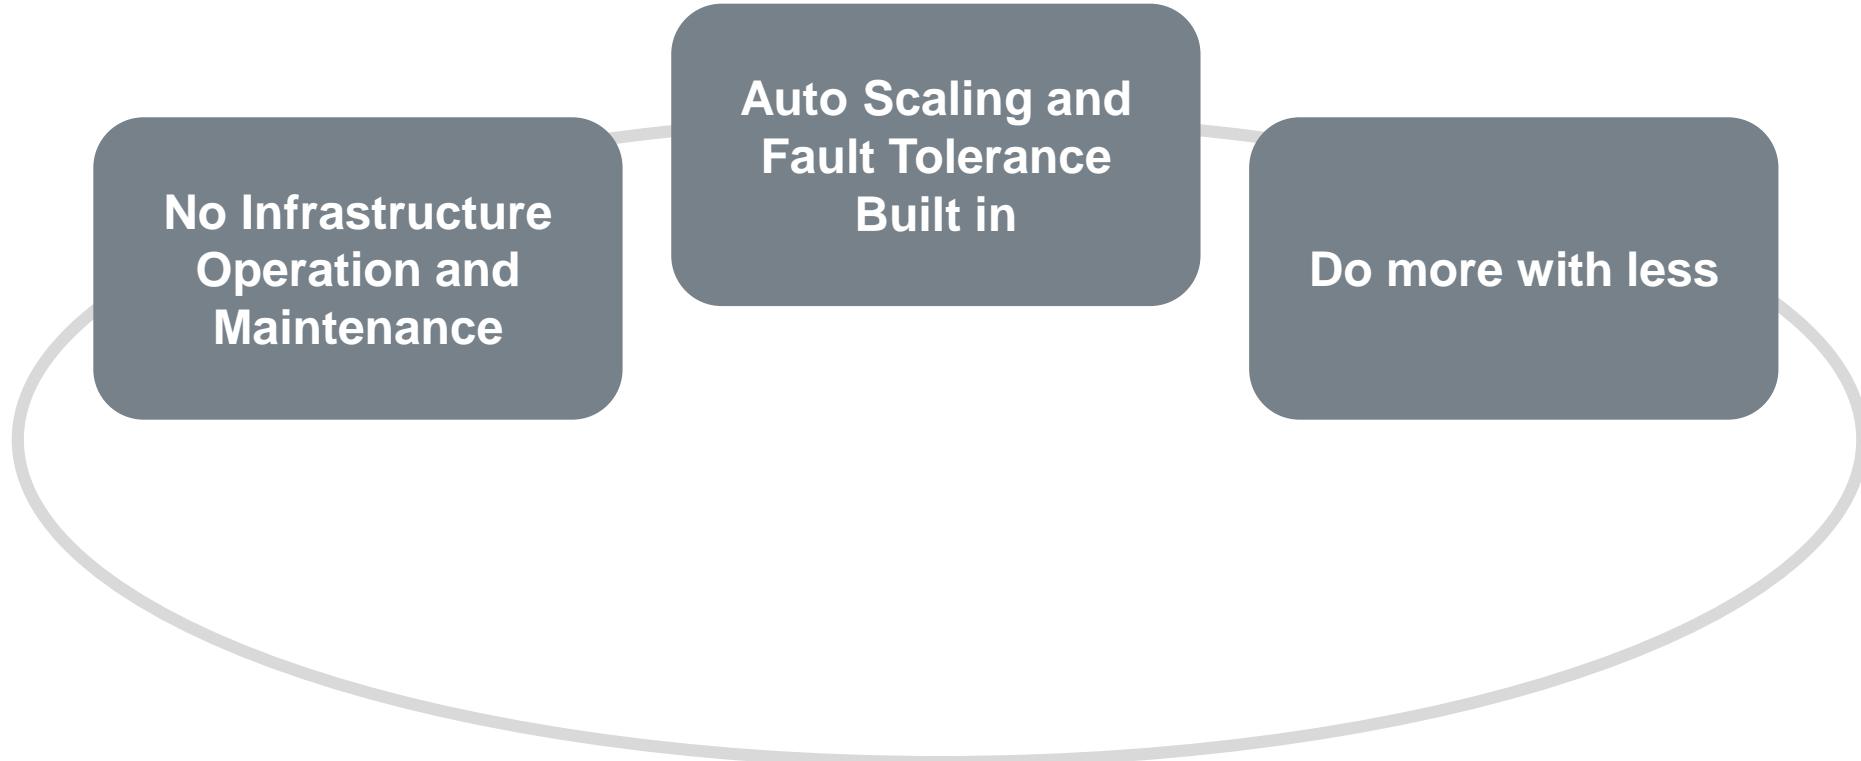
Auto Scaling And Fault Tolerance Built In

- Can you get capacity planning and auto scaling right?
- Do you want to solve the hard problem of fault tolerance by yourself?

AREAS OF COST RELATING TO POOR IT/SOFTWARE QUALITY IN THE US



TCO Full Picture

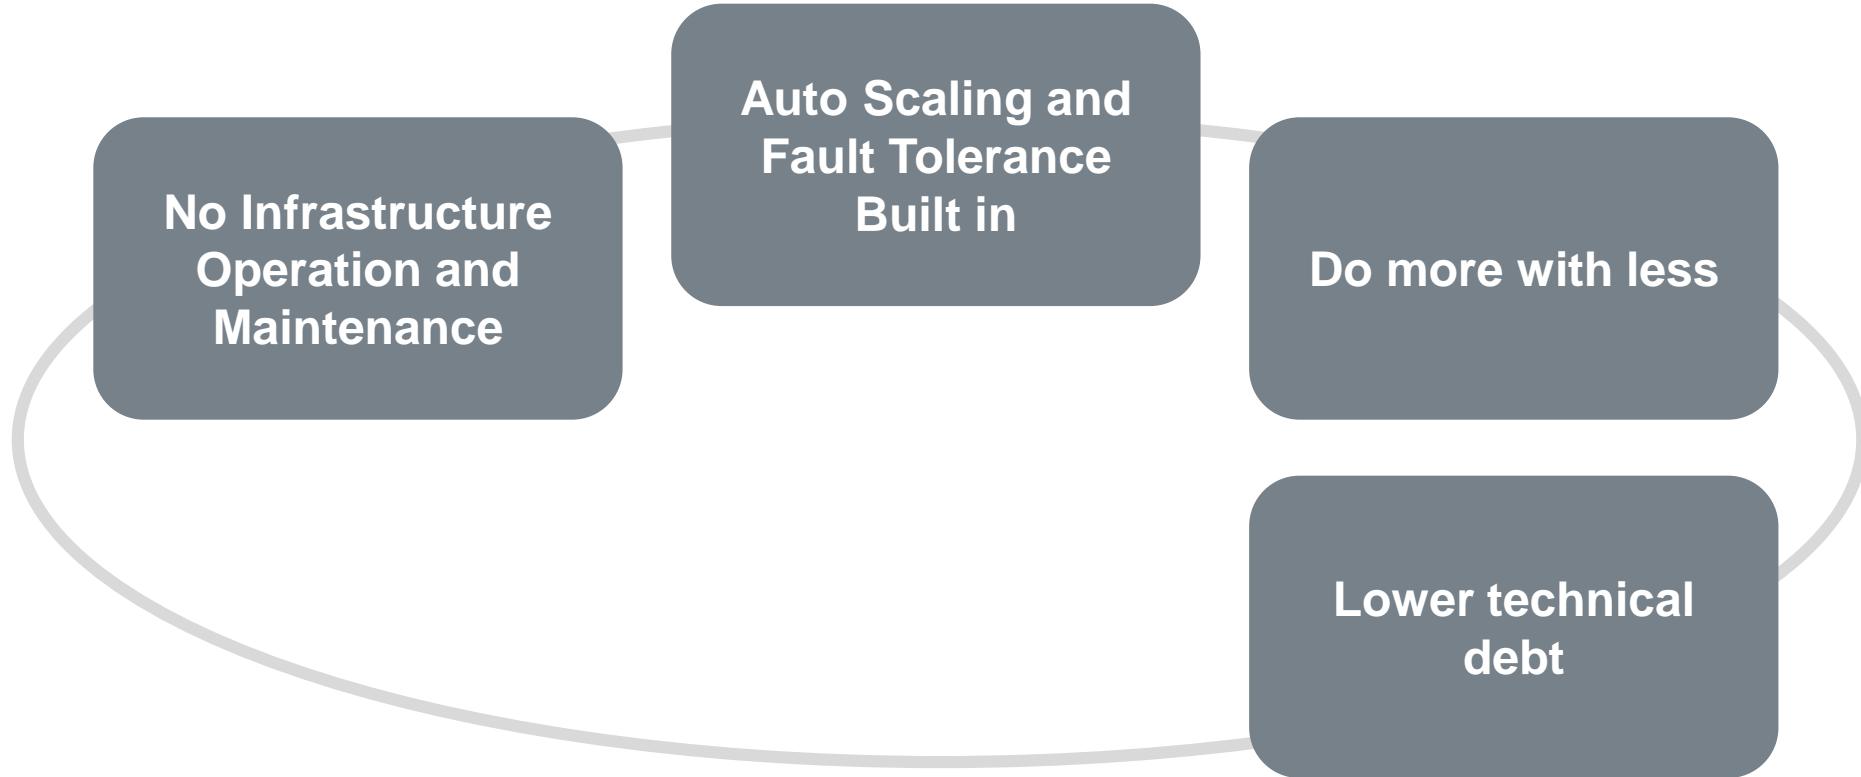


Do more with less

By heavily relying on the managed Serverless services you

- Need fewer engineers to start implementing your new product idea
- Can do more with the same amount of people

TCO Full Picture



Lower technical debt

- Whatever code you write today is always tomorrow's technical debt © Paul Johnston
- Less code means lower technical debt
- Time and effort required for maintaining the solution over its whole lifecycle is by far much more than for developing it



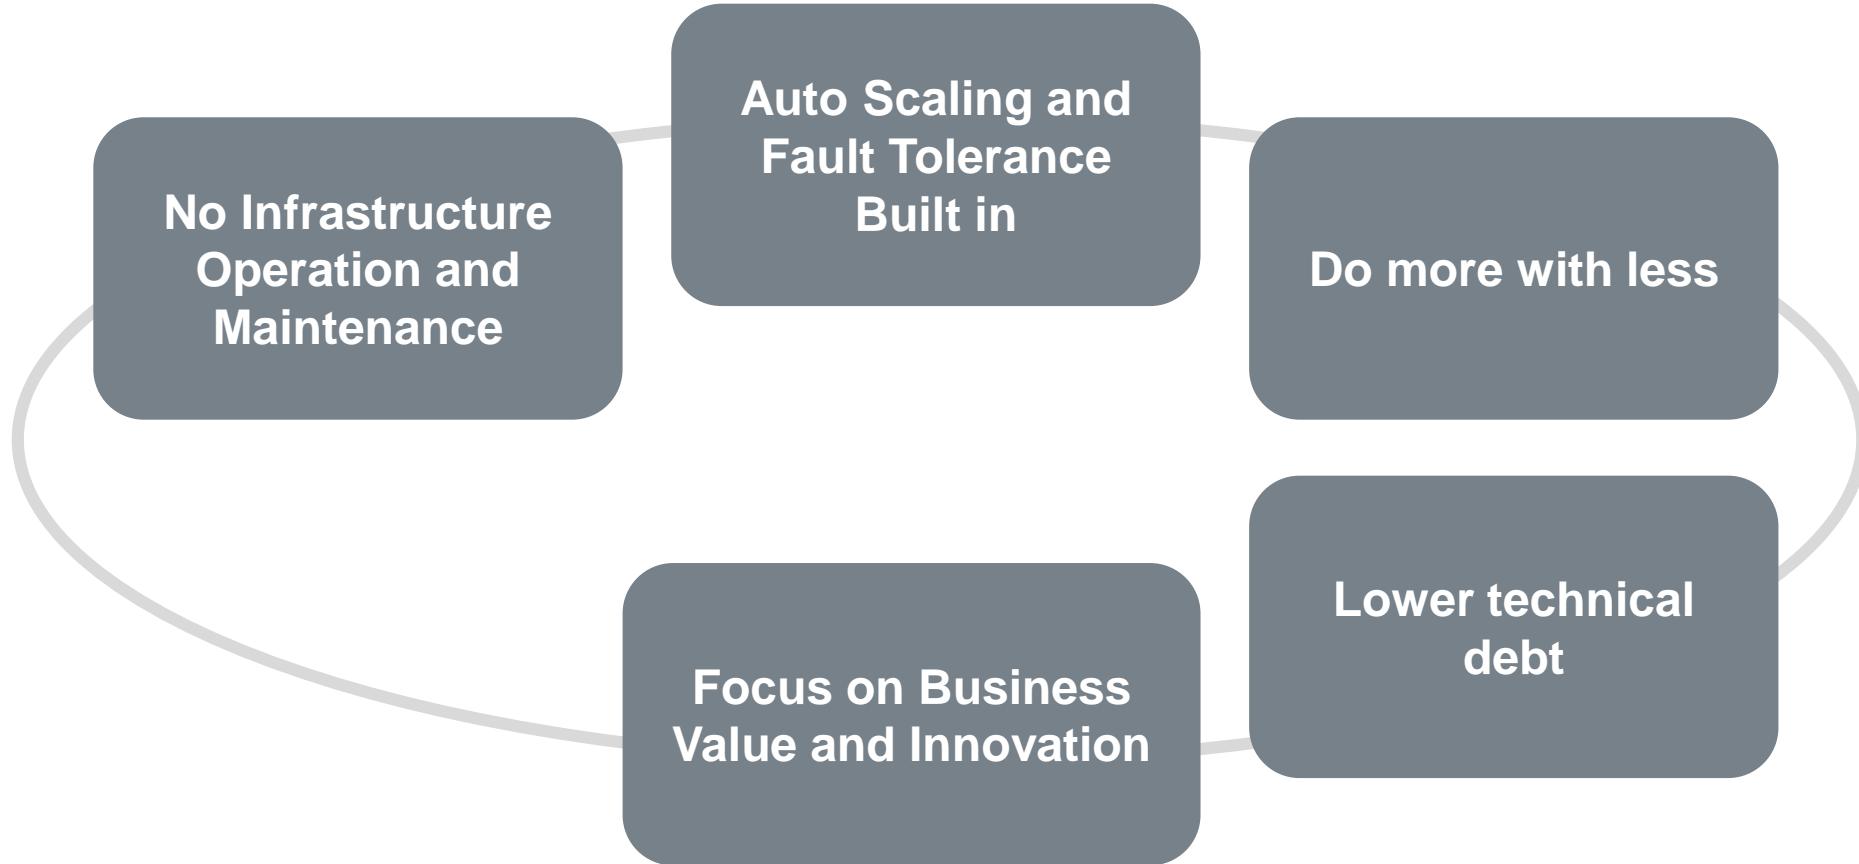
Jeff Atwood „The Best Code is No Code At All“ <https://blog.codinghorror.com/the-best-code-is-no-code-at-all/>

Paul Johnston “Cloud 2.0: Code is no longer King — Serverless has dethroned it”

<https://medium.com/@PaulDJohnston/cloud-2-0-code-is-no-longer-king-serverless-has-dethroned-it-c6dc955db9d5>



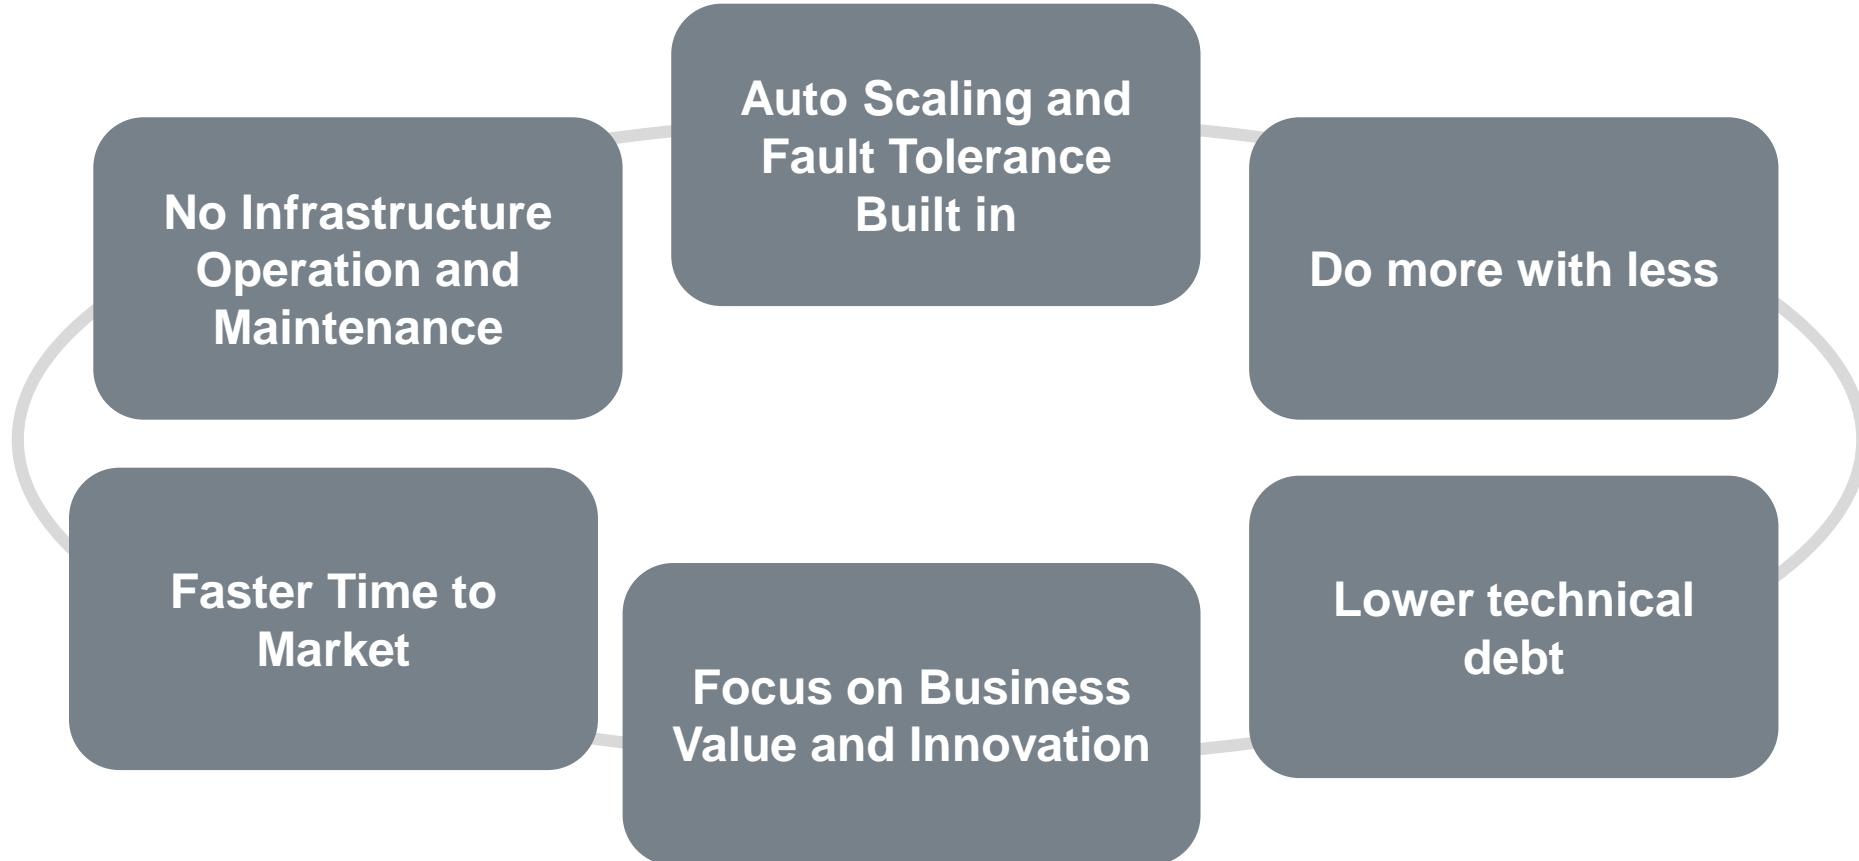
TCO Full Picture



Focus On Business Value and Innovation

Every organization wants exactly this!

TCO Full Picture

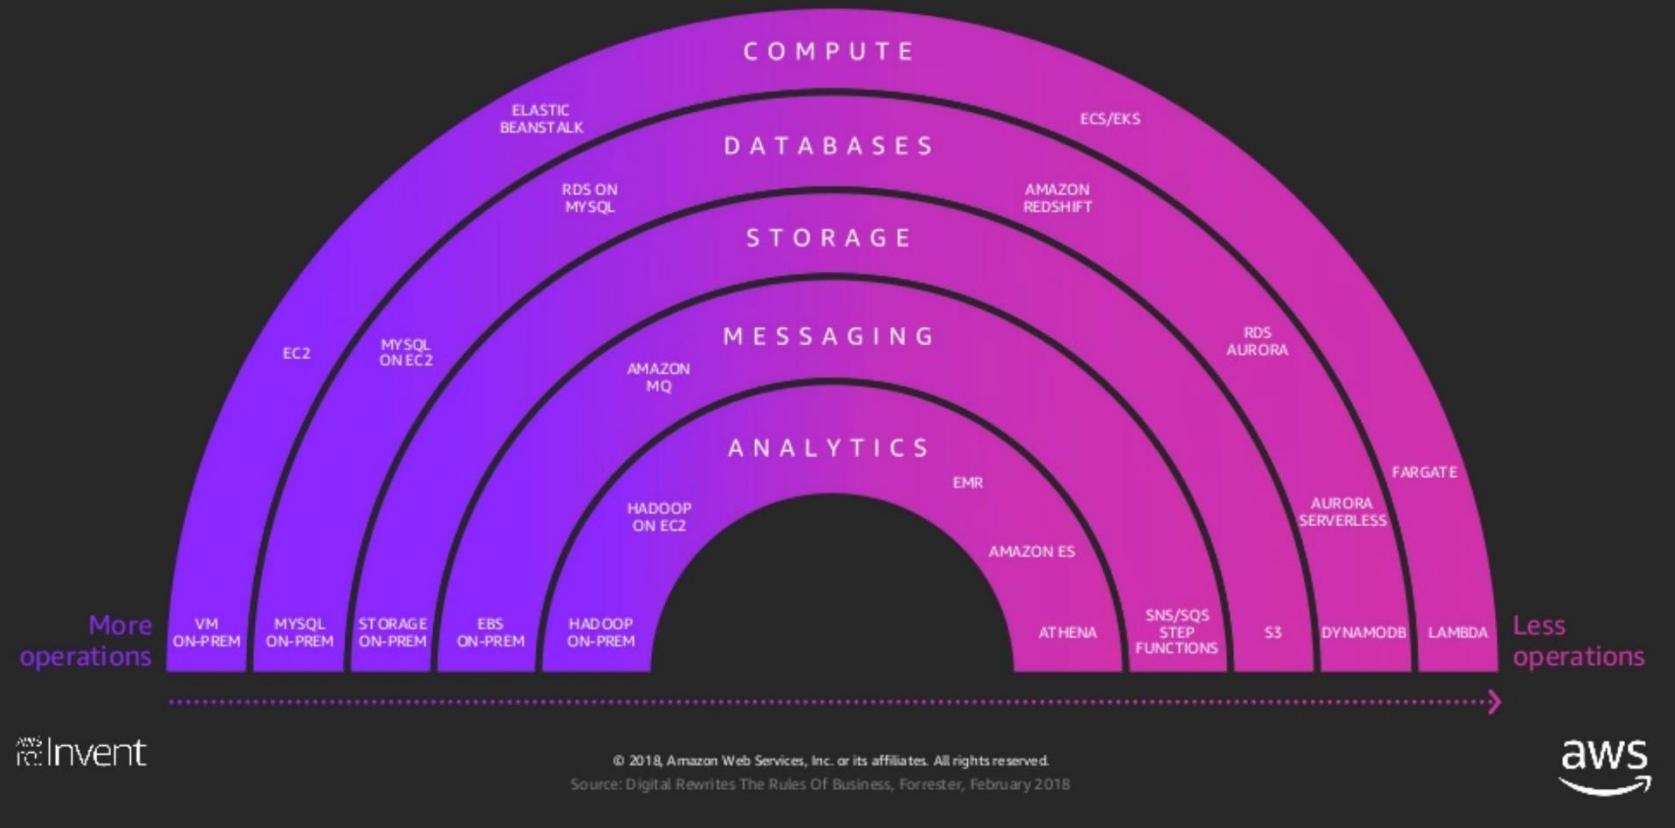


Faster Time To Market

- Time To Market is the key differentiator in today's business!
- Ask yourself: what is core for your business and what you can get as Commodity +(Utility) as a Service?

Serverless Mindset at ip.labs

Serverless is an **operational construct**



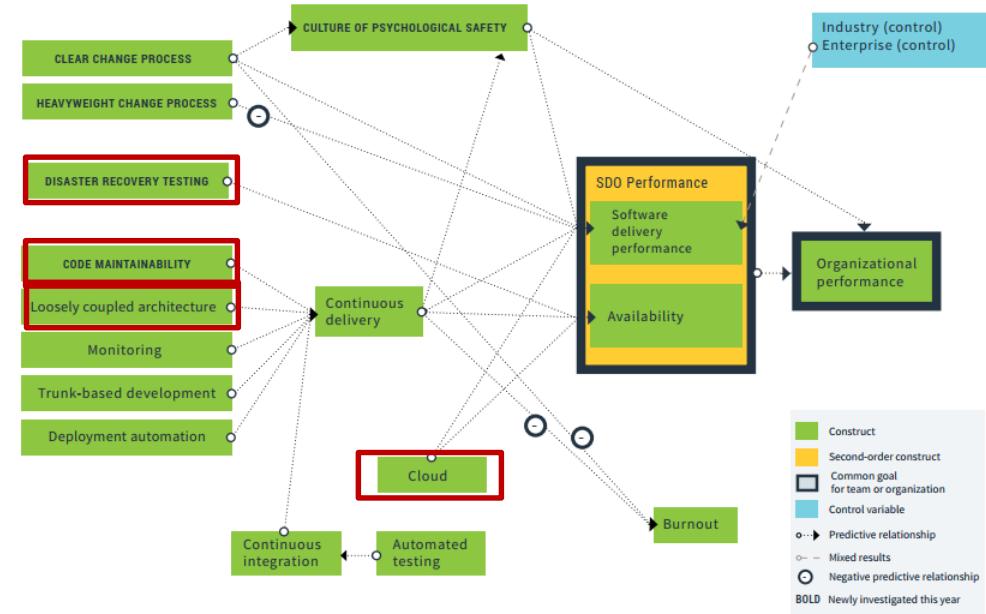
How to measure success

Aspect of Software Delivery Performance*	Elite	High	Medium	Low
Deployment frequency For the primary application or service you work on, how often does your organization deploy code to production or release it to end users?	On-demand (multiple deploys per day)	Between once per day and once per week	Between once per week and once per month	Between once per month and once every six months
Lead time for changes For the primary application or service you work on, what is your lead time for changes (i.e., how long does it take to go from code committed to code successfully running in production)?	Less than one day	Between one day and one week	Between one week and one month	Between one month and six months
Time to restore service For the primary application or service you work on, how long does it generally take to restore service when a service incident or a defect that impacts users occurs (e.g., unplanned outage or service impairment)?	Less than one hour	Less than one day ^a	Less than one day ^a	Between one week and one month
Change failure rate For the primary application or service you work on, what percentage of changes to production or released to users result in degraded service (e.g., lead to service impairment or service outage) and subsequently require remediation (e.g., require a hotfix, rollback, fix forward, patch)?	0-15% ^{b,c}	0-15% ^{b,d}	0-15% ^{c,d}	46-60%



Software Delivery and Operational Performance

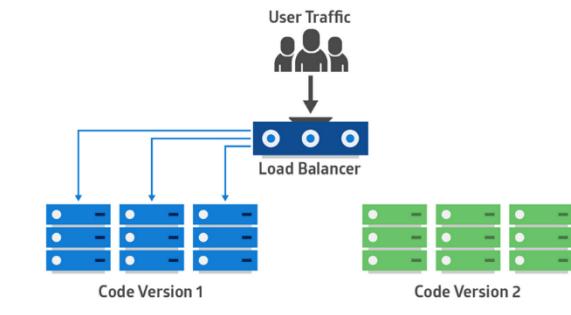
SOFTWARE DELIVERY & OPERATIONAL PERFORMANCE



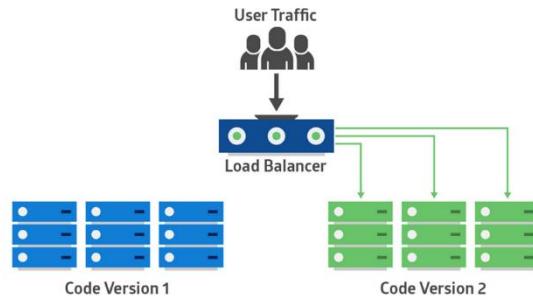
Example: strategies to reduce time to restore service

Aspect of Software Delivery Performance*	Elite	High	Medium	Low
Deployment frequency For the primary application or service you work on, how often does your organization deploy code to production or release it to end users?	On-demand (multiple deploys per day)	Between once per day and once per week	Between once per week and once per month	Between once per month and once every six months
Lead time for changes For the primary application or service you work on, what is your lead time for changes (i.e., how long does it take to go from code committed to code successfully running in production)?	Less than one day	Between one day and one week	Between one week and one month	Between one month and six months
Time to restore service For the primary application or service you work on, how long does it generally take to restore service when a service incident or a defect that impacts users occurs (e.g., unplanned outage or service impairment)?	Less than one hour	Less than one day ^a	Less than one day ^a	Between one week and one month
Change failure rate For the primary application or service you work on, what percentage of changes to production or released to users result in degraded service (e.g., lead to service impairment or service outage) and subsequently require remediation (e.g., require a hotfix, rollback, fix forward, patch)?	0-15% ^{b,c}	0-15% ^{b,d}	0-15% ^{c,d}	46-60%

Blue-Green deployment



Before a blue-green deployment



After a blue-green deployment

Canary deployment

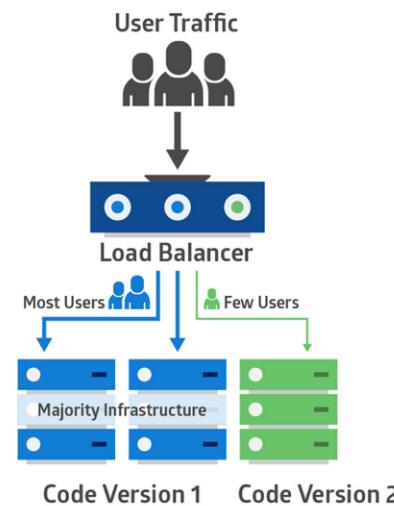
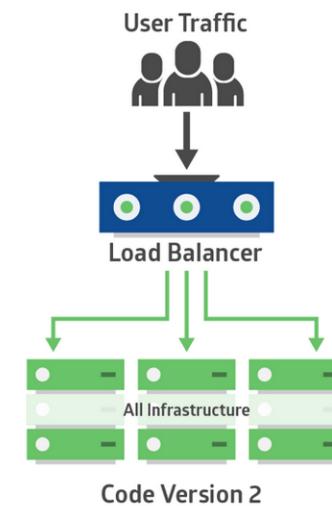


Illustration of a canary deployment



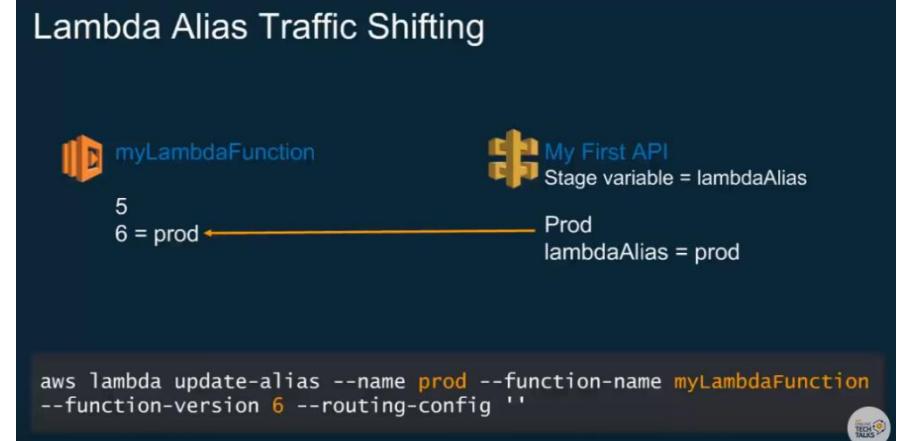
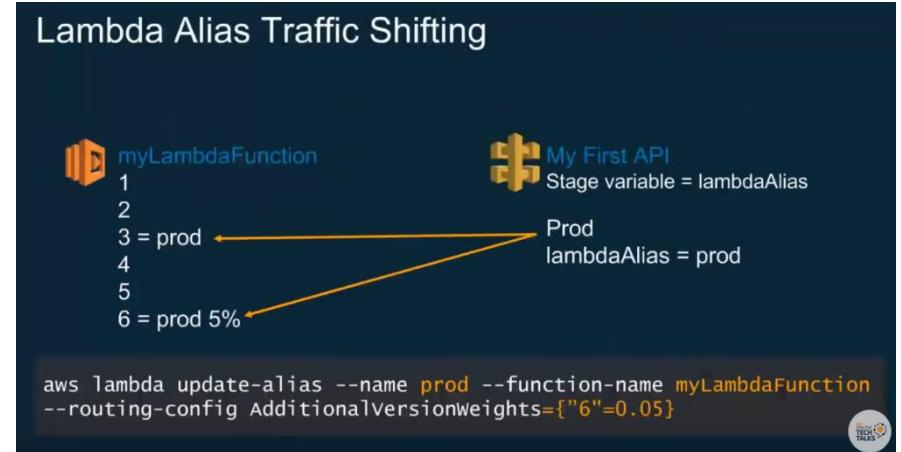
AWS Lambda Deployment

Best Practices:

- API Gateway Stage variables and Lambda Aliases
- Lambda Alias Traffic Shifting

Stage Variables and AWS Lambda Aliases

Using Stage Variables in API Gateway together with Lambda function Aliases you can manage a single API configuration and Lambda function for multiple environment stages



AWS Lambda Deployment Best Practices

- AWS Lambda Alias *Canary and Linear Traffic Shifting & AWS SAM Safe Deployments*
- CloudWatch Rollback Alarms & Lambda hooks

AWS Lambda Alias Traffic Shifting & AWS SAM

In SAM:

AutoPublishAlias

By adding this property and specifying an alias name, AWS SAM will do the following:

- Detect when new code is being deployed based on changes to the Lambda function's Amazon S3 URI.
- Create and publish an updated version of that function with the latest code.
- Create an alias with a name you provide (unless an alias already exists) and points to the updated version of the Lambda function.

Deployment Preference Type

Canary10Percent30Minutes
Canary10Percent5Minutes
Canary10Percent10Minutes
Canary10Percent15Minutes
Linear10PercentEvery10Minutes
Linear10PercentEvery1Minute
Linear10PercentEvery2Minutes
Linear10PercentEvery3Minutes
AllAtOnce

AWS Lambda Alias Traffic Shifting & AWS SAM

In SAM:

Alarms: # A list of alarms that you want to monitor

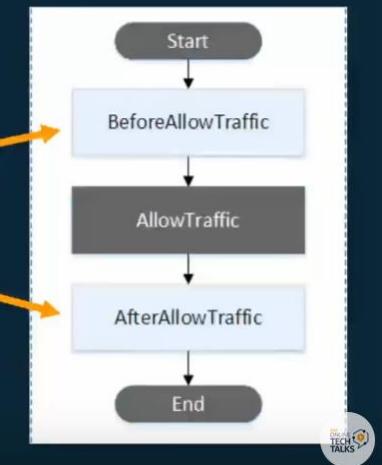
-!Ref AliasErrorMetricGreaterThanZeroAlarm
-!Ref LatestVersionErrorMetricGreaterThanZeroAlarm

Hooks: # Validation Lambda functions that are run before & after traffic shifting

PreTraffic: !Ref PreTrafficLambdaFunction

PostTraffic: !Ref PostTrafficLambdaFunction

Note: You can specify a maximum of 10 alarms



Time Spent

Time Spent	Elite	High	Medium	Low
NEW WORK	50%	50%	40%	30%
Unplanned work and rework	19.5%	20% ^a	20% ^a	20% ^a
Remediating security issues	5%	5% ^b	5% ^b	10%
Working on defects identified by end users	10%	10% ^c	10% ^c	20%
Customer support work	5%	10%	10%	15%



See DORA State of DevOps 2018-2019 Reports

Using Serverless ecosystem will

with the right engineering practices in place will significantly reduce

- extraneous and germane cognitive load
- the amount of code written

How to write less code with AWS Serverless services 1/3

- Write fewer Lambda functions
 - use direct AWS service integrations in case the Lambda only calls the AWS service itself

Less Lambda functions means less:

- code to write, test, run and maintain
- CI/CD (deploy, rollback strategies) to maintain
- Infrastructure as a Code (IAM policies, permission) to write and test
- cold-start worries
- point of failures and retries
- security concerns
- worries about Lambda limits (e.g. concurrency settings per AWS account)
- spending on Lambda, CloudWatch and 3rd party SaaS on (AWS) (monthly) bill



Sheen Brisals "Don't wait for Functionless. Write less Functions instead"

<https://medium.com/lego-engineering/dont-wait-for-functionless-write-less-functions-instead-8f2c331cd651>



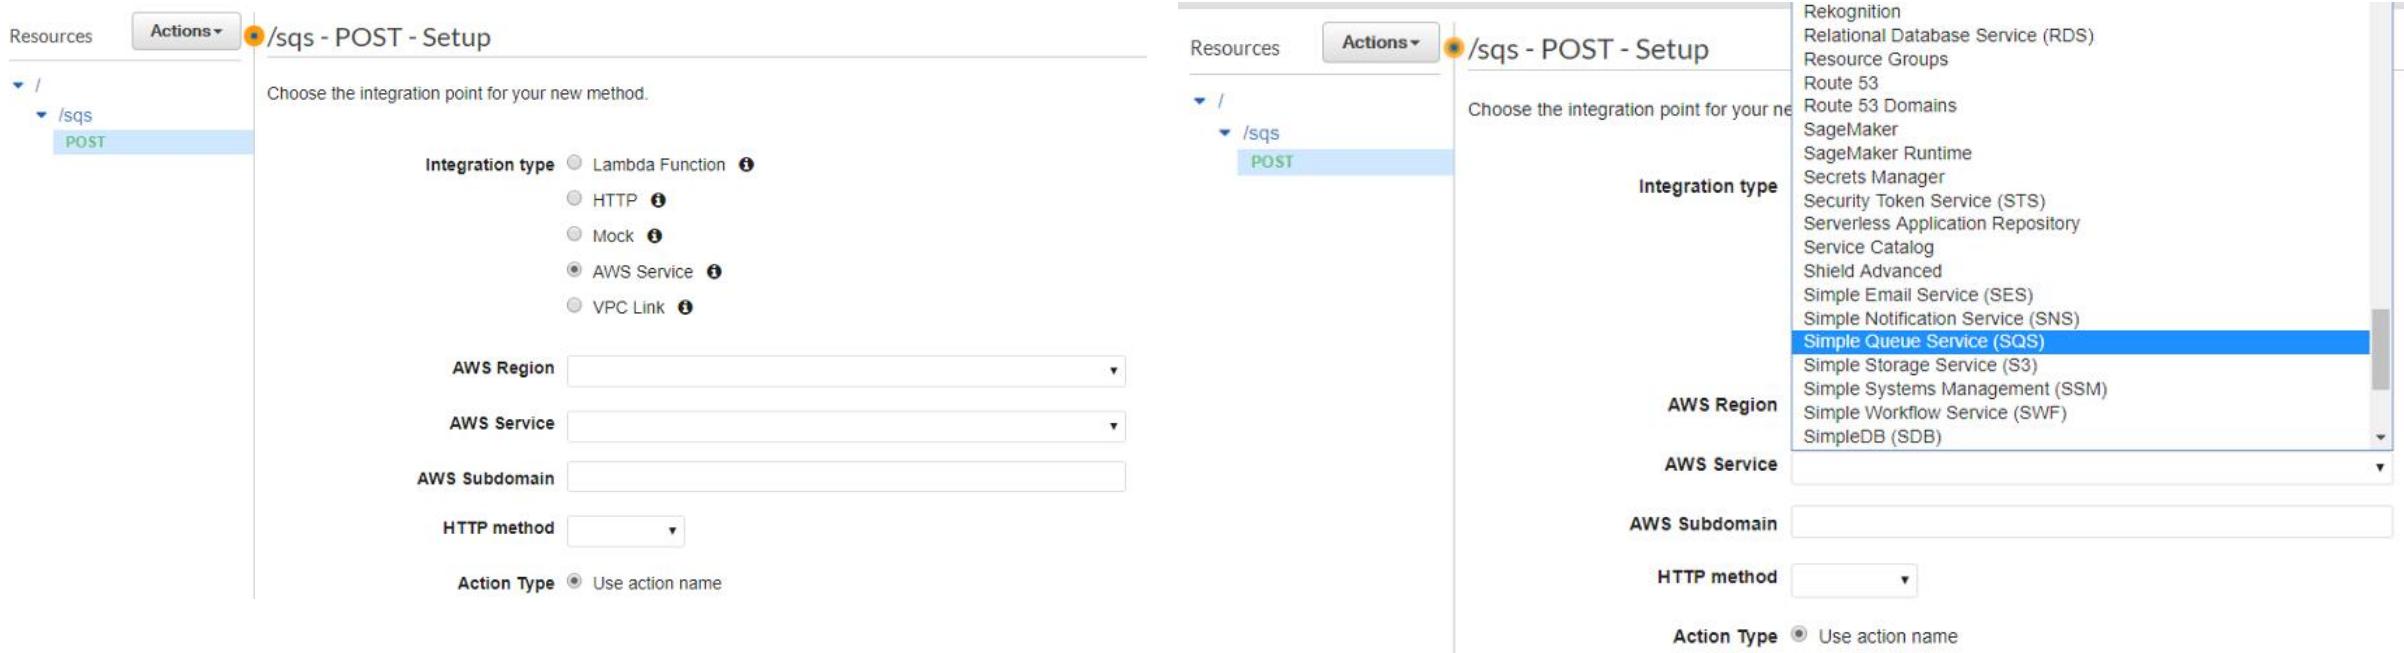
How to write less code with AWS Serverless services 2/3

- Write fewer Lambda functions
 - use more direct service integrations in case the Lambda only calls the Service itself
- Write less Infrastructure as a Code
 - applies not only for Lambda but also to other services like AppSync
- Use (AWS) Services/Frameworks which provide meaningful abstractions
 - Amplify Framework or Serverless Framework Components

How to write less code with AWS Serverless services 3/3

- API Gateway Service Integration
- HTTP APIs Storage-First Service Integration
- Step Functions Service Integration
- Event Bridge Filtering and Routing
- Lambda Destinations
- AppSync and Direct Lambda Resolvers
- Amplify Framework
- Serverless Framework Components

API Gateway Service Integration



The screenshot shows two side-by-side views of the AWS API Gateway 'Actions' interface for setting up a new method.

Left View: The path is /sqss - POST - Setup. The 'Integration type' dropdown is set to 'AWS Service'. Other options include Lambda Function, HTTP, Mock, and VPC Link. Below it are fields for 'AWS Region', 'AWS Service', 'AWS Subdomain', 'HTTP method' (set to POST), and 'Action Type' (set to 'Use action name').

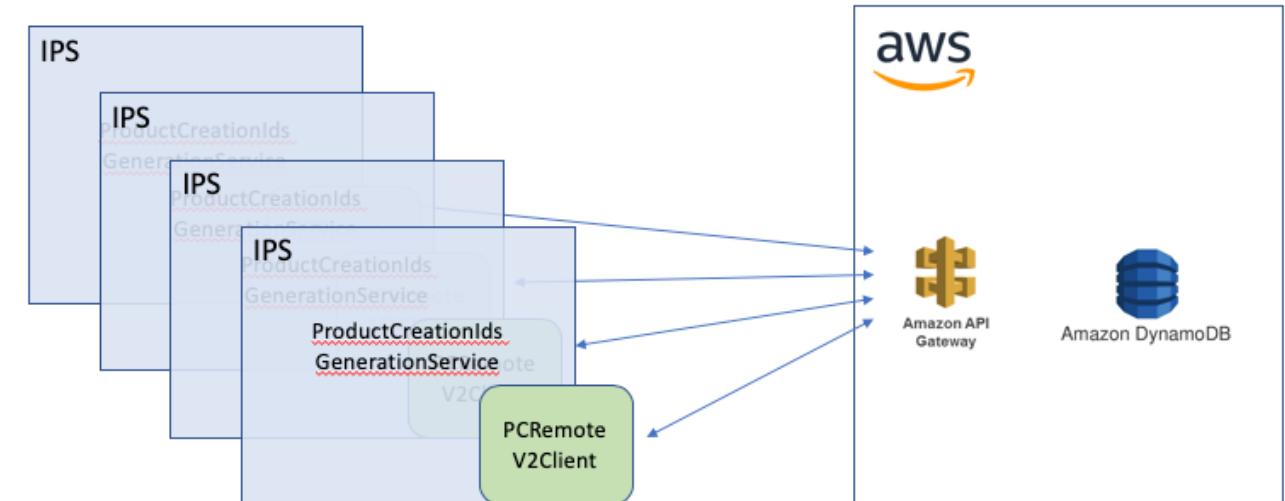
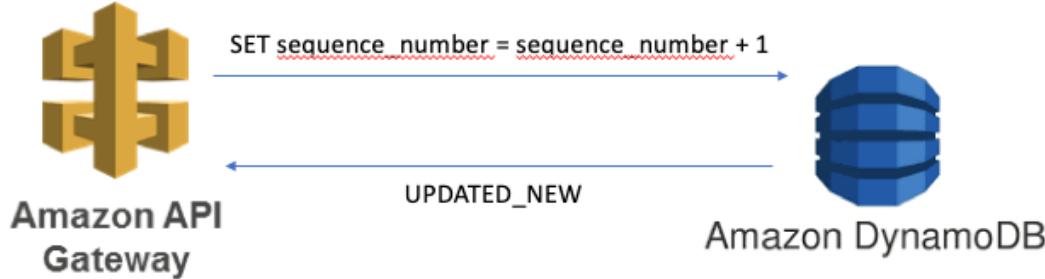
Right View: The path is /sqss - POST - Setup. The 'Integration type' dropdown is also set to 'AWS Service'. The 'AWS Service' dropdown is expanded, showing a list of services. The 'Simple Queue Service (SQS)' option is highlighted with a blue selection bar.

Services List (Right Side):

- Rekognition
- Relational Database Service (RDS)
- Resource Groups
- Route 53
- Route 53 Domains
- SageMaker
- SageMaker Runtime
- Secrets Manager
- Security Token Service (STS)
- Serverless Application Repository
- Service Catalog
- Shield Advanced
- Simple Email Service (SES)
- Simple Notification Service (SNS)
- Simple Queue Service (SQS)** (highlighted)
- Simple Storage Service (S3)
- Simple Systems Management (SSM)
- Simple Workflow Service (SWF)
- SimpleDB (SDB)



API Gateway Service Integration with Dynamo DB : Example ID generator 1/2



API Gateway Service Integration with Dynamo DB : Example ID generator

The screenshot shows the AWS API Gateway interface for setting up a new POST method under the resource '/number'. The 'Integration type' is set to 'AWS Service' (DynamoDB), and the 'Action' is 'UpdateItem'. The 'Content-Type' mapping template is defined as 'application/json'.

Mapping templates

Request body passthrough

- When no template matches the request Content-Type header
- When there are no templates defined (recommended)
- Never

Content-Type

application/json

Add mapping template

application/json

Generate template:

```
1: {  
2:   "TableName": "order-sequence-number",  
3:   "Key": {  
4:     "id": {  
5:       "S": "100"  
6:     }  
7:   },  
8:   "ExpressionAttributeValues": {  
9:     "val": {  
10:      "N": "1"  
11:    }  
12:  },  
13:  "UpdateExpression": "SET sequence_number = sequence_number +  
14:    :val",  
15:  "ReturnValues": "UPDATED_NEW"
```



Sheen Brisals “Sequence Numbering in Serverless via API Gateway”

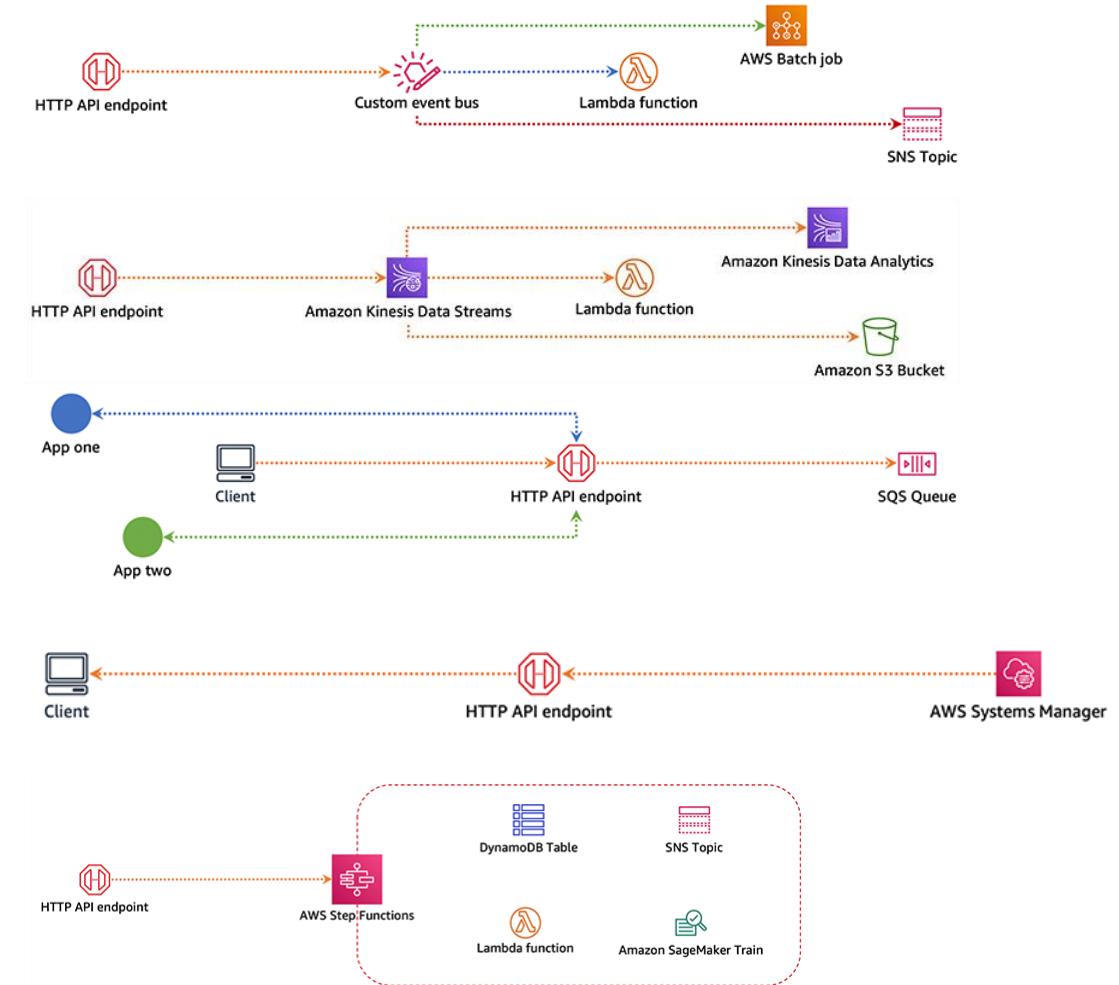
<https://medium.com/lego-engineering/sequence-numbering-in-serverless-via-api-gateway-40e5f6c83e93>

<https://github.com/ToQoz/api-gateway-mapping-template>

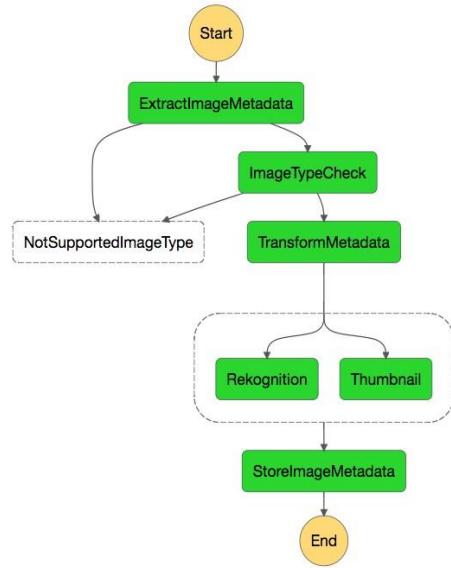


HTTP APIs Storage-First Service Integration

- Event Bridge
- Kinesis Data Streams
- SQS
- AppConfig
- Step Functions



Step Functions Service Integration



Supported Service Integrations			
Service	Request Response	Run a Job (.sync)	Wait for Callback (.waitForTaskToken)
Lambda	✓		✓
AWS Batch	✓	✓	
DynamoDB	✓		
Amazon ECS/AWS Fargate	✓	✓	✓
Amazon SNS	✓		✓
Amazon SQS	✓		✓
AWS Glue	✓	✓	
SageMaker	✓	✓	
Amazon EMR	✓	✓	
AWS CodeBuild	✓	✓	
AWS Step Functions	✓	✓	✓



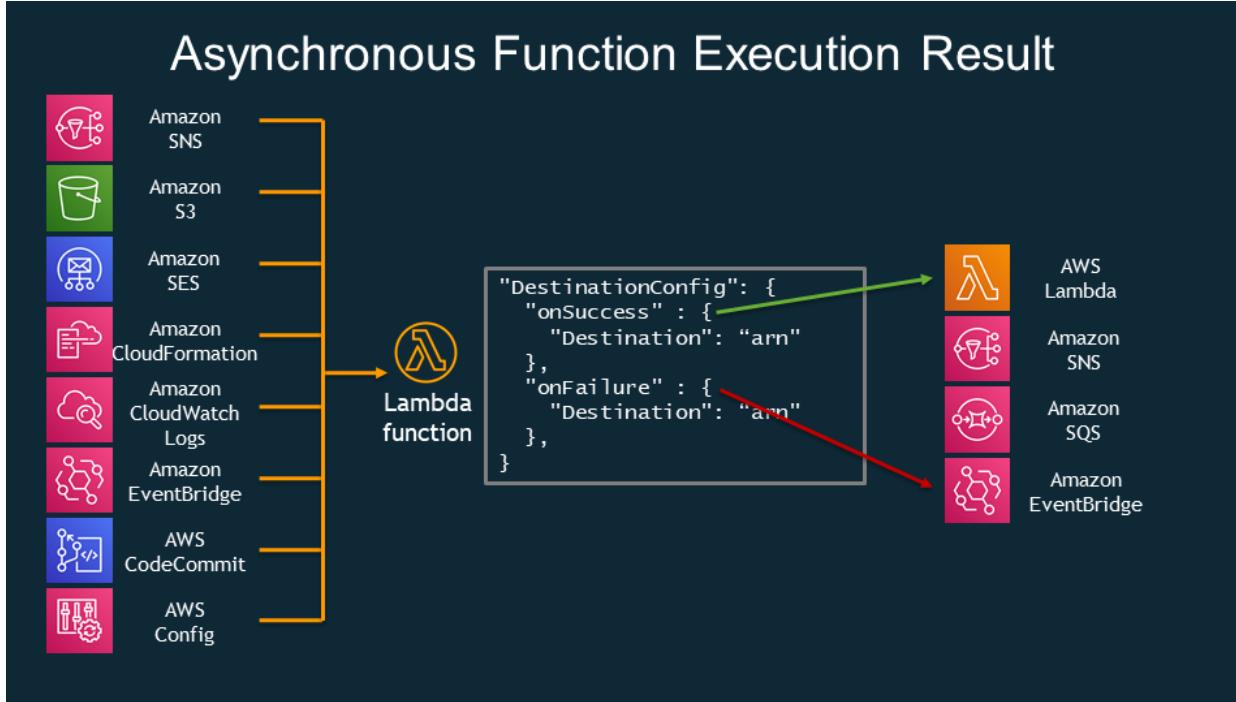
Event Bridge Filtering and Routing



“Reducing custom code by using advanced rules in Amazon EventBridge”

<https://aws.amazon.com/de/blogs/compute/reducing-custom-code-by-using-advanced-rules-in-amazon-eventbridge/>

Lambda Destinations



Add destination

Destination configuration
Send invocation records to a destination when your function is invoked asynchronously, or if your function processes records from a stream.

Source
The type of invocation that maps to the destination.
 Asynchronous invocation
 Stream invocation

Condition
The condition for using the destination.
 On failure
 On success

Destination
An SQS queue, SNS topic, Lambda function, or EventBridge event bus.
arn:aws:sns:sa-east-1:594035263019:event-destinations-test

Designer

event-destinations

Layers (0)

Amazon SNS

+ Add trigger + Add destination

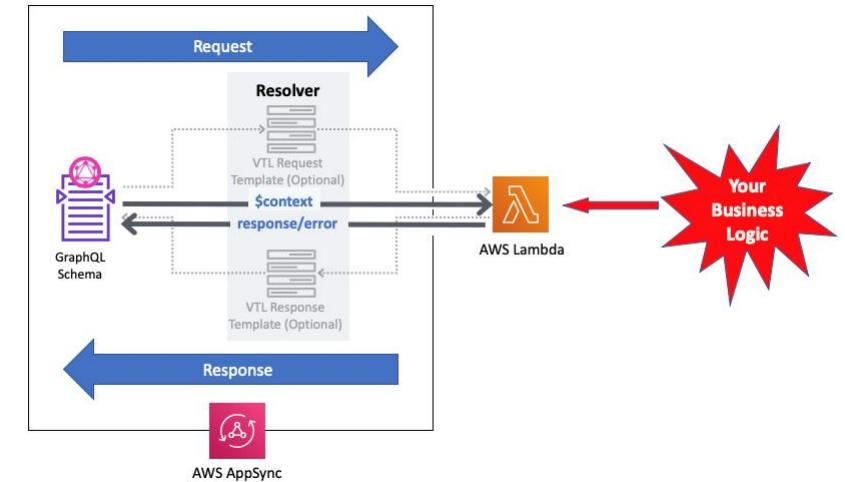
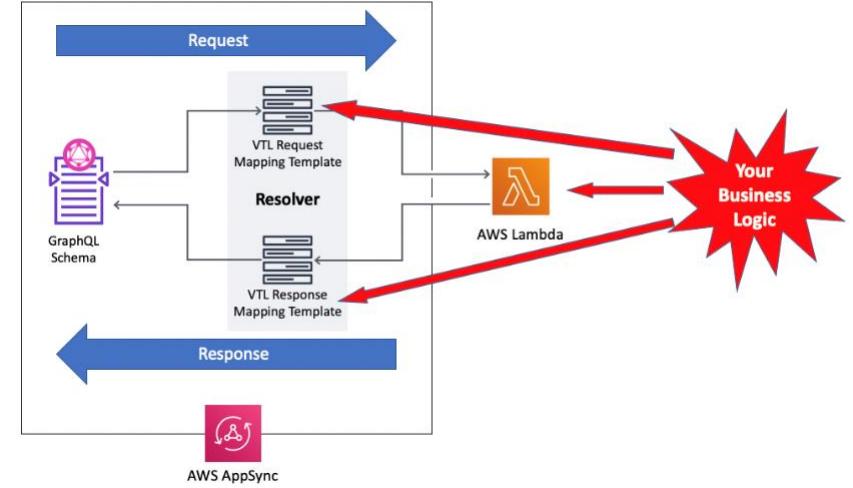
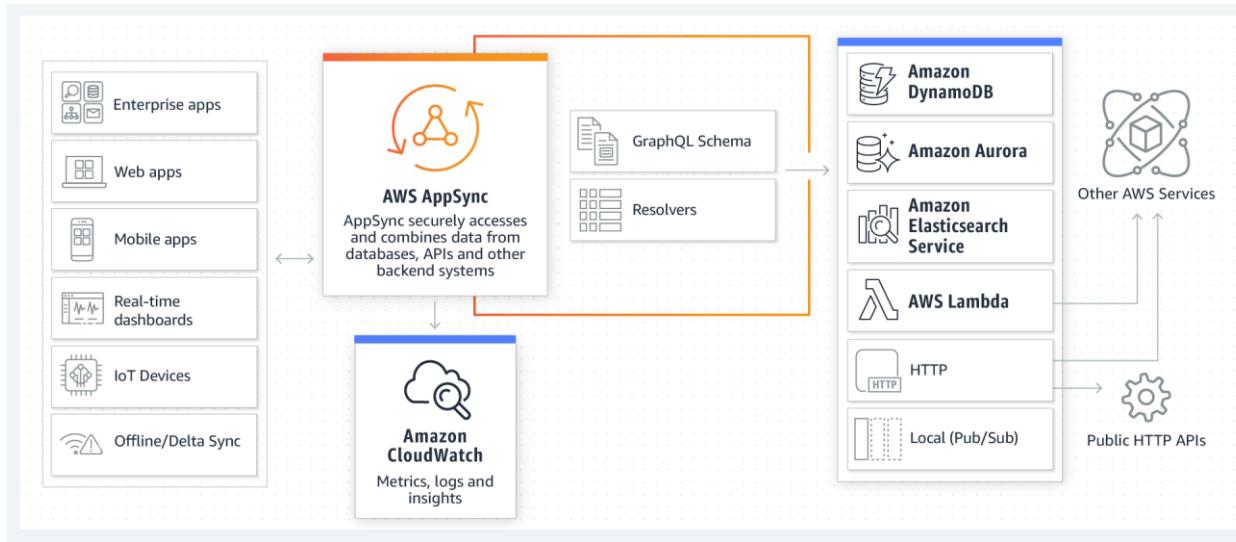
Amazon SNS Info

Source	Stream	Condition	Destination
<input type="radio"/> Asynchronous invocation	-	On success	arn:aws:sns:sa-east-1:594035263019:event-destinations-test



"Introducing AWS Lambda Destinations" <https://aws.amazon.com/de/blogs/compute/introducing-aws-lambda-destinations/>
<https://www.trek10.com/blog/lambda-destinations-what-we-learned-the-hard-way>

AppSync and Direct Lambda Resolvers

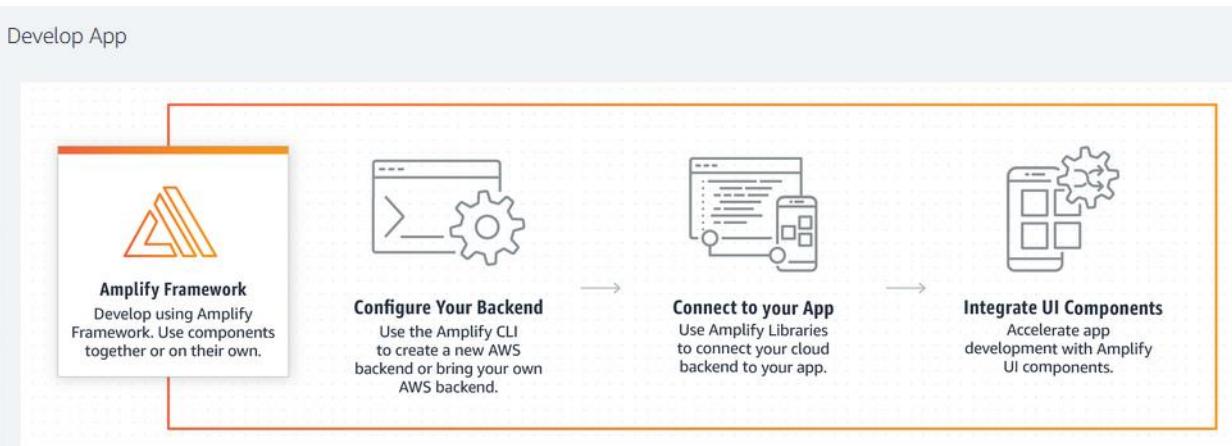


<https://aws.amazon.com/de/appsync>

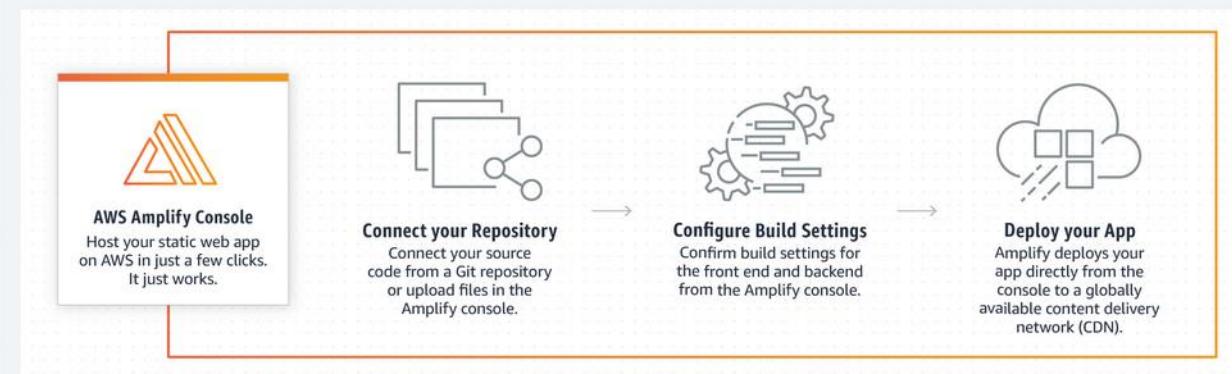
"Introducing Direct Lambda Resolvers: AWS AppSync GraphQL APIs without VTL" <https://aws.amazon.com/de/blogs/mobile/appsync-direct-lambda/>

Amplify Framework

Develop App



Host Web App



https://aws.amazon.com/amplify/?nc1=h_ls

Serverless Framework Components



Serverless Components is now Generally Available. [Click Here for the old Beta Version](#)

English | [简体中文](#)

Serverless Components are simple abstractions that enable developers to deploy serverless applications and use-cases easily, via the [Serverless Framework](#).

- Ease** - Deploy entire serverless applications/use-cases via Components, without being a cloud expert.
- Instant Deployments** - Components deploy in 2-4 seconds, making rapid development on the cloud possible.
- Streaming Logs** - Many Components stream logs from your app to your console in real-time, for fast debugging.
- Automatic Metrics** - Many Components auto-set-up metrics upon deployment.
- Build Your Own** - Components are easy to build.
- Registry** - Share your Components with you, your team, and the world, via the Serverless Registry.



<https://www.serverless.com/blog/what-are-serverless-components-how-use>



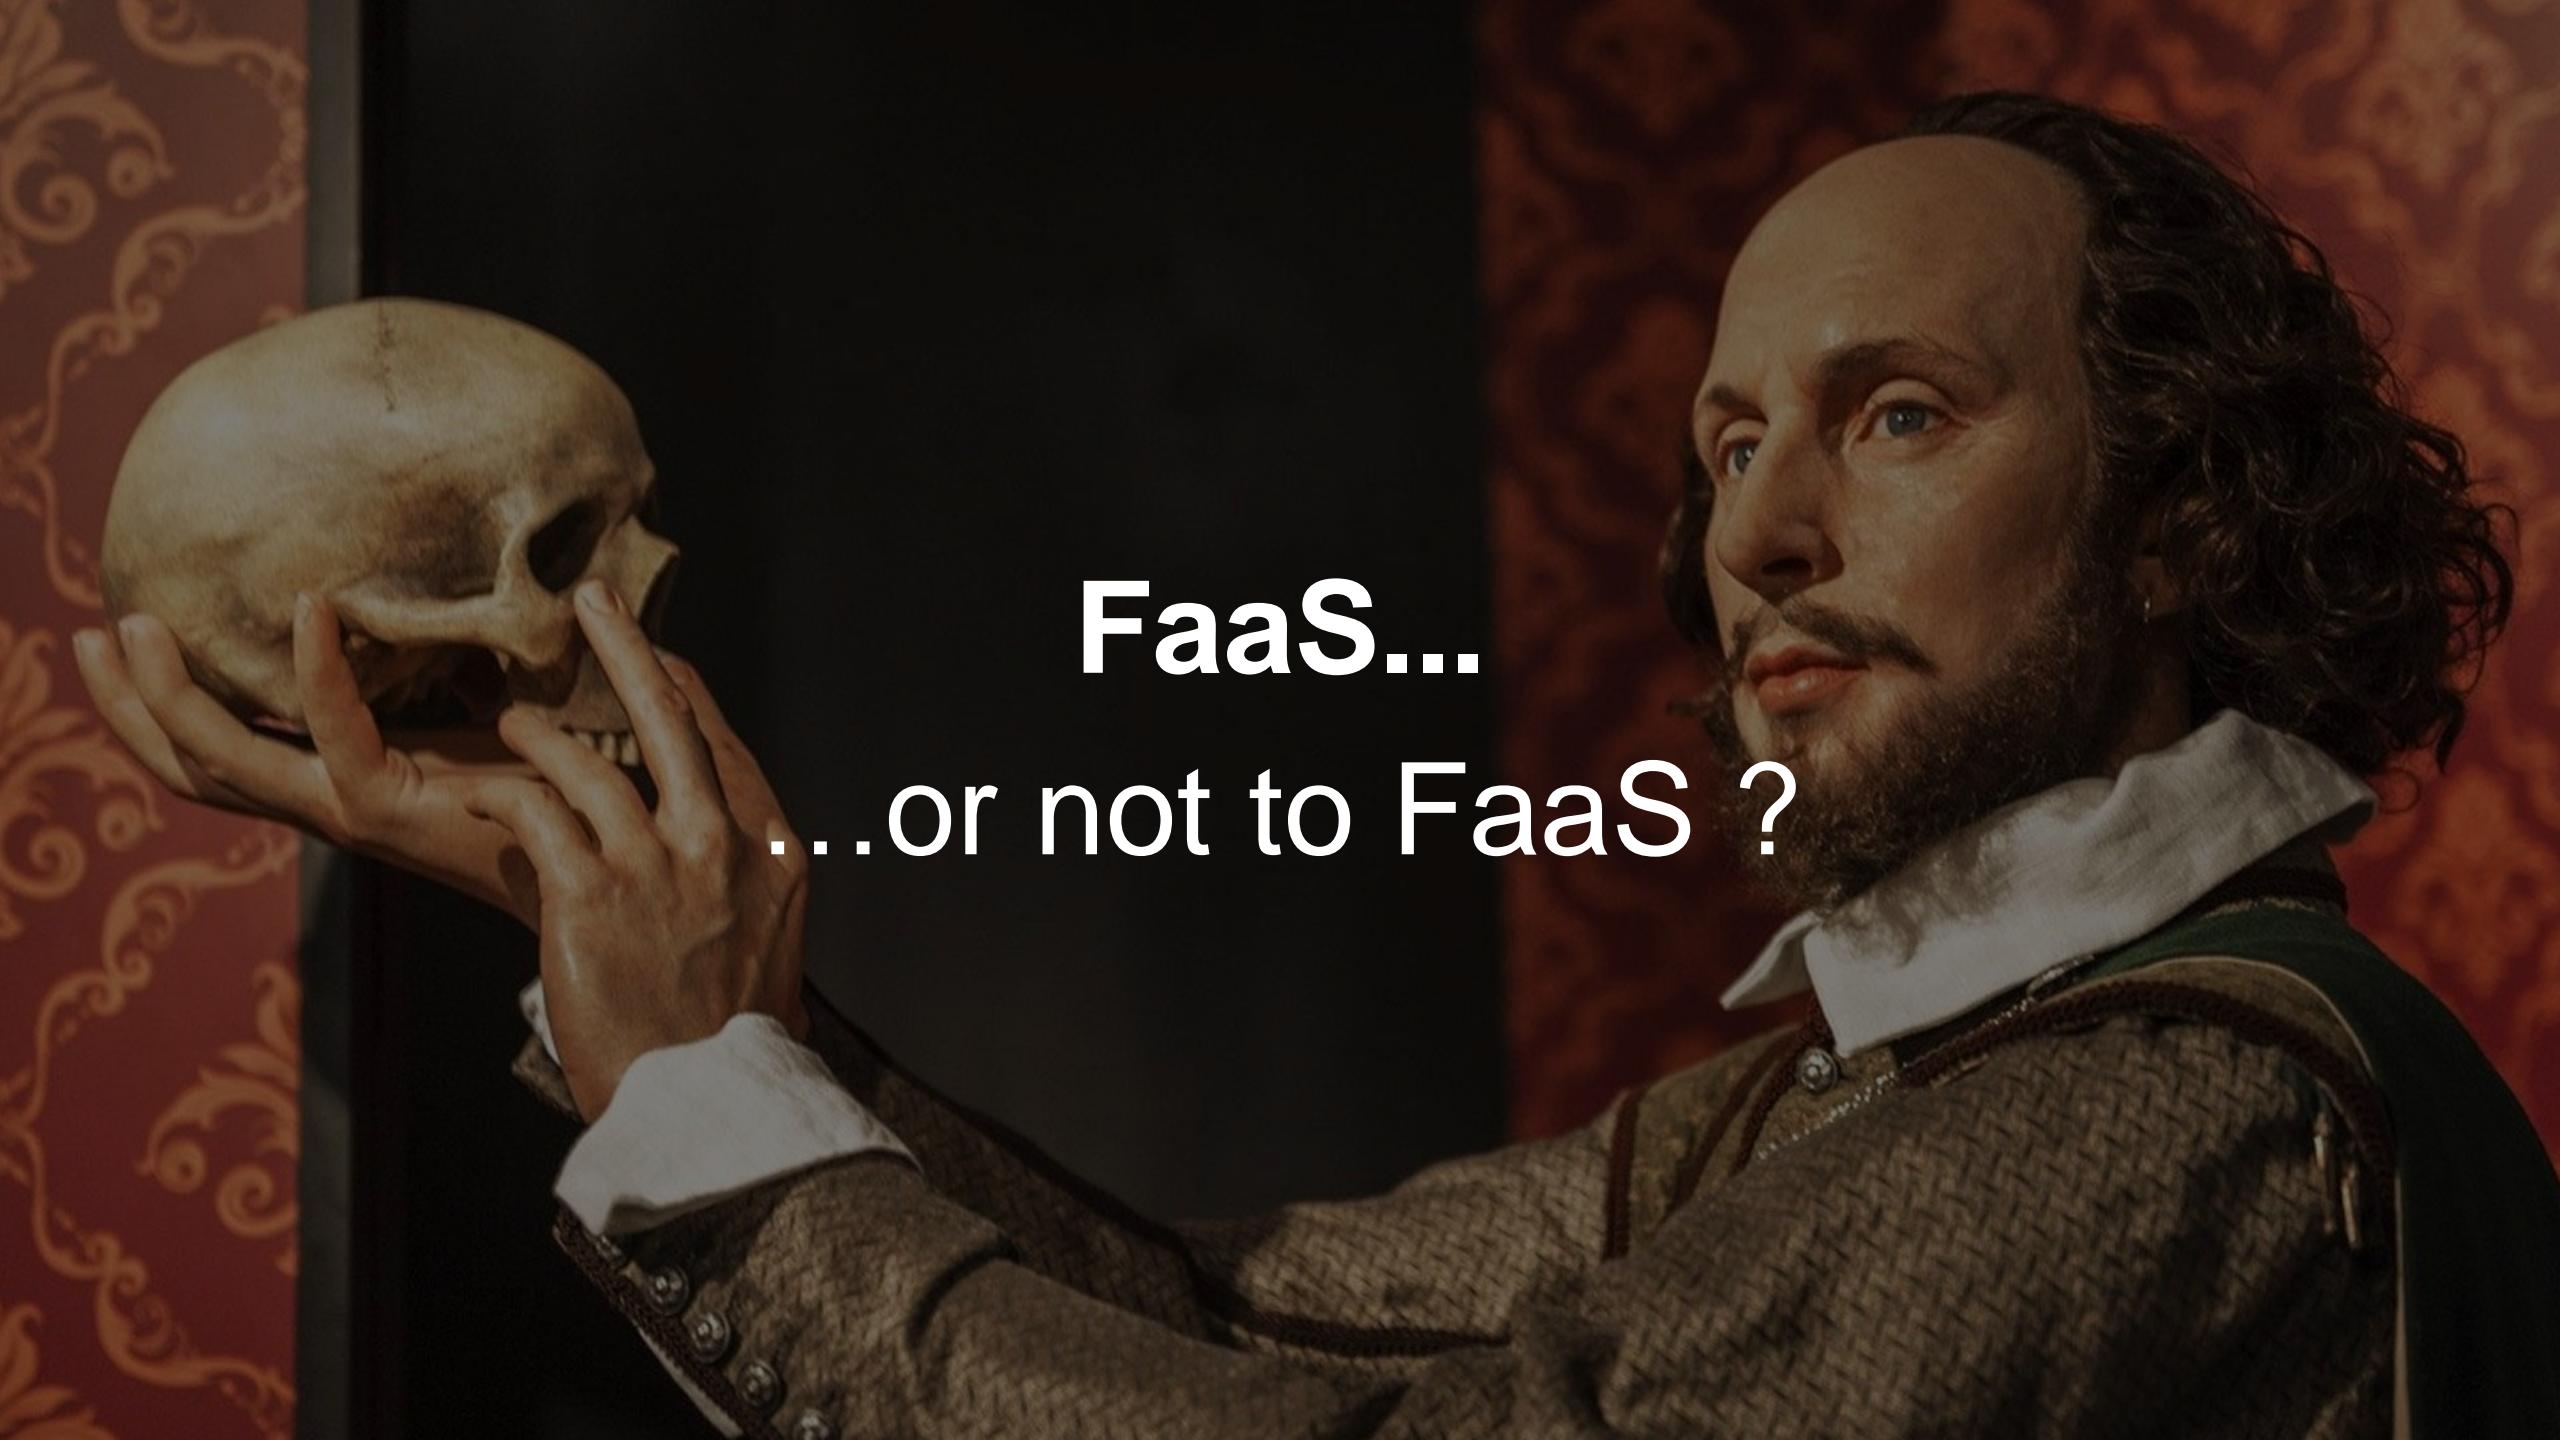
Using Serverless ecosystem will

with the right engineering practices in place will significantly reduce

- extraneous and germane cognitive load
- the amount of code written
- amount of dependencies (no execution runtime and web application server to take care of)

Serverless with the focus on your core domains will enable

- iterative development mind set
- experimentation culture
- focus on business value innovation and faster time to market
- evolutionary architectures

A painting of a man with a beard and curly hair, wearing a brown jacket over a white collared shirt. He is holding a human skull in his hands, looking thoughtfully at it. The background is dark red with gold scrollwork on the left.

FaaS...
...or not to FaaS ?



Image: <https://stock.adobe.com>

Decision Checklist: understand your...

1. Application lifecycle
2. Workloads
3. Programming Model
4. Platform limitations
5. Cost at scale
6. Organizational environment
7. Platform and tooling maturity

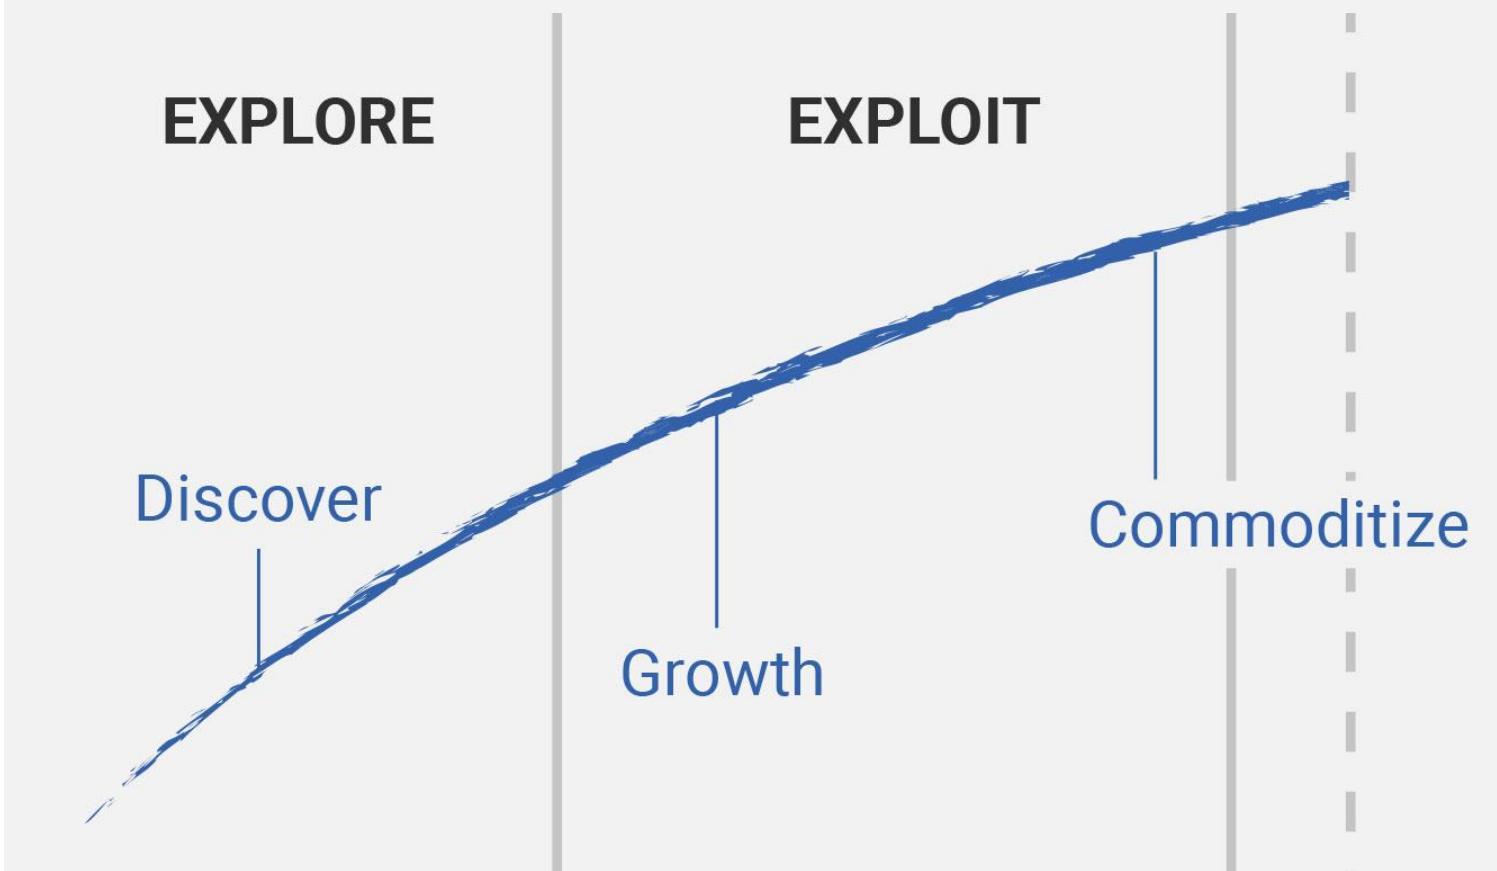


Decision Checklist: understand your...

1. Application lifecycle
2. Workloads
3. Programming Model
4. Platform limitations
5. Cost at scale
6. Organizational environment
7. Platform and tooling maturity



Understand Application lifecycle



Explore phase

- Quickly validate hypotheses
- Rapidly experiment
- Run experiments as cheaply as possible



Serverless is a perfect fit



Image: burst.shopify.com/photos/a-look-across-the-landscape-with-view-of-the-sea

Christian Bannes and Vadym Kazulkin @VKazulkin , ip.labs GmbH



Exploit phase

- Built something that does provide customer value
- Build it on scale
- Build a profitable product around it



partly serverless and partly not serverless architecture



Image: Robert Scoble via Flickr

Christian Bannes and Vadym Kazulkin @VKazulkin , ip.labs GmbH

Application lifecycle

- How much of my stack should I own to be able to deliver business value?
- Outsource SLA, regulatory compliance, price, and roadmap to my service provider?



Existing applications

- You can't magically move that all off to service providers
- You can try to modernize parts of them



Strangler Pattern

- Add a proxy (API Gateway or Application Loadbalancer), which sits between the legacy application and the user
- Add new services and link it to the proxy



Marin Fowler „StranglerFigApplication“ <https://martinfowler.com/bliki/StranglerFigApplication.html>

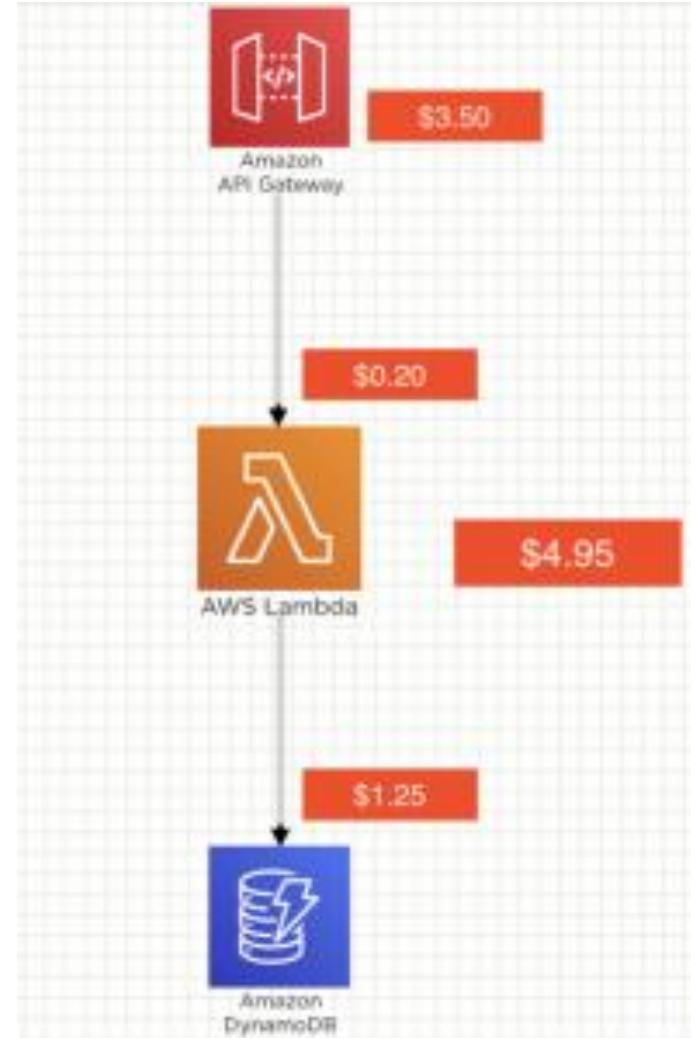
Christian Bannes and Vadym Kazulkin @VKazulkin , ip.labs GmbH



FinDev Concept

Activity-based costing on a digital operation-by-operation basis

- Figure out features which deliver business value comparing to their cost



1. Application lifecycle
2. Workloads
3. Programming Model
4. Platform limitations
5. Cost at scale
6. Organizational environment
7. Platform and tooling maturity

The reality is...

Lambda is often just a small percentage of your total cost

Details

AWS Service Charges	\$29.59
▶ API Gateway	\$2.70
▶ CloudFront	\$0.03
▶ CloudWatch	\$5.08
▶ Data Transfer	\$2.22
▶ DynamoDB	\$4.19
▶ Elasticsearch Service	\$8.45
▶ Key Management Service	\$1.00
▶ Lambda	\$0.25
▶ Simple Storage Service	\$0.93

API Gateway

\$ 3.50

Per million API calls

HTTP APIs in Beta

70% cheaper as API Gateway

- Fewer configuration options
- Well-suited for most use-cases

DynamoDB On-Demand

Read/write capacity mode

Select on-demand if you want to pay only for the read and writes you perform, with no capacity planning required. Select provisioned to save on throughput costs if you can reliably estimate your application's throughput requirements. See the [DynamoDB pricing page](#) and [DynamoDB Developer Guide](#) to learn more.

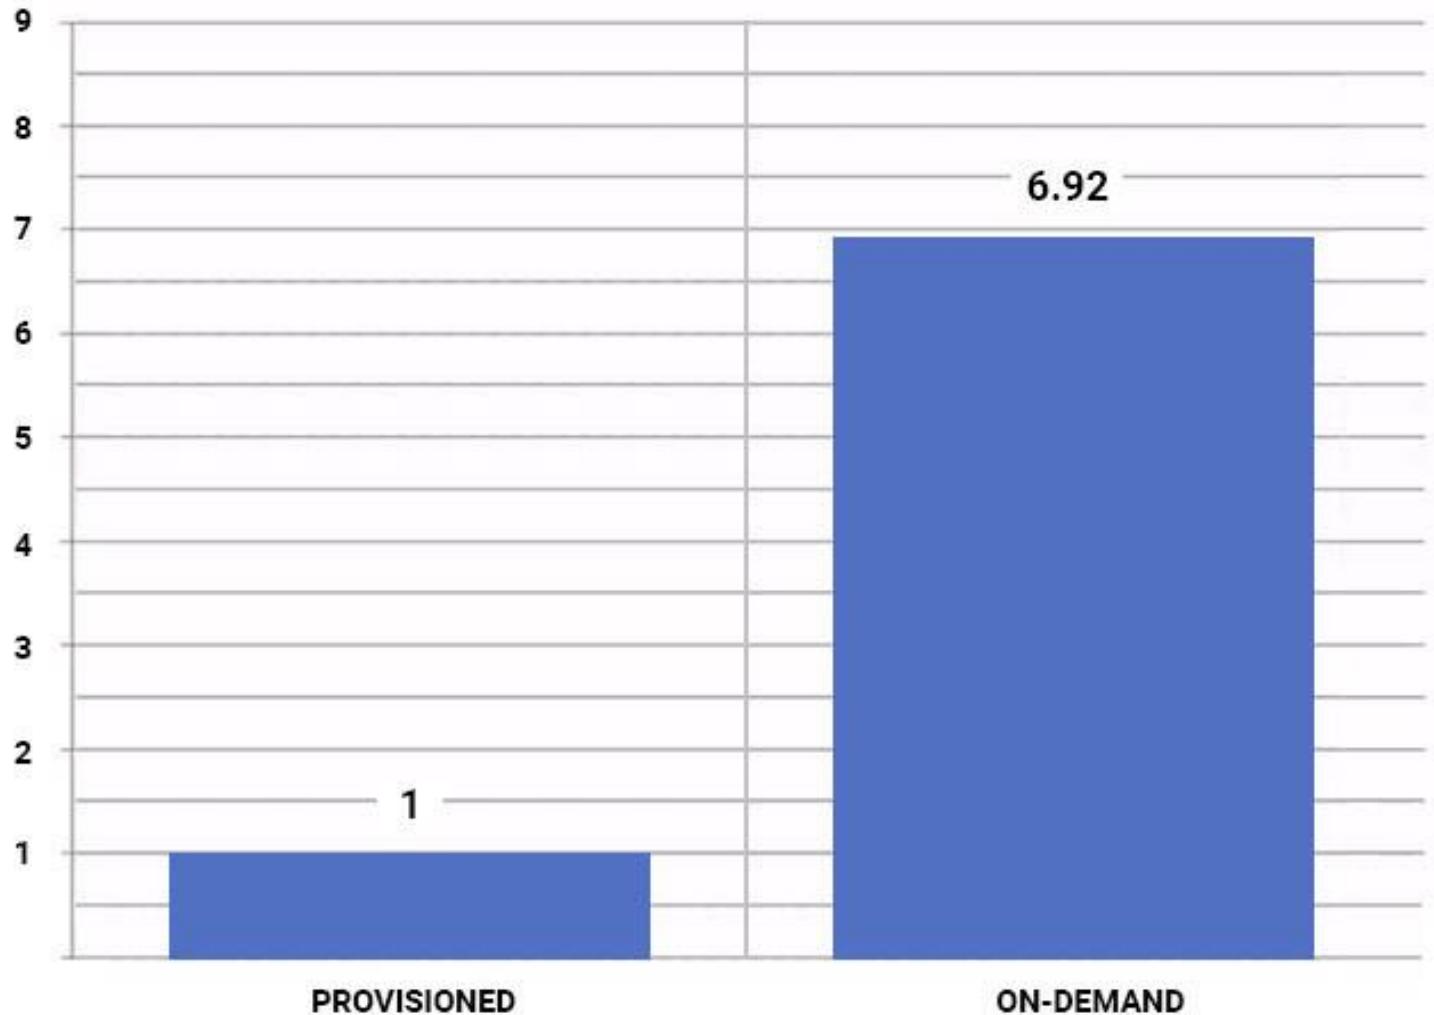
Read/write capacity mode can be changed later.

- Provisioned (free-tier eligible)
- On-demand



Provisioned vs On-Demand

- Use On-Demand for spiky workloads
- Use Provisioned for constantly high workload



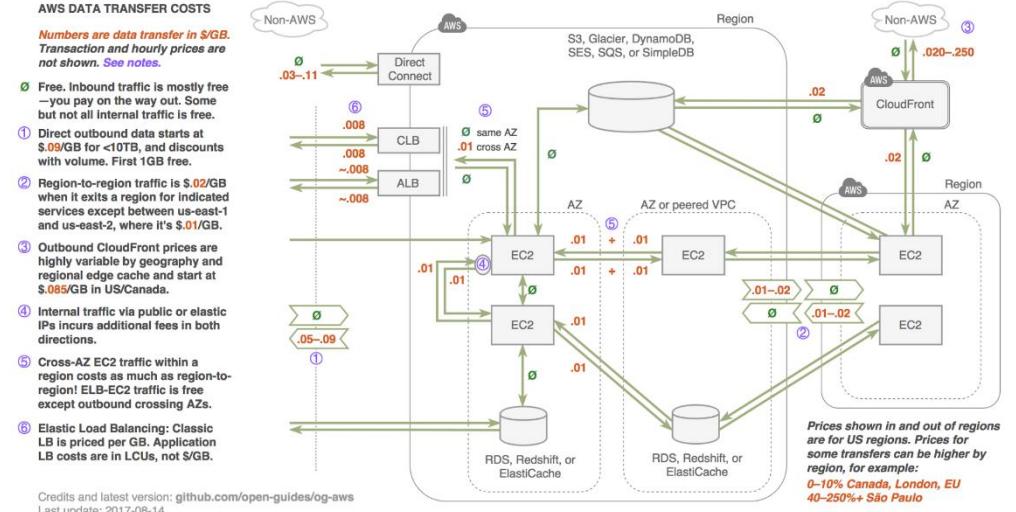
Understand your cost at scale

- Lambda
- API Gateway
- Dynamo DB capacity choices
- Logging costs
- Monitoring costs



Understand your cost at scale

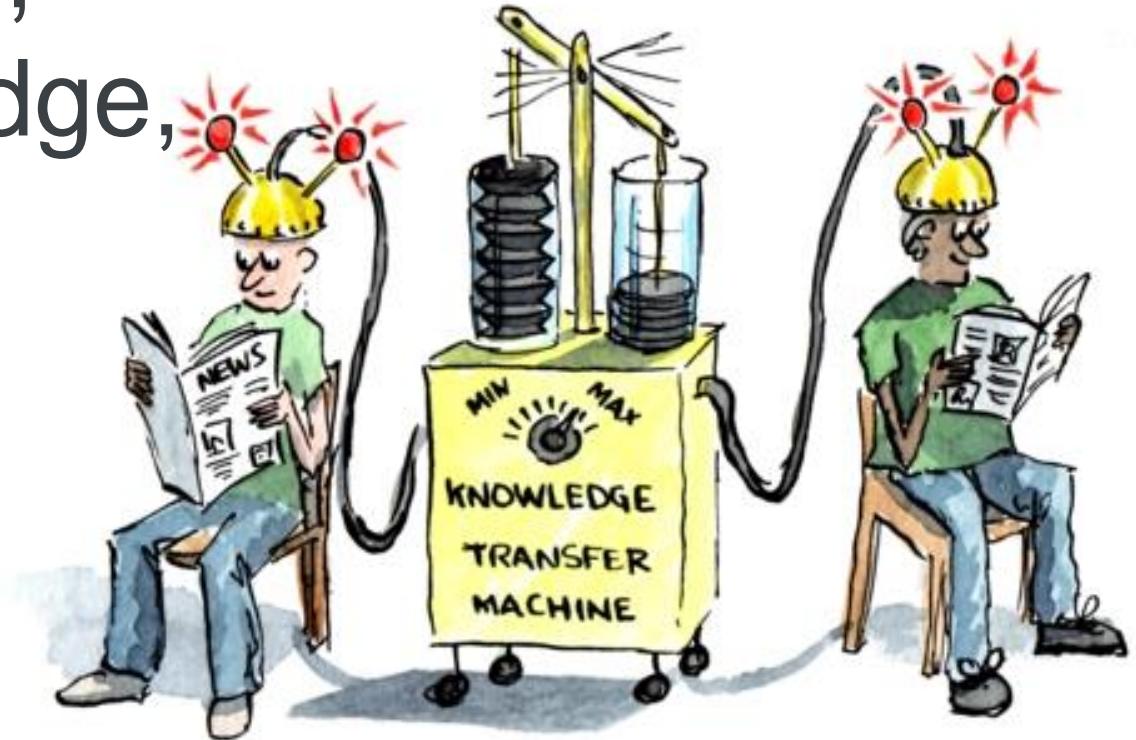
- Data transfer costs
- Step functions
- Remote API calls
- 3rd party Services price models



1. Application lifecycle
2. Workloads
3. Programming Model
4. Platform limitations
5. Cost at scale
6. Organizational environment
7. Platform and tooling maturity

Understand organizational environment (structure, responsibilities, knowledge, culture)

- Do you already embrace DevOps best practices?



Charity Majors: „The Future of Ops careers“

- Advocates for the internal observability team (even if you pay for SaaS observability solution and you use Serverless-first approach)
 - team should write libraries, generate examples, and drive standardization; ushering in consistency, predictability, and usability
 - team should partner with internal teams to evaluate use cases. They might also write glue code and helper modules to connect different data sources and create cohesive visualizations
 - team becomes an integration point between your organization and the outsourced work



Tom McLaughlin Talk: „What do we do when the server goes away?“

- Observability (Monitoring, Logging, Tracing & Alerting)
- Chaos Engineering & Game Days
- Infrastructure as Code & Testing
- Help understand constraints of AWS services & choose the right one



Tom McLaughlin „What do we do when the server goes away“
<https://speakerdeck.com/tmclaugh/serverless-devops-what-do-we-do-when-the-server-goes-away>

Christian Bannes and Vadym Kazulkin @VKazulkin , ip.labs GmbH



Help understand constraints of AWS services & choose the right one. Example **Event Sources:**



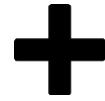
Kinesis Streams



SNS

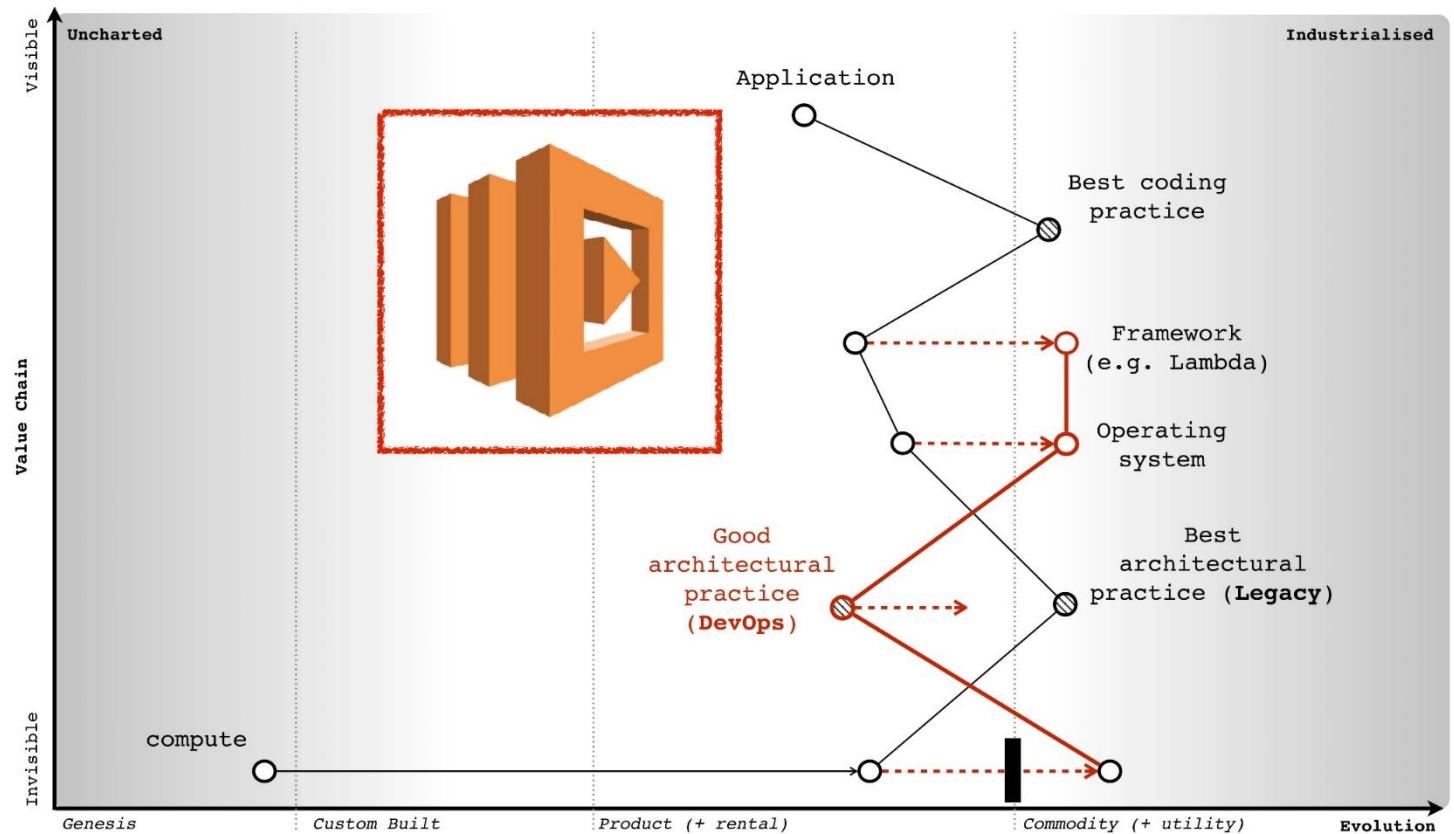


SQS



AWS EventBridge

Wardley Map



Co-evolution of practices with Serverless 1/2

- True DevOps (even DevSecOps)
- FinDev responsibilities in the teams
- Complete infrastructure automation
- Chaos Engineering



Type 2: Fully Shared Ops Responsibilities

Where operations people have been integrated in product development teams, we see a Type 2 topology. There is so little separation between Dev and Ops that all people are highly focused on a shared purpose; this is arguably a form of [Type 1 \(Dev and Ops Collaboration\)](#), but it has some special features.

Organisations such as Netflix and Facebook with effectively a single web-based product have achieved this Type 2 topology, but I think it's probably not hugely applicable outside a narrow product focus, because the budgetary constraints and context-switching typically present in an organisation with multiple product streams will probably force Dev and Ops further apart (say, back to a [Type 1 model](#)). This topology might also be called 'NoOps', as there is no distinct or visible Operations team (although the Netflix NoOps might also be [Type 3 \(Ops as IaaS\)](#)).



Co-evolution of practices with Serverless 2/2

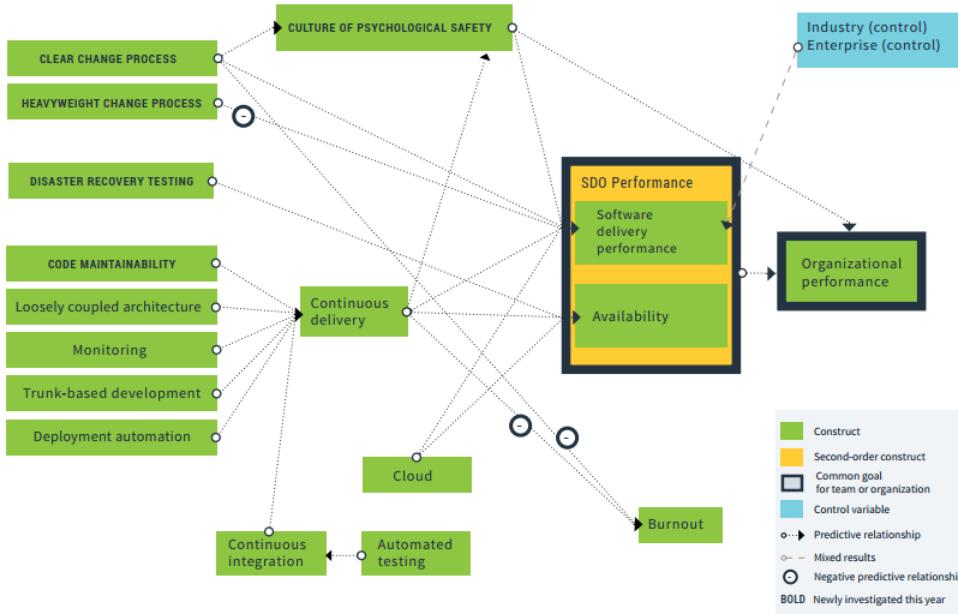
- Each team or even developer can get its own (AWS) test account
- No local testing environment or only for quick functional tests
- Testing in production



Michael Bryzek “What do you know about testing in production?” <https://www.youtube.com/watch?v=z-ATZTUgaAo>

Invest in Software Delivery and Operational Performance Excellence

SOFTWARE DELIVERY & OPERATIONAL PERFORMANCE



1. Application lifecycle
2. Workloads
3. Programming Model
4. Platform limitations
5. Cost at scale
6. Organizational environment
7. Platform and tooling maturity

Serverless platform and tooling maturity

- Infrastructure-as-a-Code solutions maturity (e.g. Cloud Formation, CDK, Terraform)
- Development environment & framework maturity (e.g. AWS SAM, AWS Amplify, Serverless Framework)
- Integration with 3rd party SaaS

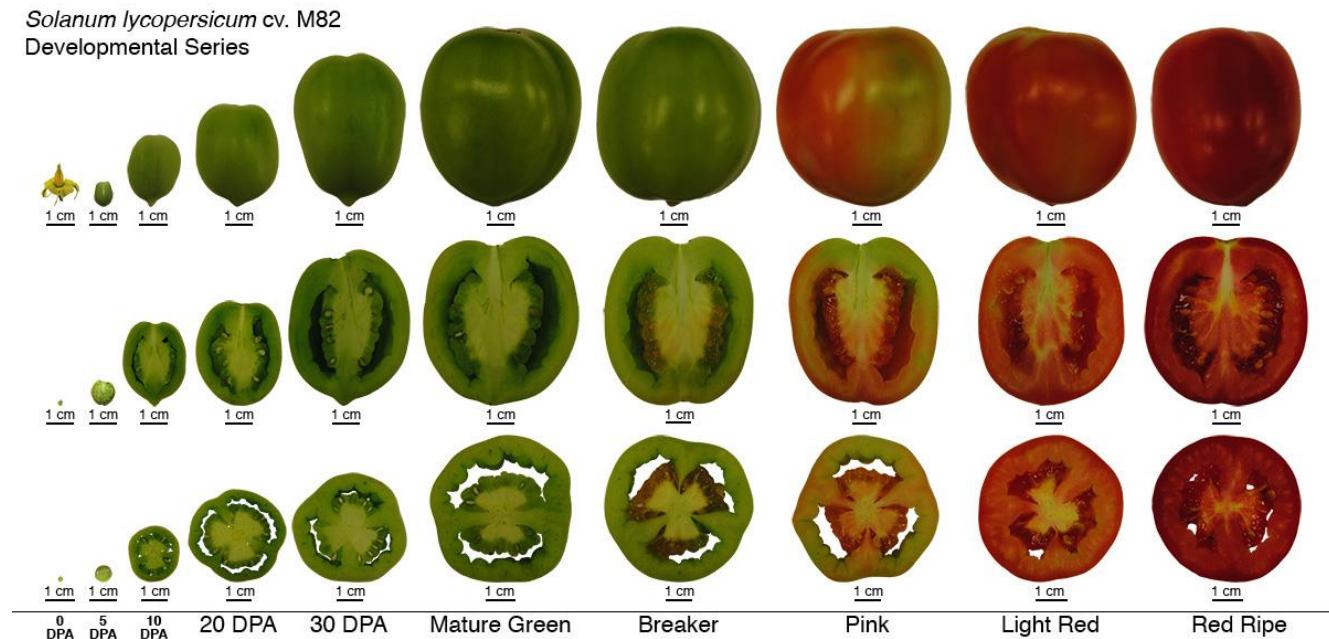


Image: http://tea.solgenomics.net/anatomy_viewer/microscopy/slm82_fruit

Christian Bannes and Vadym Kazulkin @VKazulkin , ip.labs GmbH



Serverless platform and tooling maturity

- CI/CD
- Observability (Logging, Monitoring, Tracing)
- Alerting
- Security

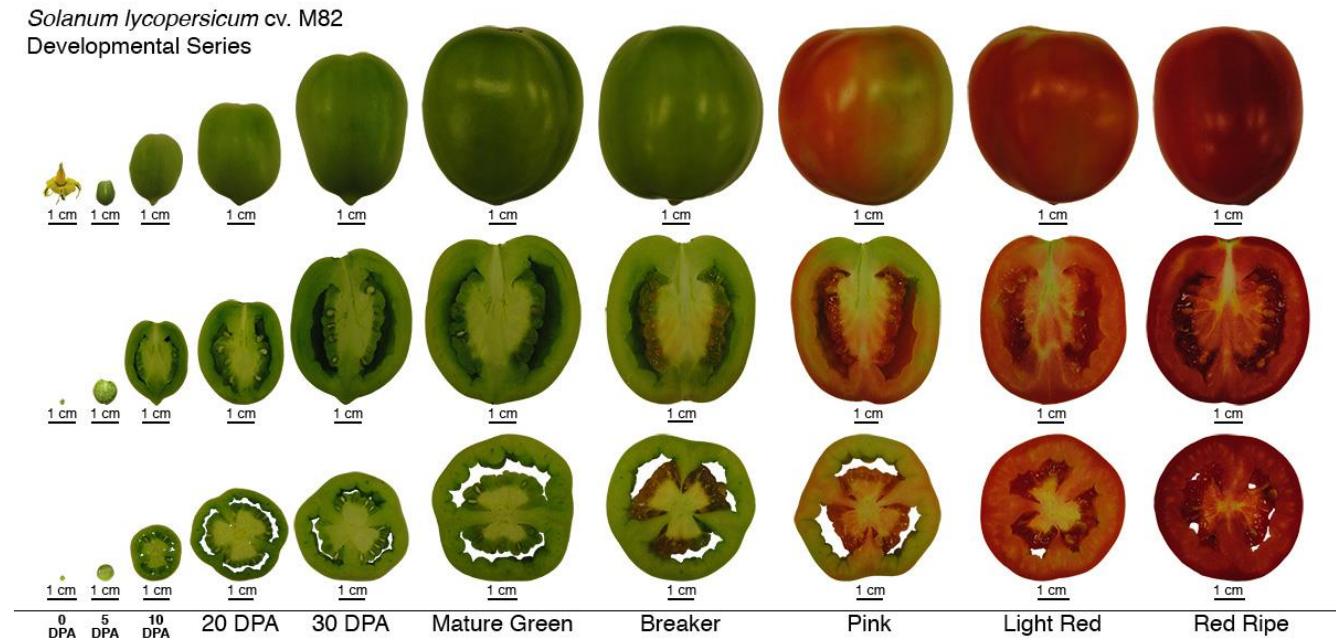


Image: http://tea.solgenomics.net/anatomy_viewer/microscopy/slm82_fruit

Christian Bannes and Vadym Kazulkin @VKazulkin , ip.labs GmbH







Thank You!